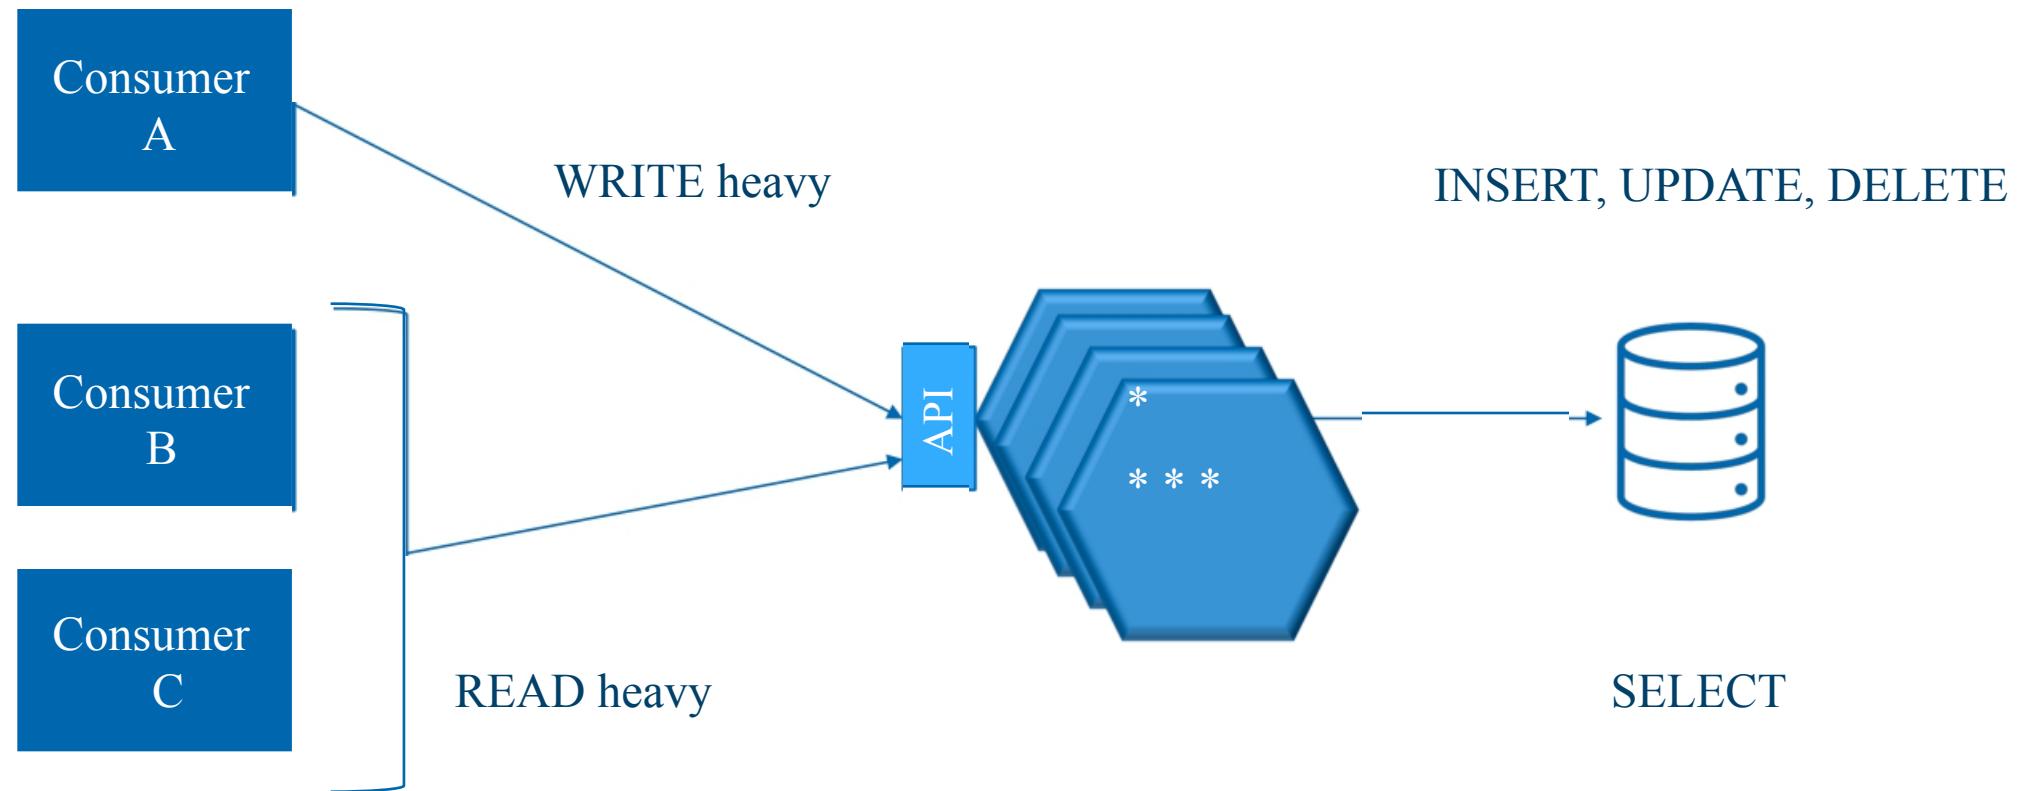
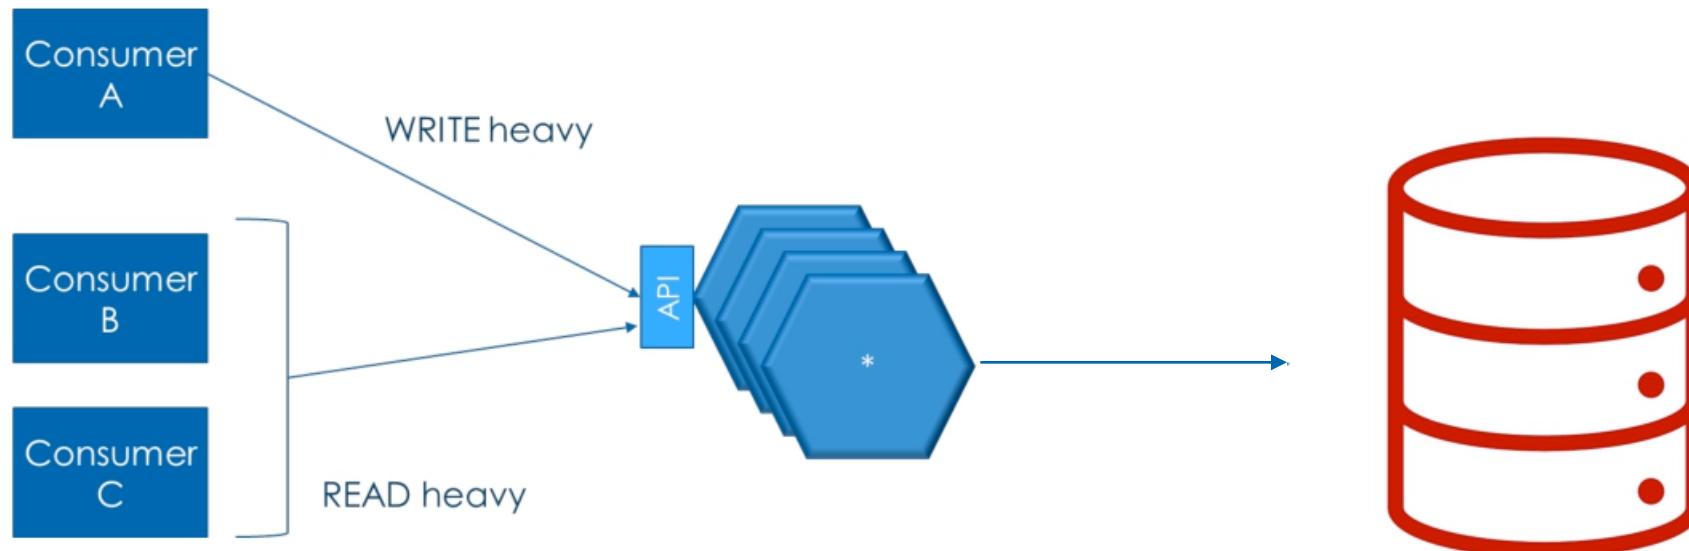


# Performance & Data Integrity

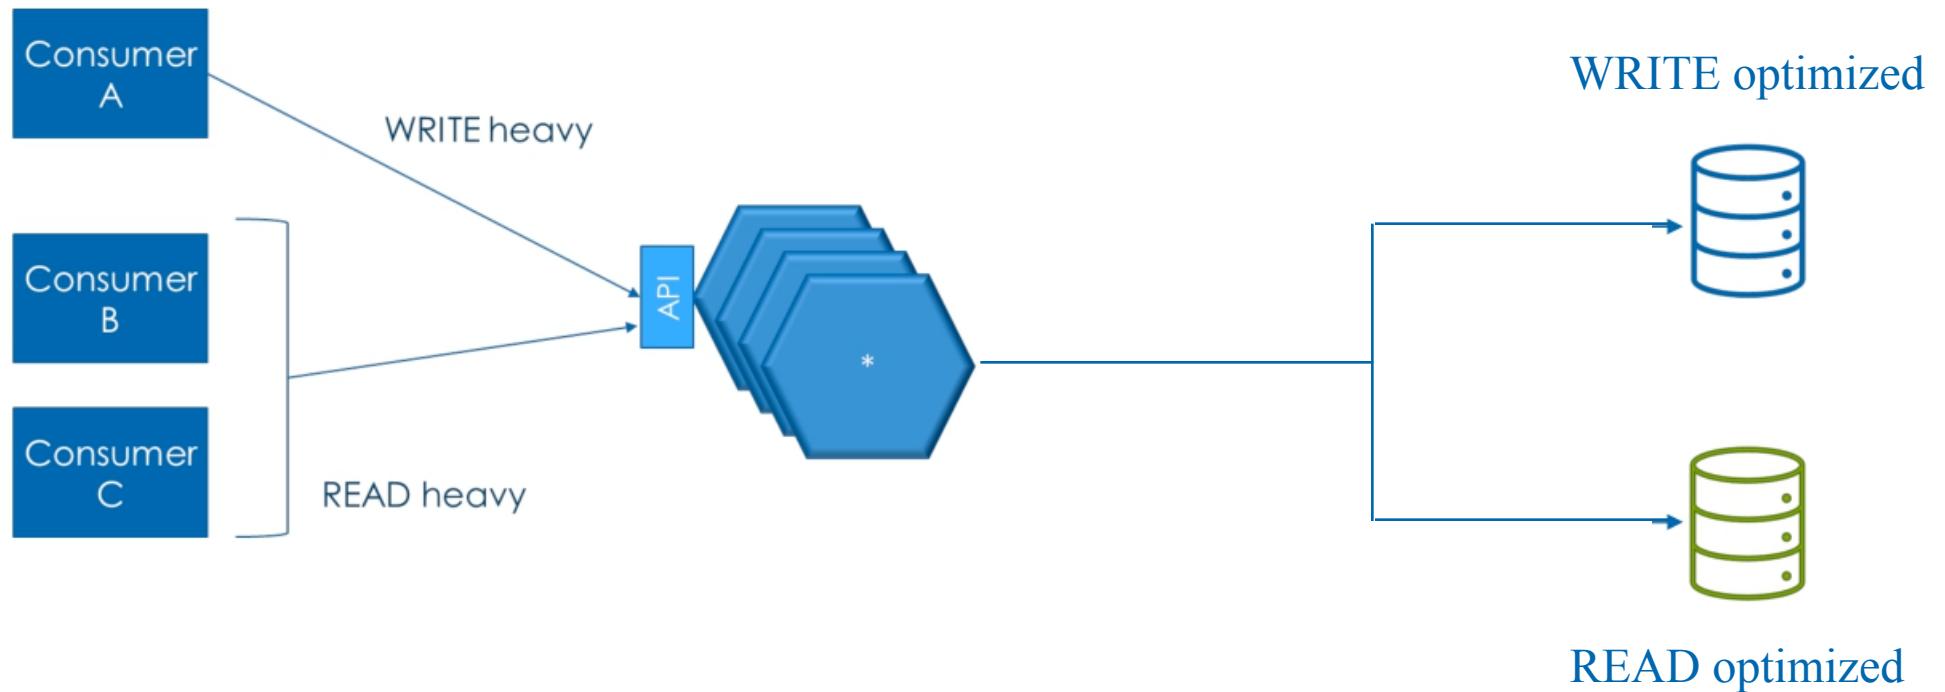
## Microservice Performance



## Database will become the Bottleneck

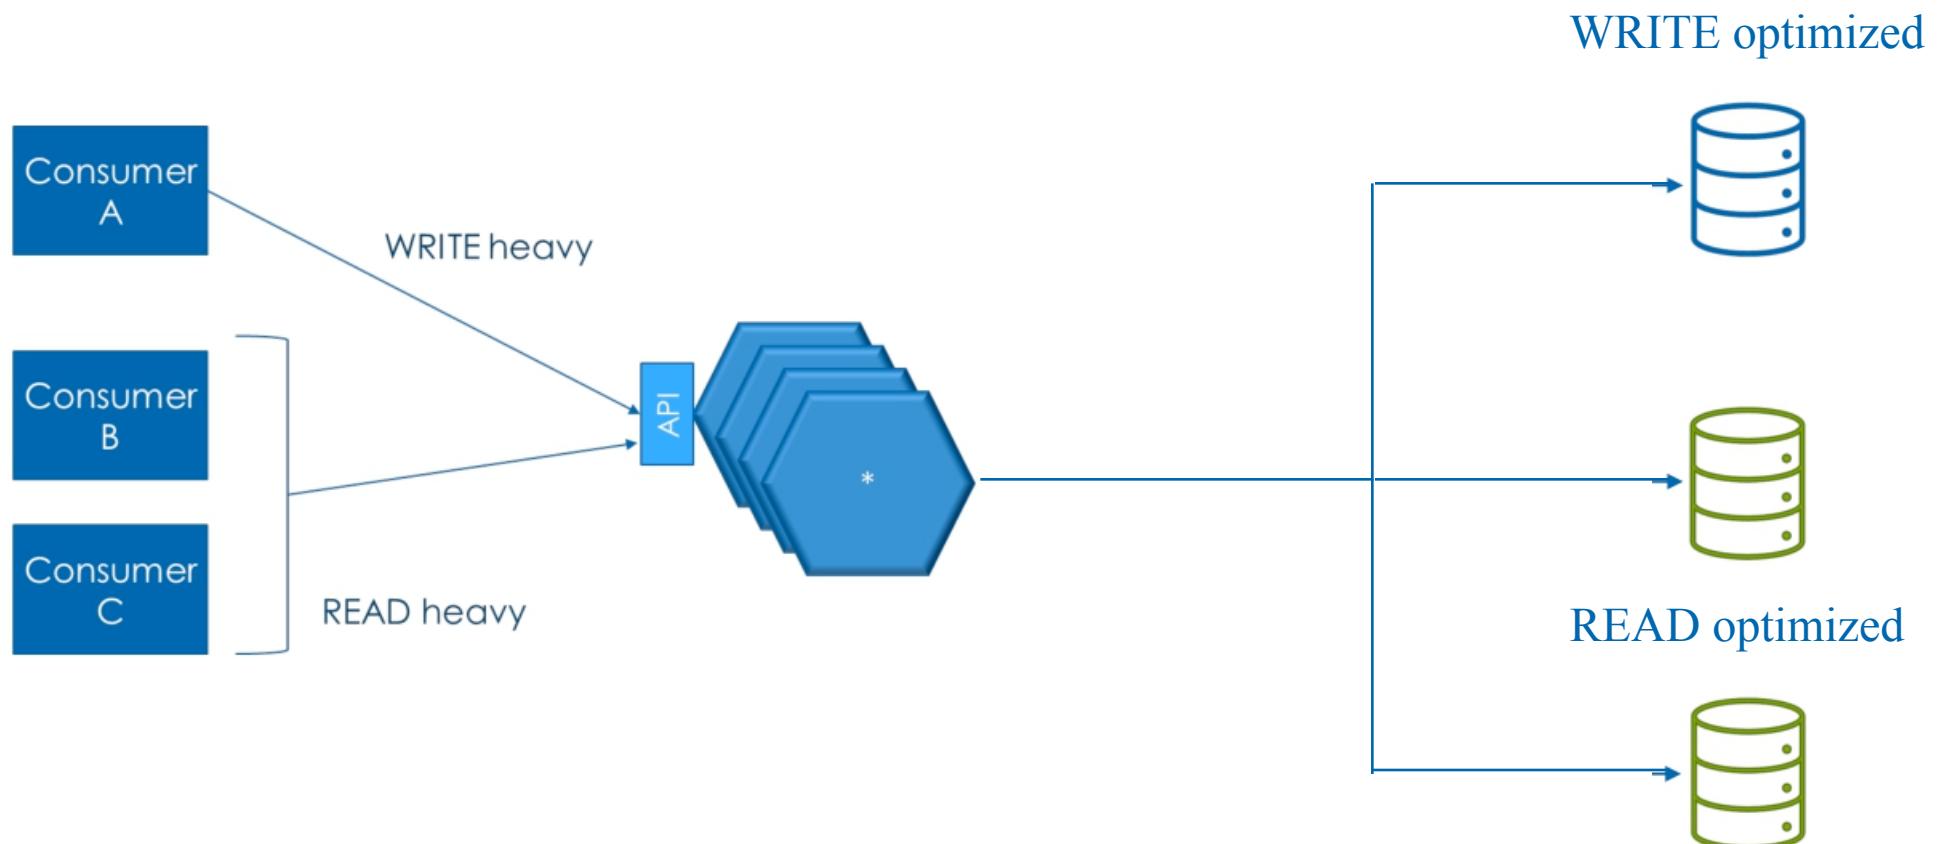


## Separate the Databases !!!

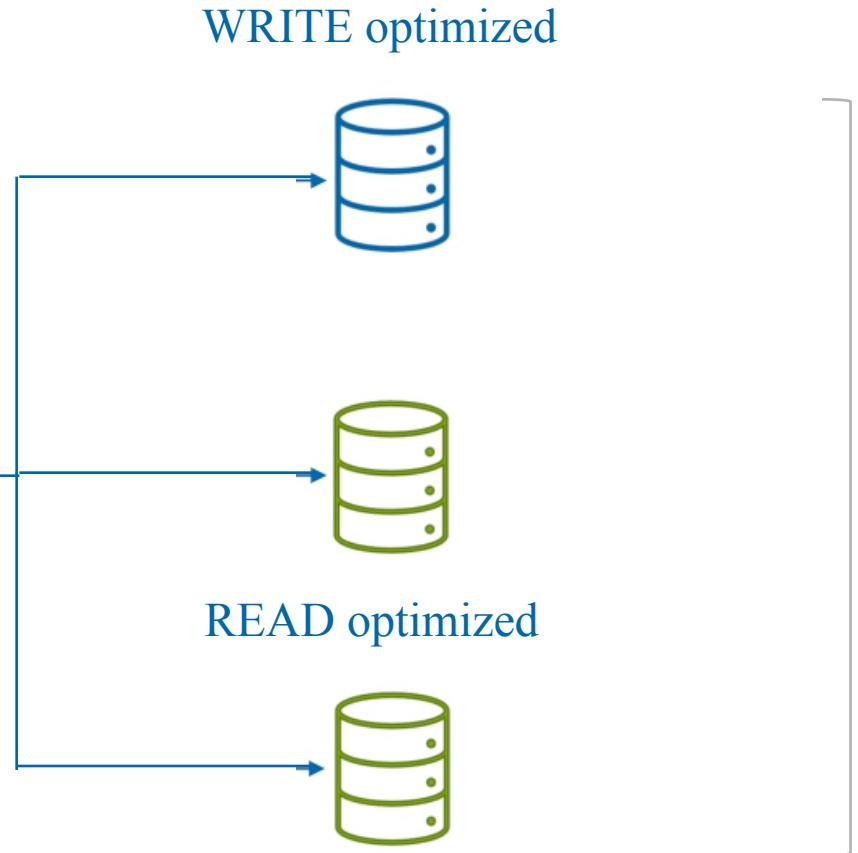


# Performance & Data Integrity

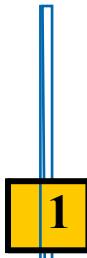
## Separate the Databases !!!



## Data Integrity MUST be maintained across Databases



Need for additional patterns 😊



1 Command Query Separation pattern (CQS)

2 Command Query Responsibility Segregation (CQRS)

3 Event Sourcing pattern

4 Case Study : Working CQRS for ACME Sales proposal capability  
(UML, JAVA, RabbitMQ, PostgreSQL, MongoDB)

# Commands & Queries



Command Query Separation Principle

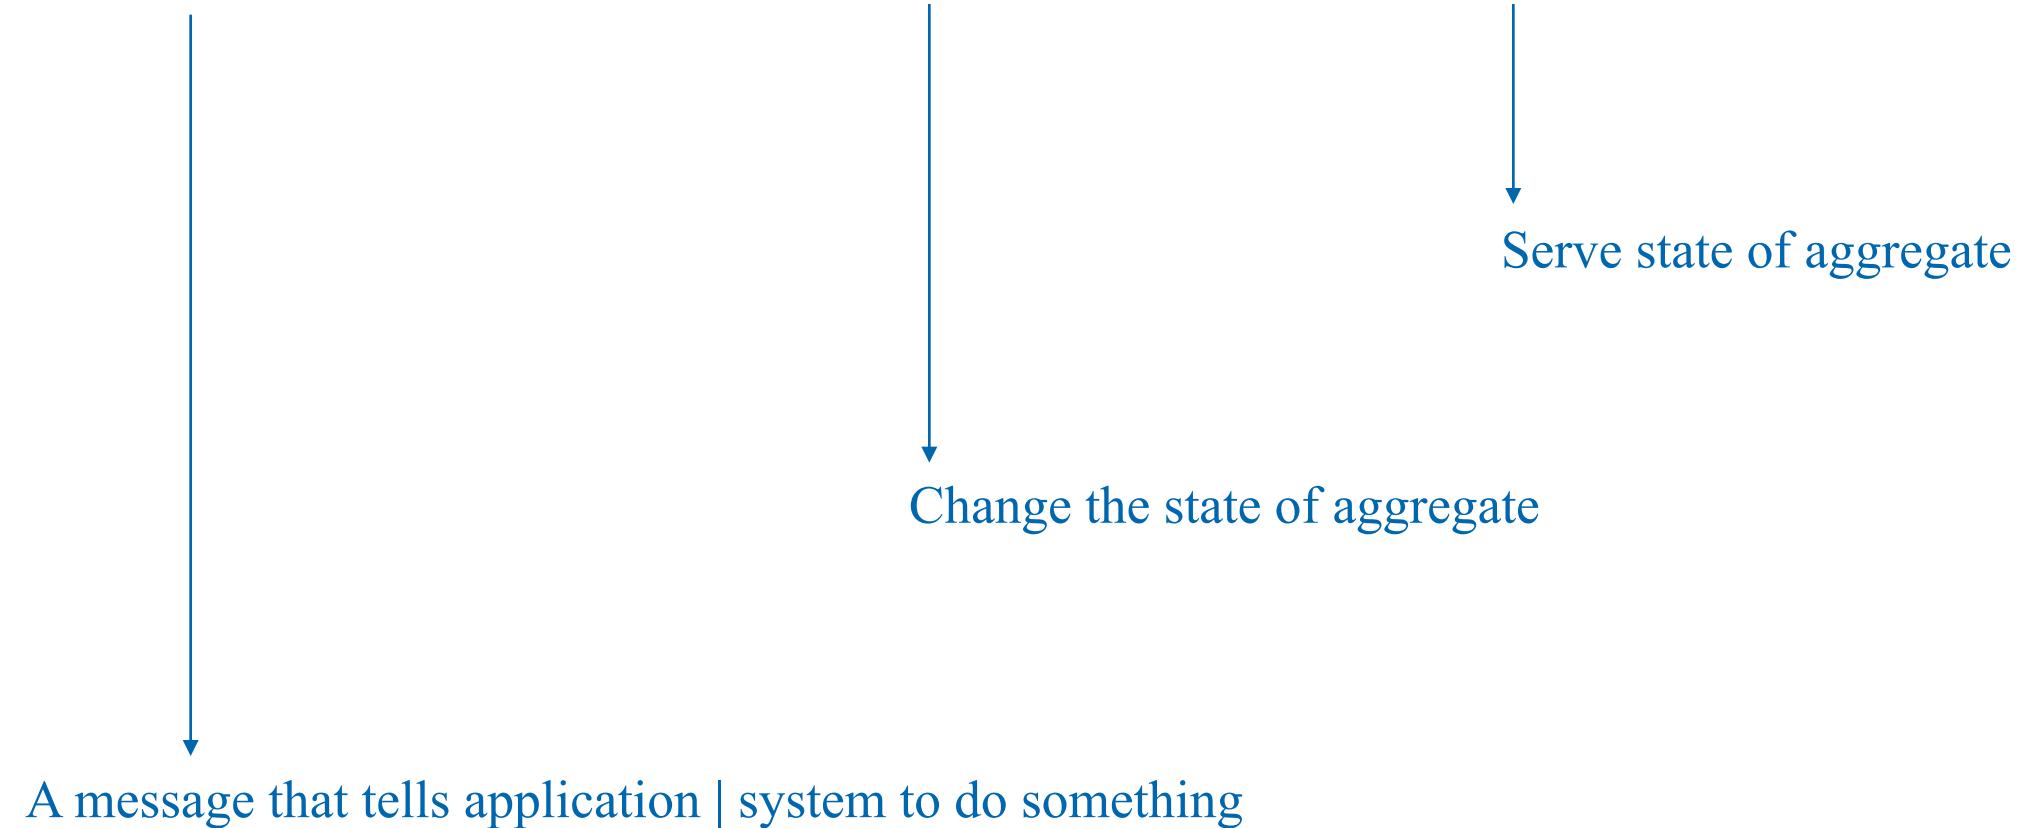
---



- 1 Commands & Queries
- 2 Command Query Separation (CQS) Principle
- 3 Realization of CQS

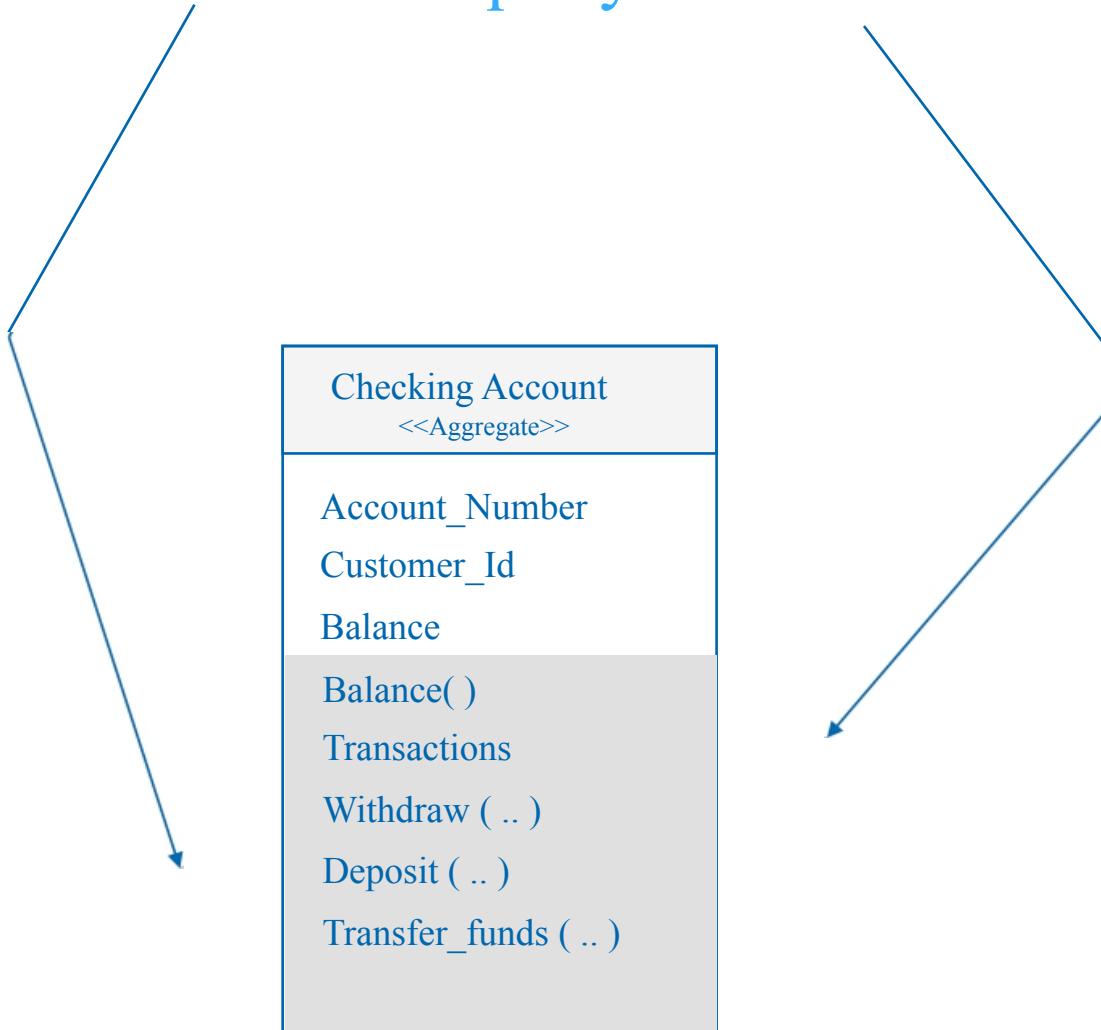
## Domain Object Operation

An Operation is either a command or a query



## Domain Object Operation

An Operation is either a command or a query



# Command Query Separation principle

a.k.a.

CQS pattern

“

Clear demarcation between methods that change the state & those that don't change the state of domain

Commands = Changes the state of object

Queries = Does NOT change the state



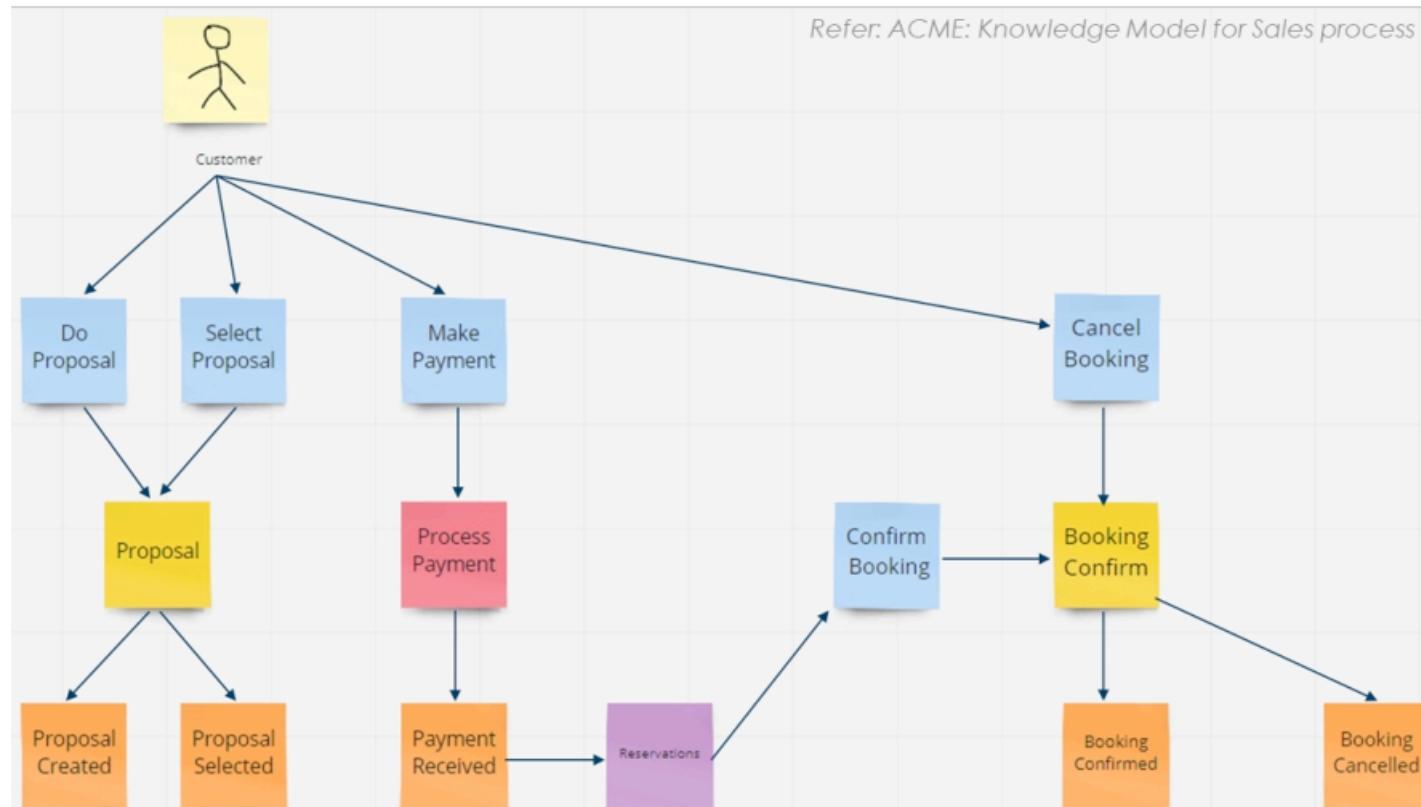
Suggested By: Bertrand Myers

[https://en.wikipedia.org/wiki/Command - query\\_separation](https://en.wikipedia.org/wiki/Command - query_separation)

# Commands

Command

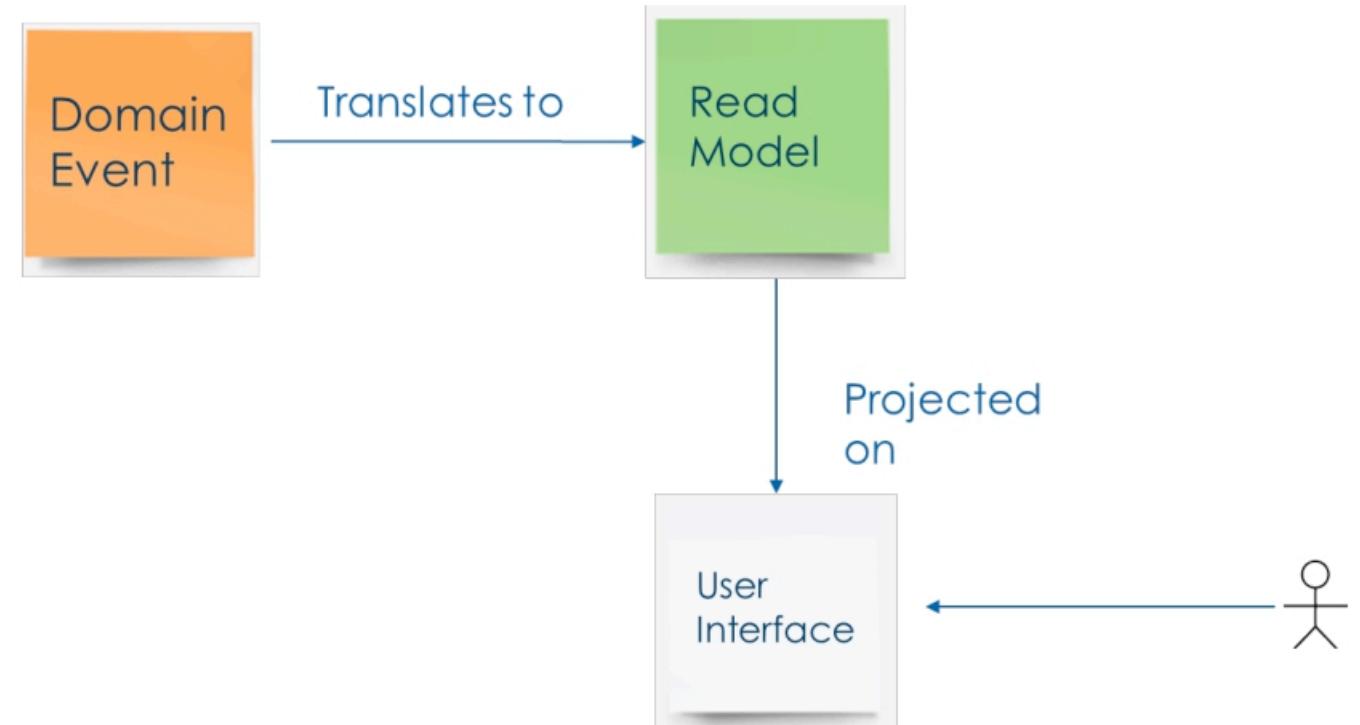
- An Action initiated by user/system/service/time
- This action MUST be carried out
- Named as verbs



# Read Models



- Also known as Query Model
- Projected on the UI



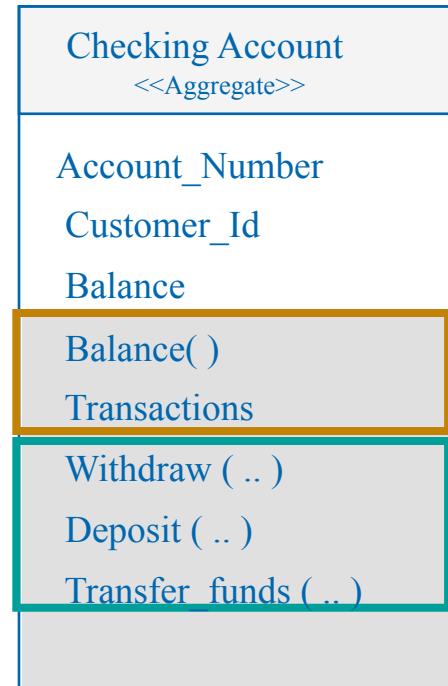
## Separation of concerns

Command & Query realized across multiple classes

Complex Logic

Unit of Work

Scalability



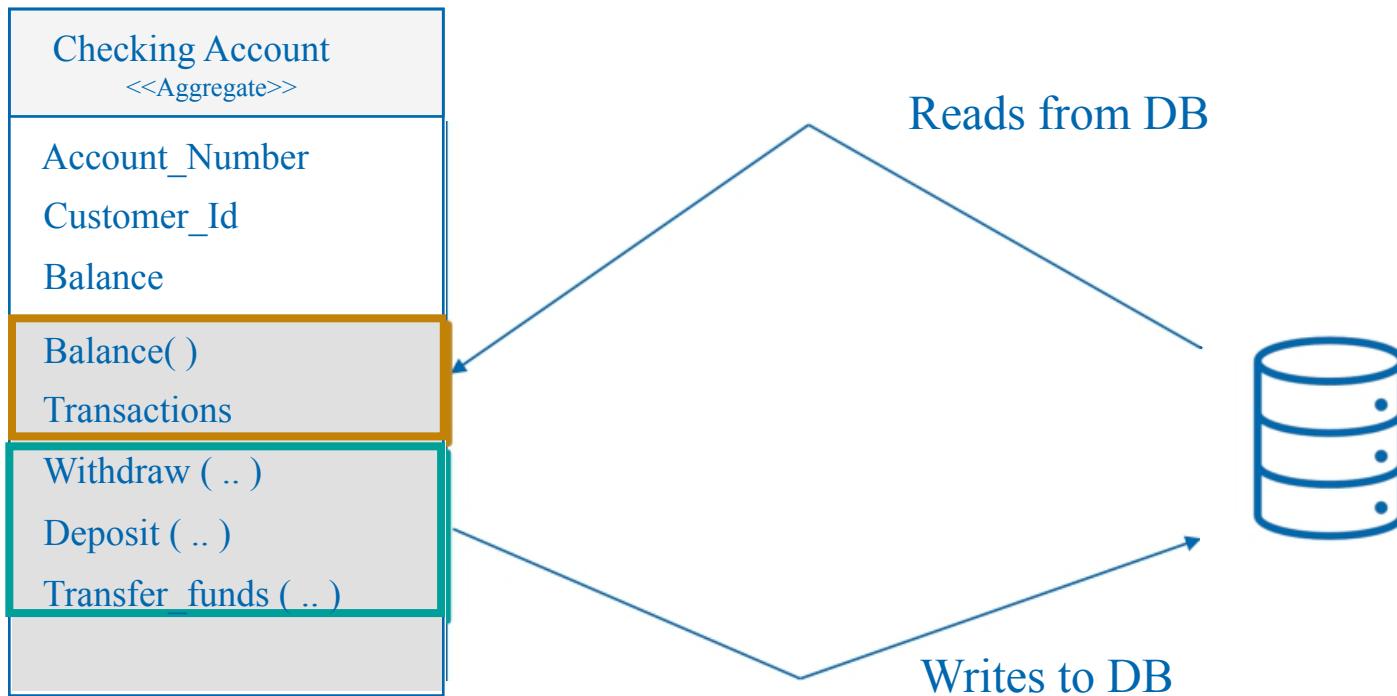
Performance

Flexibility

# Shared Datastore

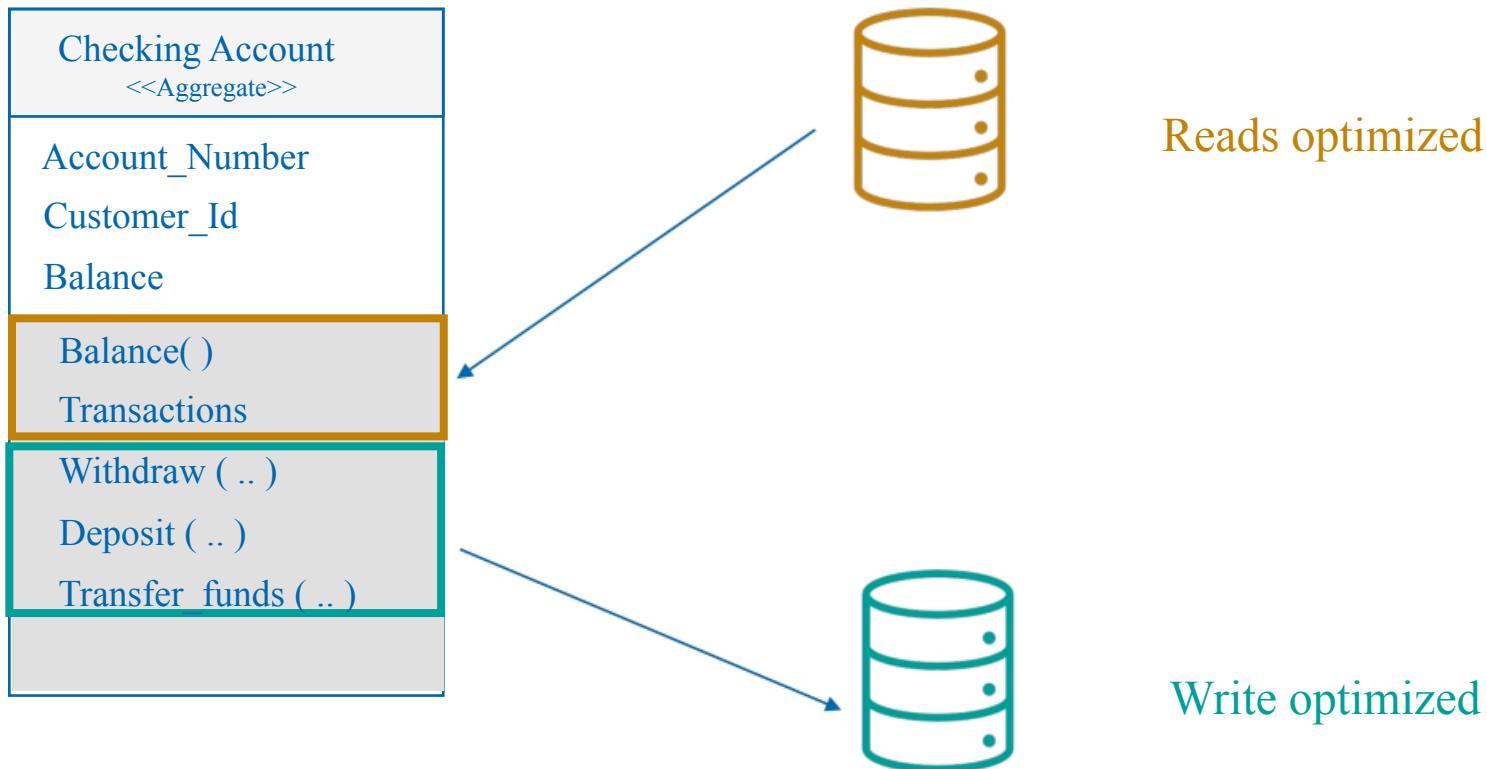
Commands - Writes

Query - Reads



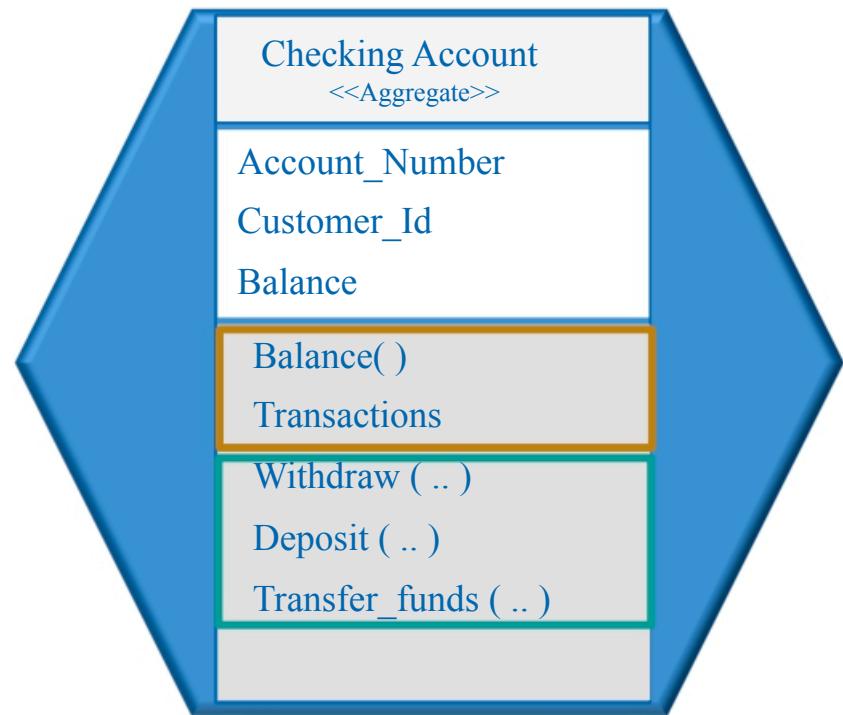
## Separate Datastores

May use Datastores optimized for Reads & Writes



## Realization of CQS

Commands & Queries in CQS are part of the same component





## Quick Review

Commands => Actions or Intent

Query => Serves the domain data

CQS => Separate operations for Command & Query

Use Shared | Separate Datastores depending on needs

# Realization : Command & Query



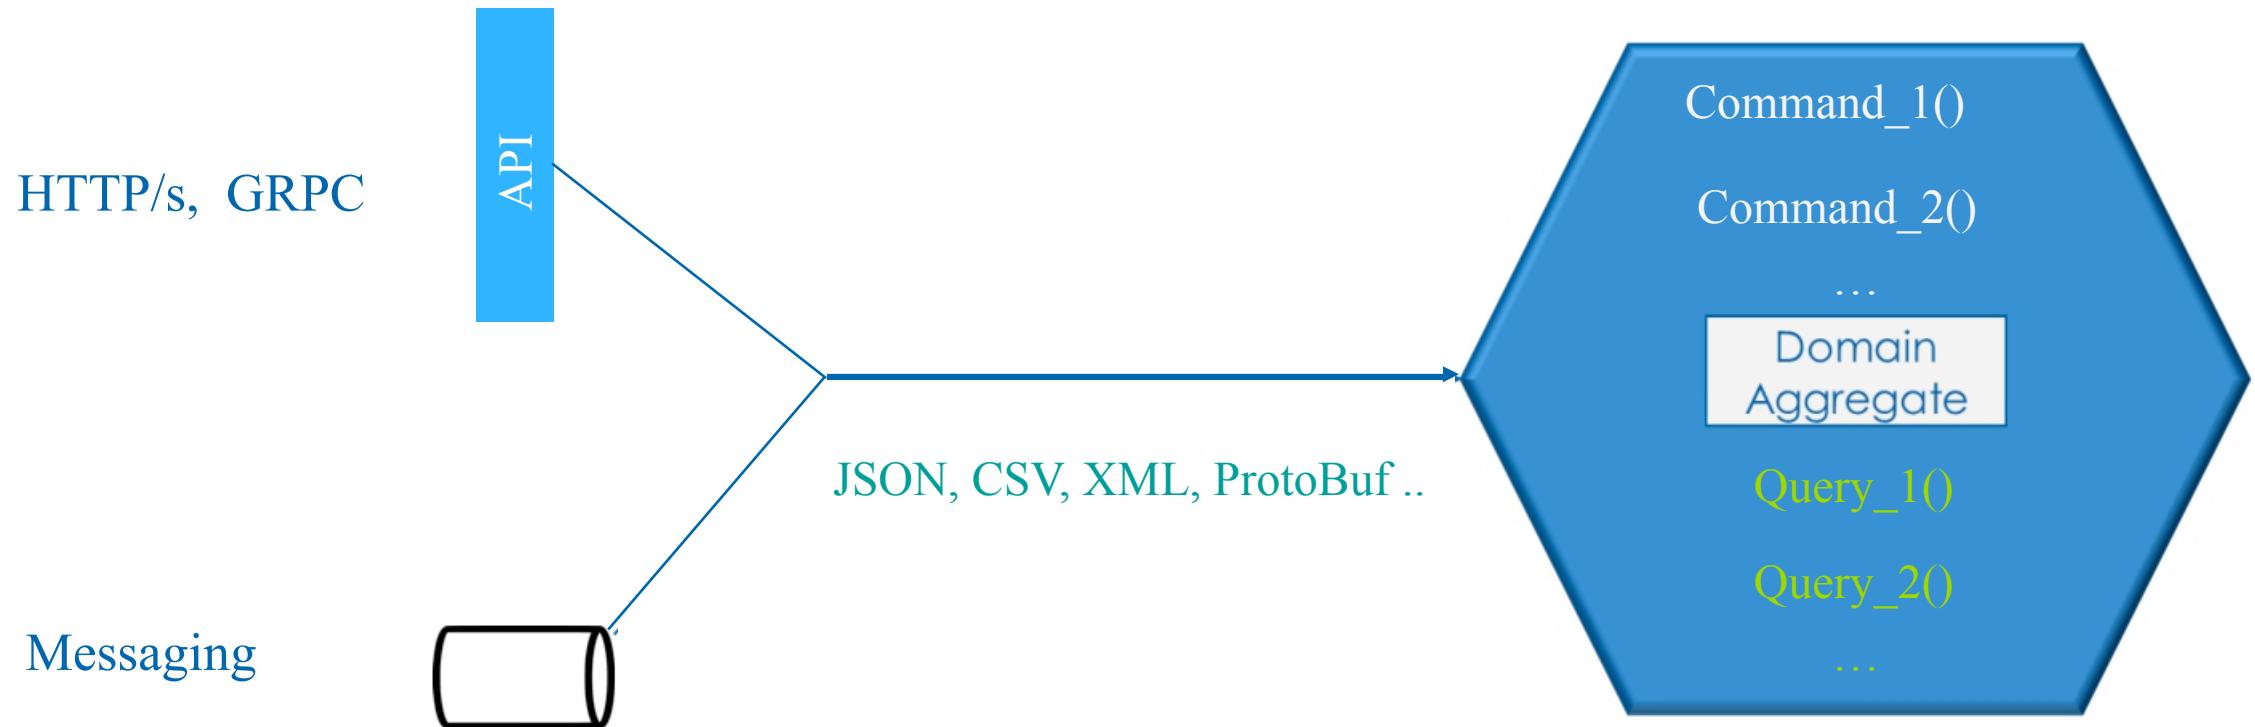
READ & WRITE performance considerations



- 1 Command | Query Realization
- 2 WRITE & READ Performance
- 3 Collaborative Domains

# Command & Query Interface

Commands & Queries exposed via Synchronous & Asynchronous



# DDD : Command & Query models

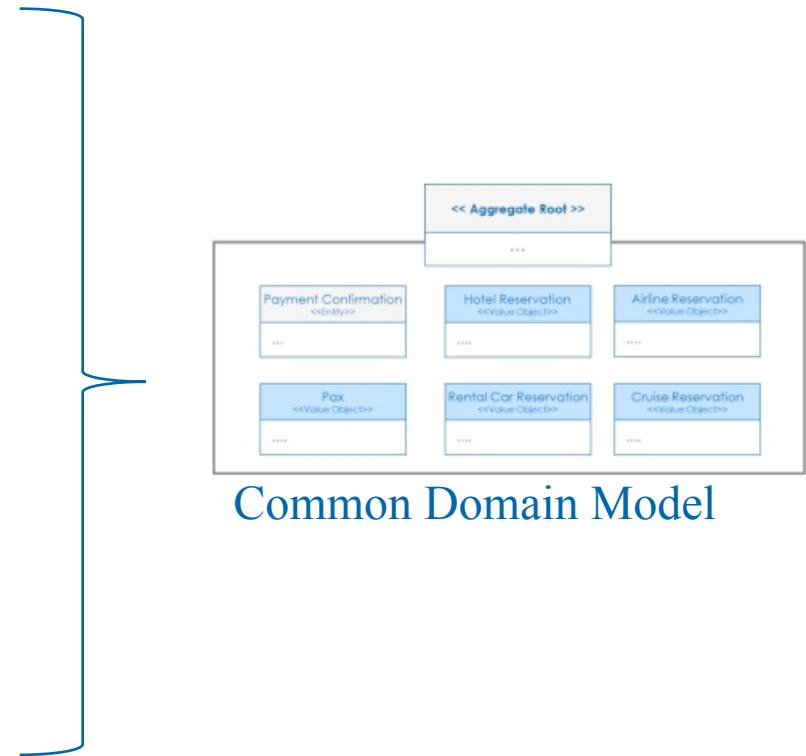
Data from/to DB translated to/from the common Model

- Commands => WRITEs to the DB

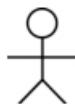
Referred: WRITE Model  
WRITE Side

- Queries => READs from the DB

Referred: READ Model  
READ Side



Common Domain Model



Command

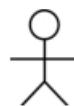
</>

1. Convert the Request to domain model
2. Execute business Logic
3. Convert to DB API/SQL Write the data

Common Model for  
READ | WRITE



Common Database for  
READ | WRITE



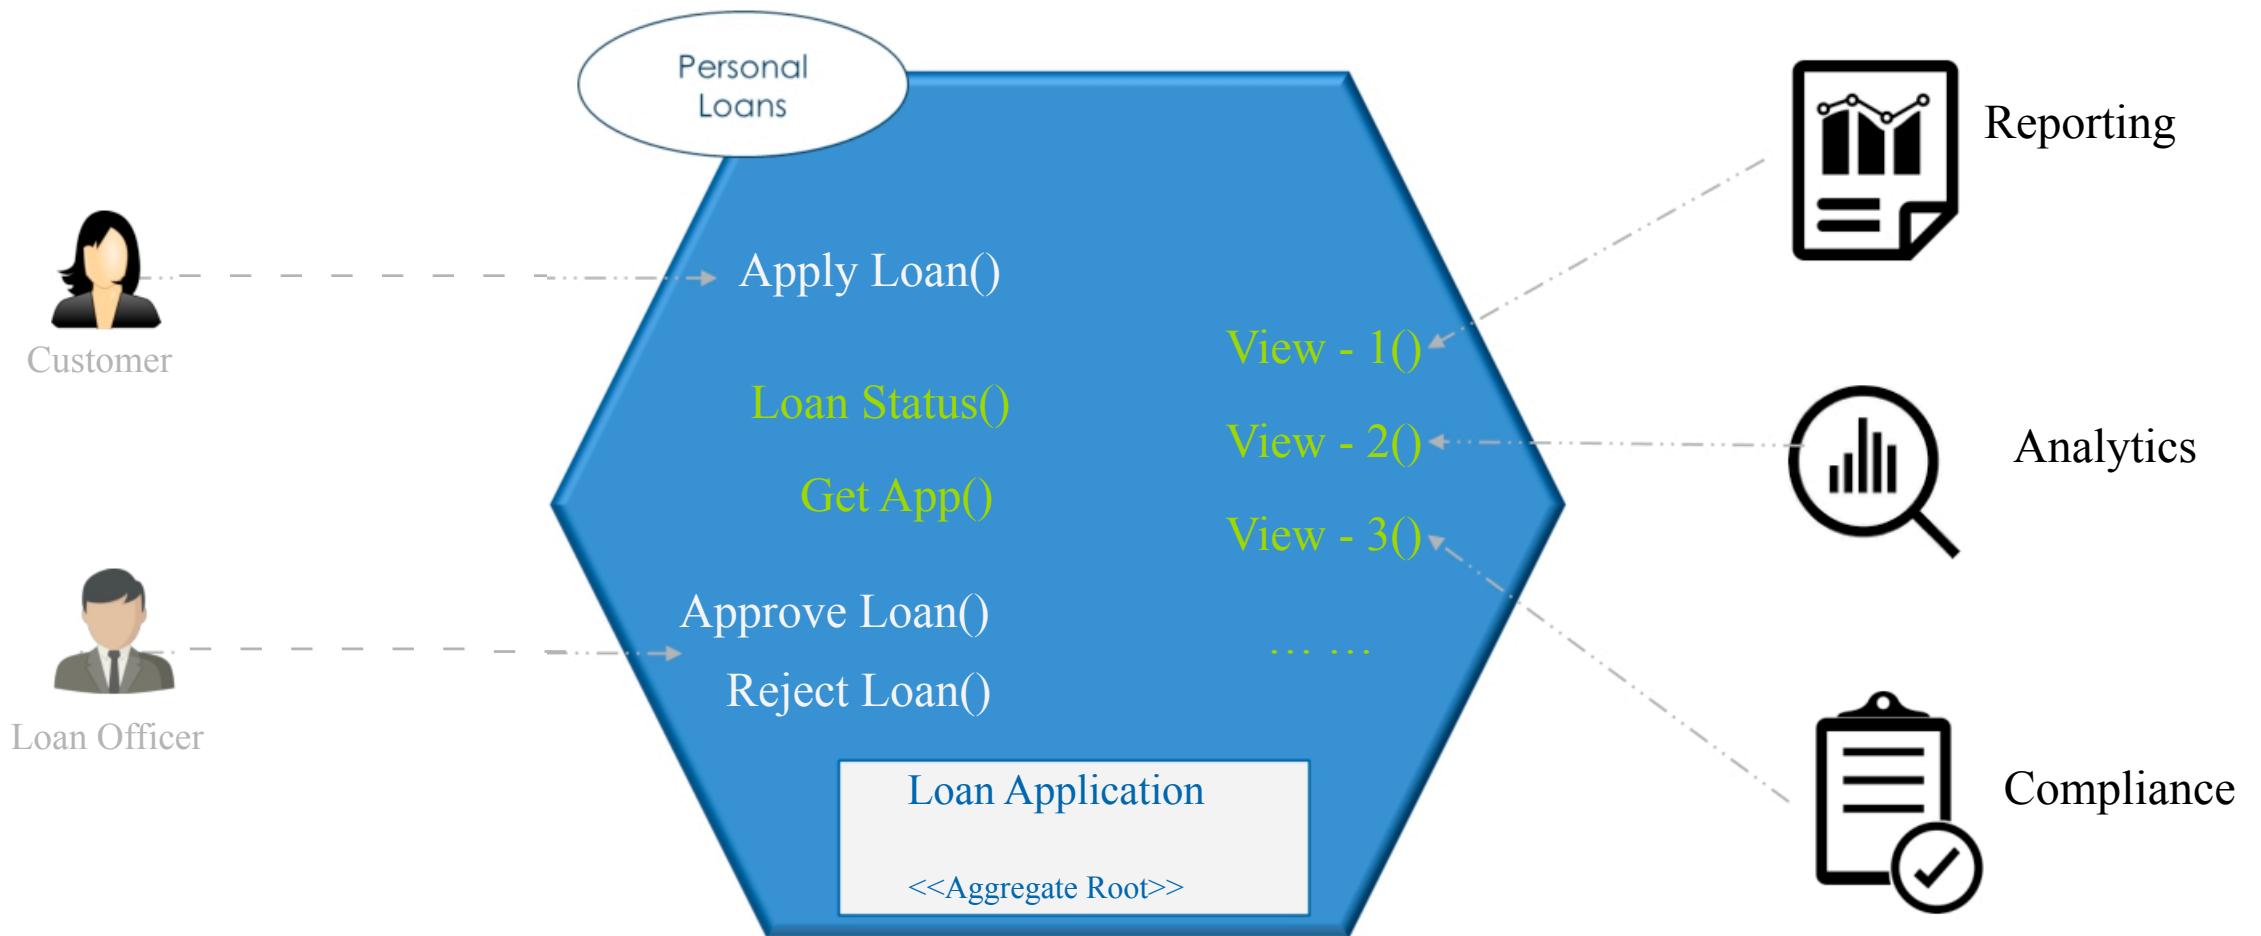
Query

</>

1. Convert Request to domain model
2. Get the Data from the database
3. Convert the Response to domain Model

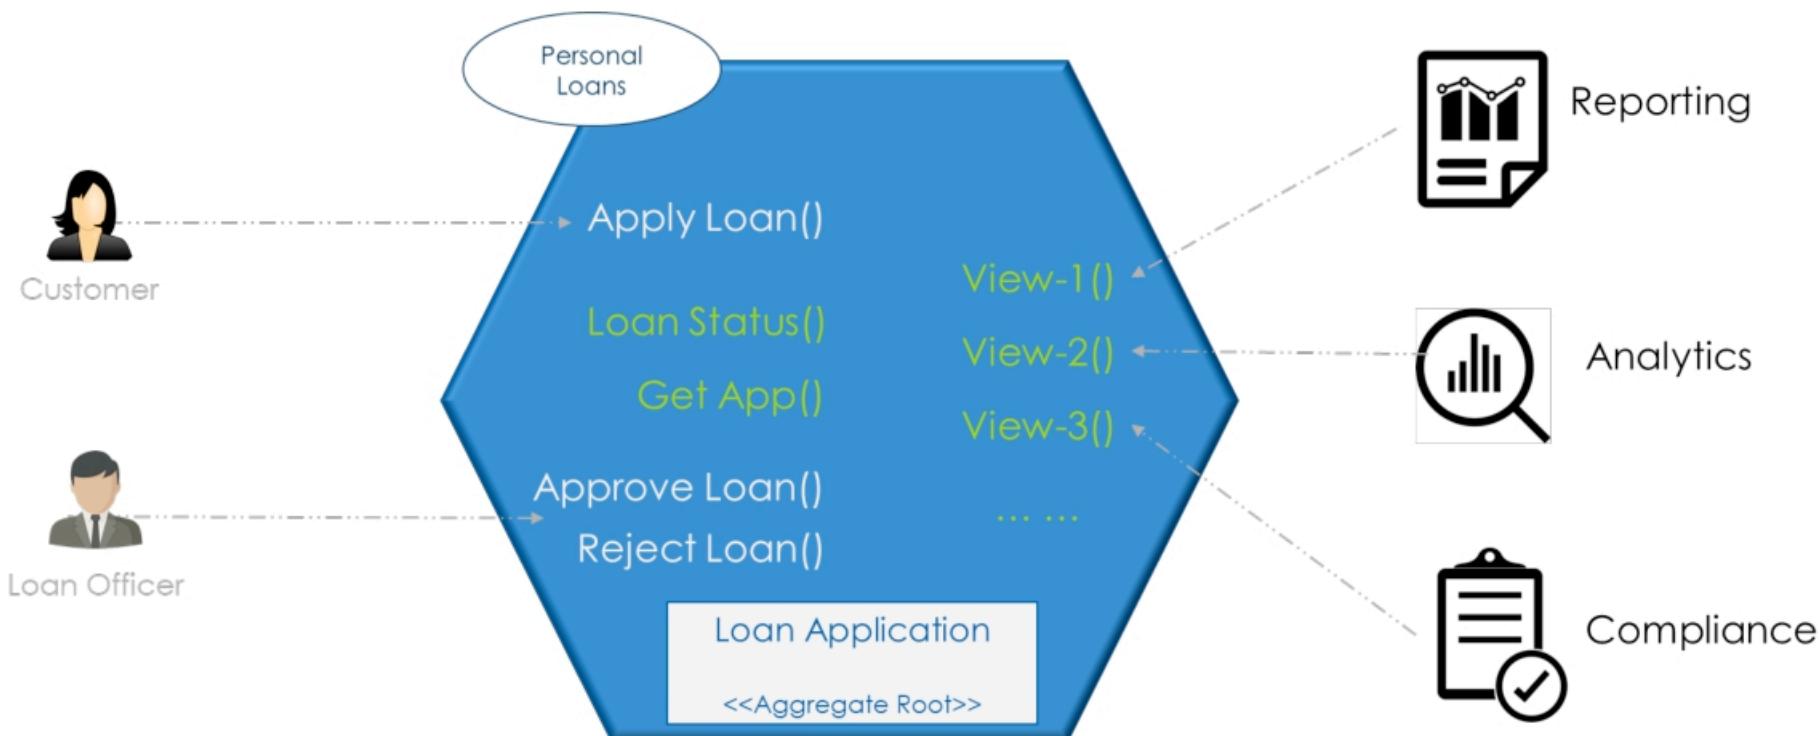
## READ model flexibility

Query consumers requests different view/model of the same data



## READ model flexibility

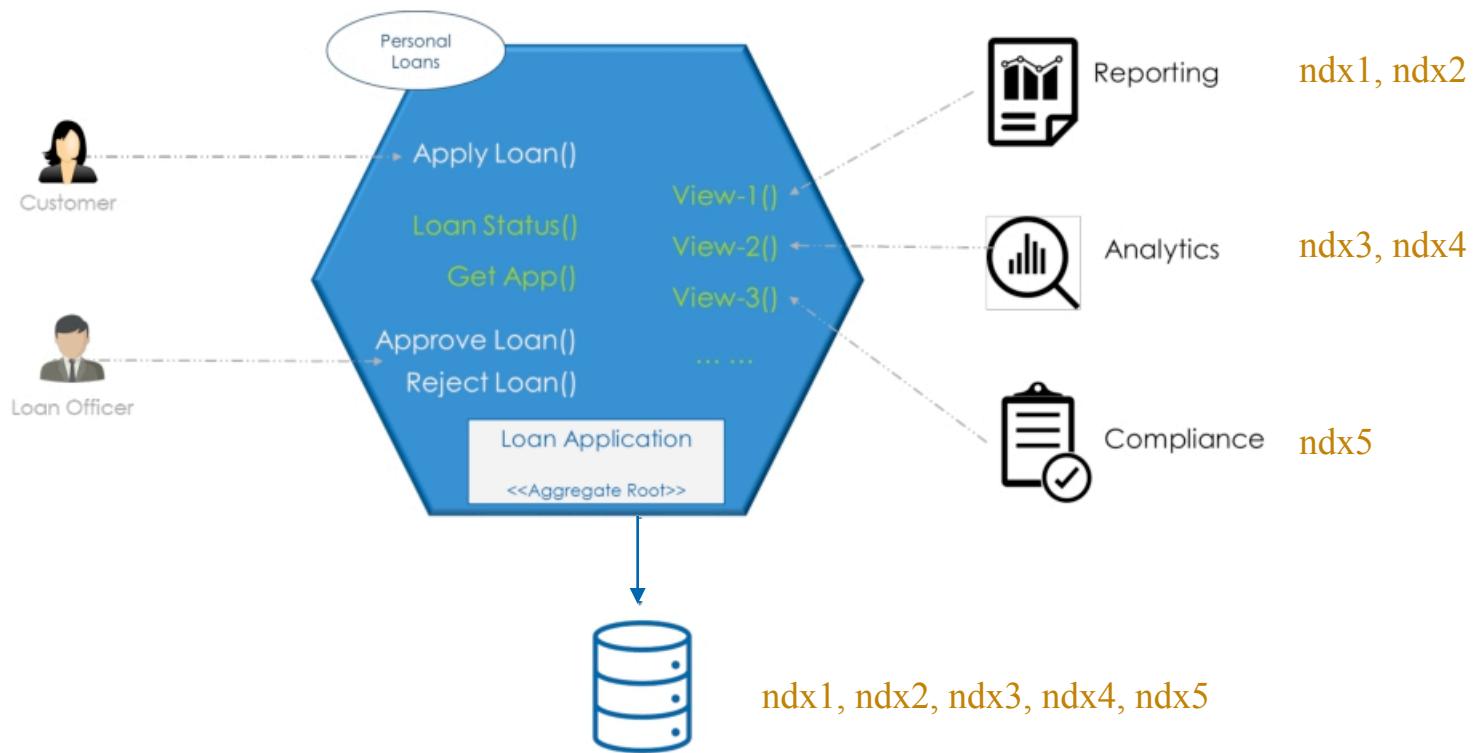
Query consumers requests different view/model of the same data



Requires changes to Microservices code

## READ performance

Multiple DB Indexes needed to achieve desired performance



Indexes negatively impacts the performance of WRITES !!!

## Collaborative domain

Multiple actors invoke commands in parallel on the same data

- Online bidding
- Concert tickets reservations
- Trading Systems

## Characteristics of Collaborative domains

Low tolerance for performance hit | degradation

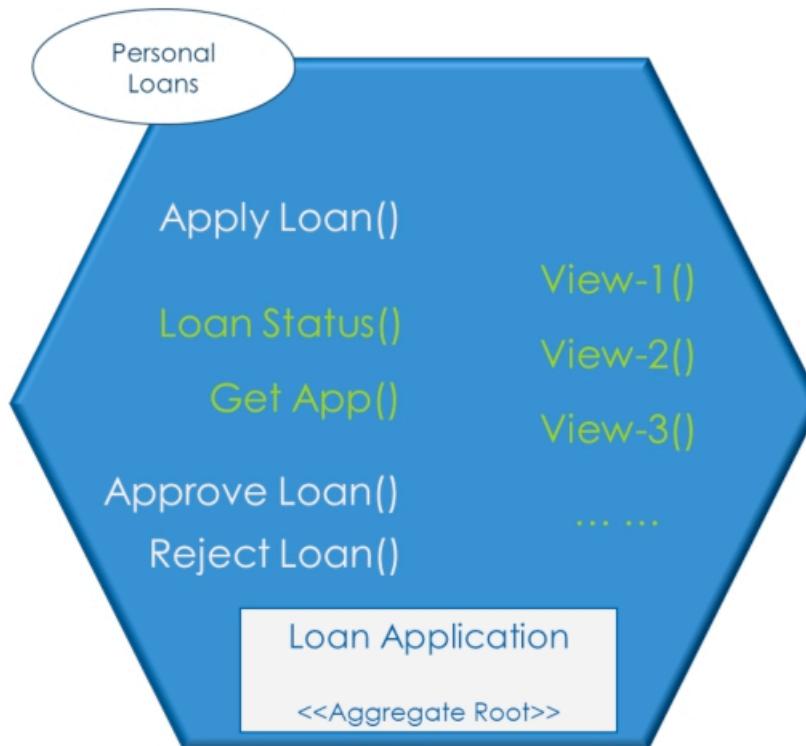
- Commands in such systems have complex business logic
- Commands are implemented with finer levels of granularity
- Business rules | logic may change frequently

WRITE & READ models need to meet these demands !!

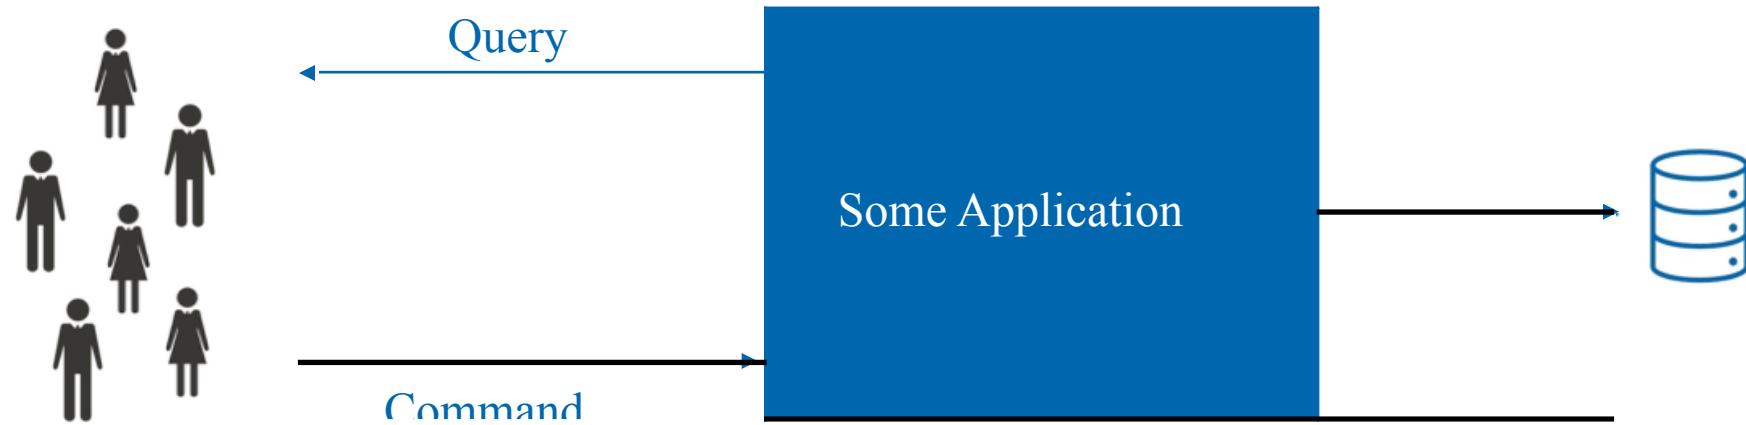
## READ vs WRITE

Typical applications | systems have more READ (s) than WRITE (s)

~15% Writes

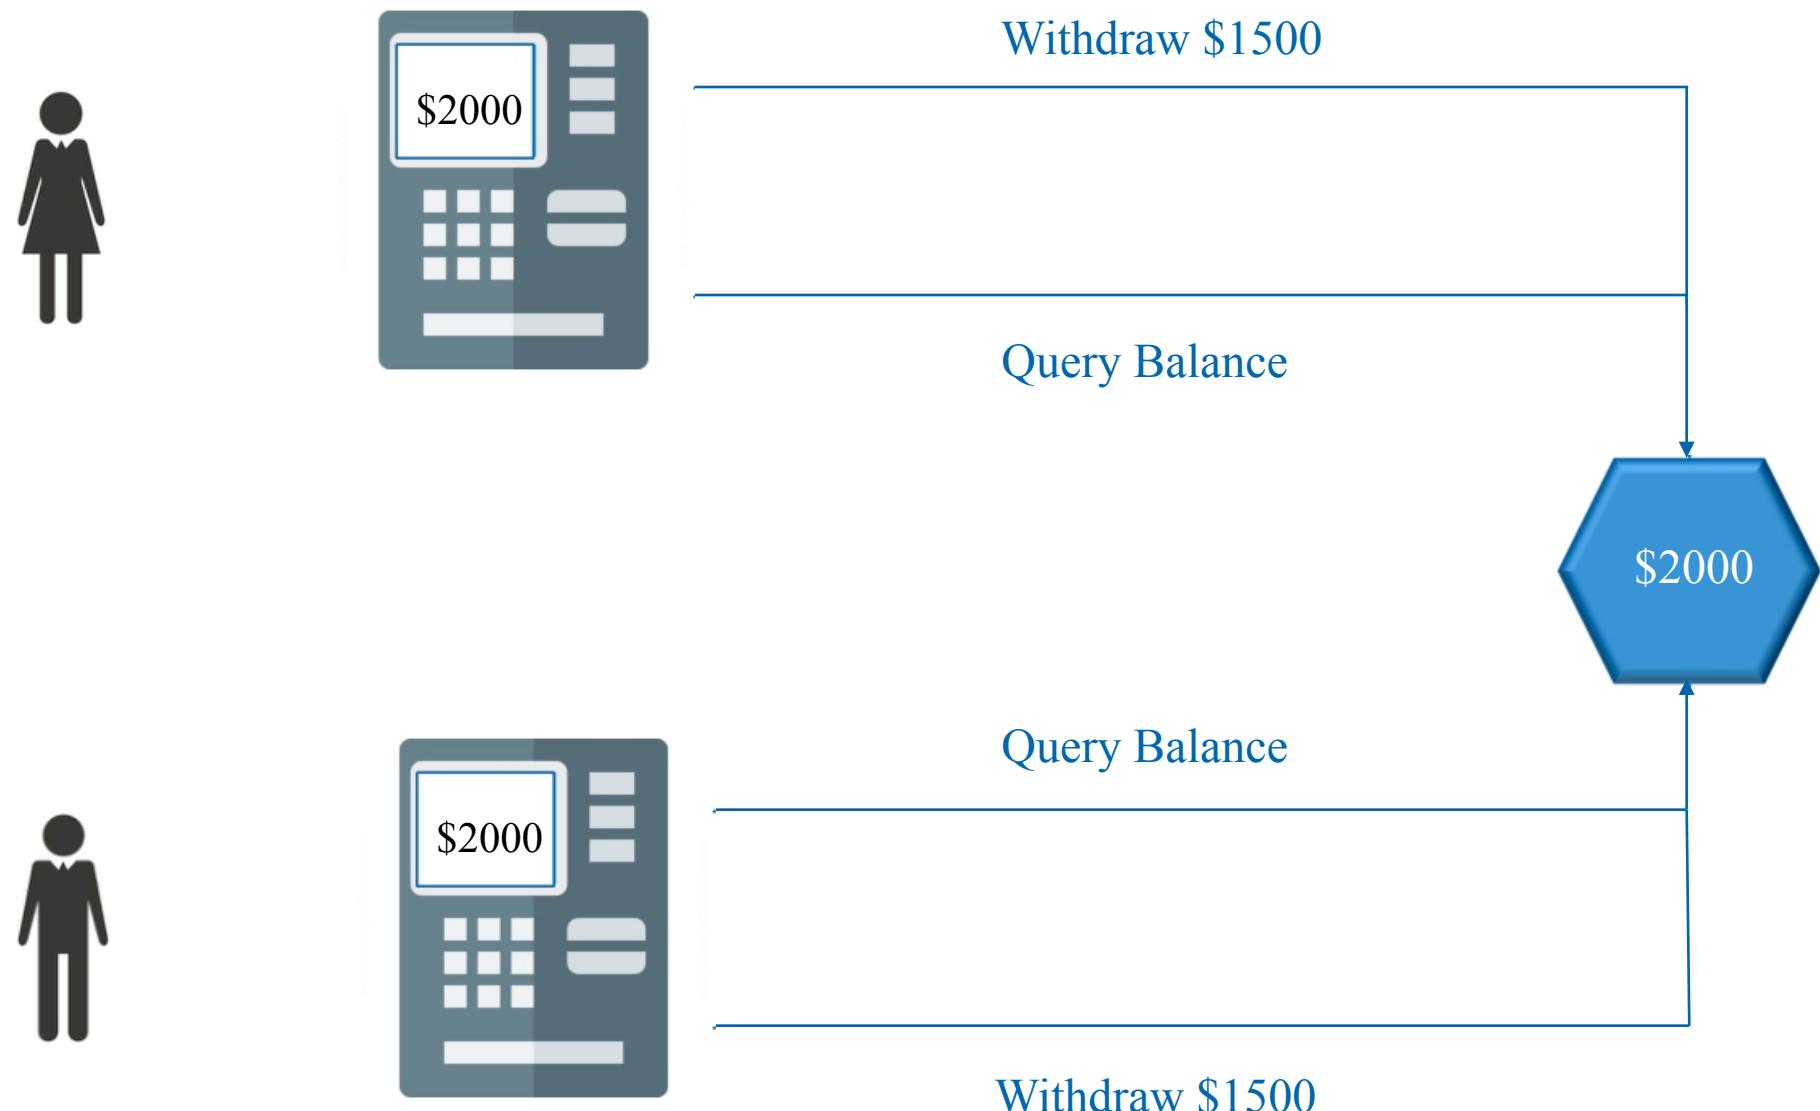


~85% Reads



Can your app deal with Stale query data?

# MOST apps can !!!



Query data is ALREADY stale when it left the microservice service !!!



## Quick Review

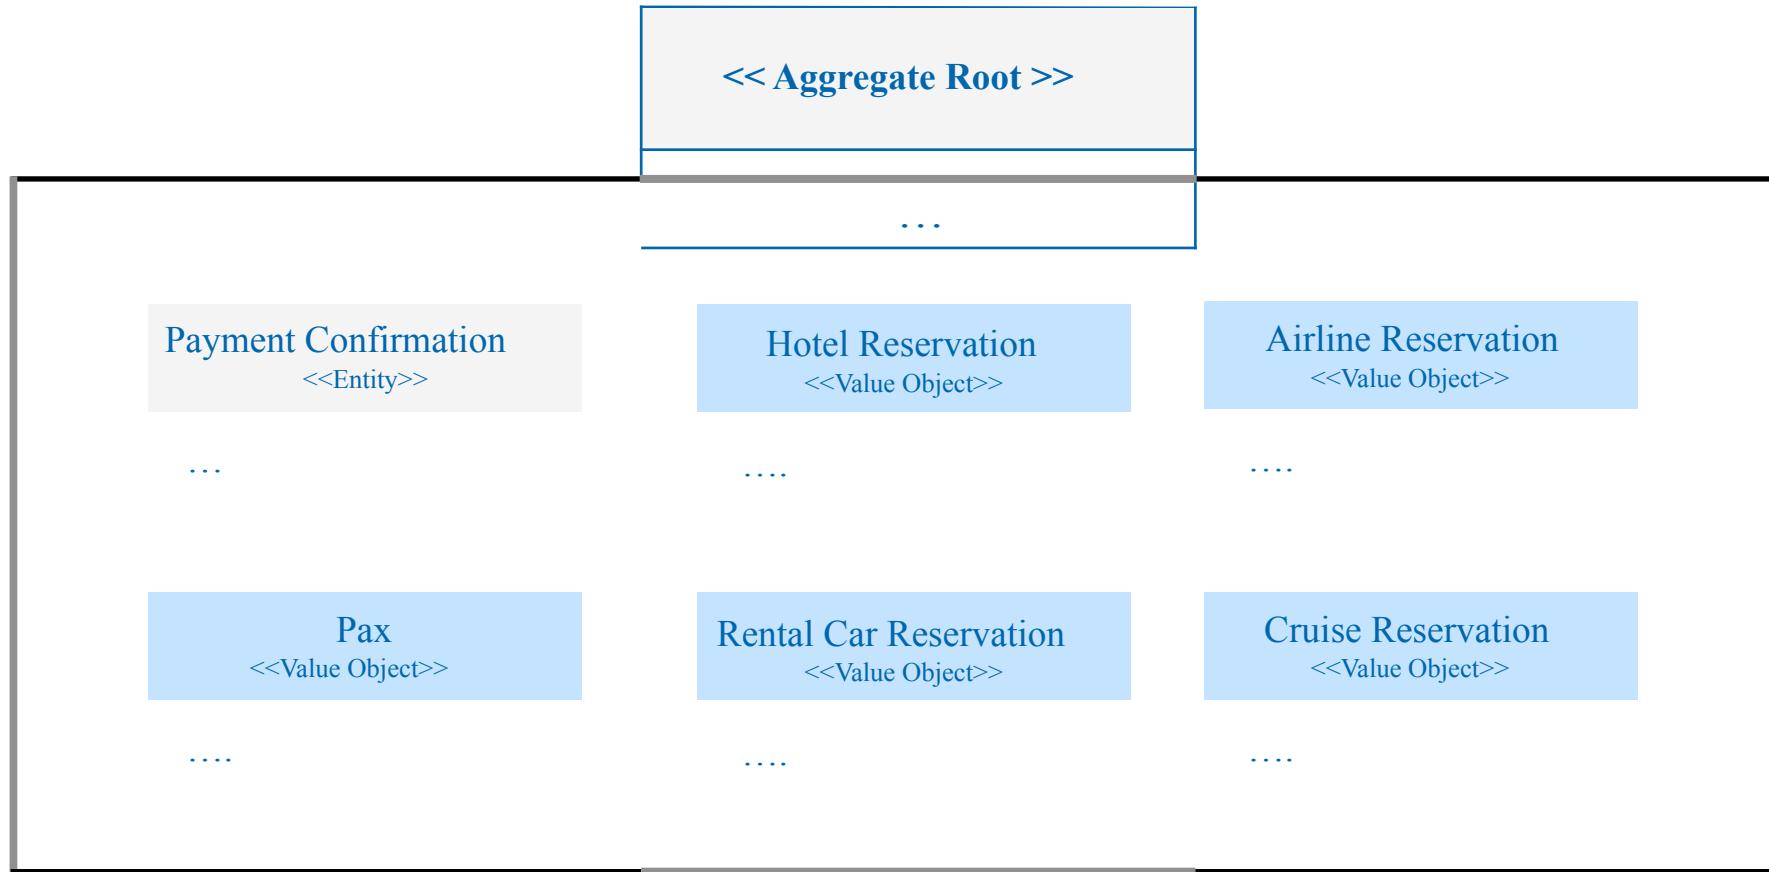
In DDD common model is used for WRITE/Command & READ/Query

Common to see MULTIPLE Query requirements

- Requires changes to the common model
- Use of indexes leads to negative impact on WRITE performance

Alternate design options needed for Collaborative domains

# Changes made only via exposed aggregate functions



# Repository for JDBC source

Managing the Data in PostgreSQL



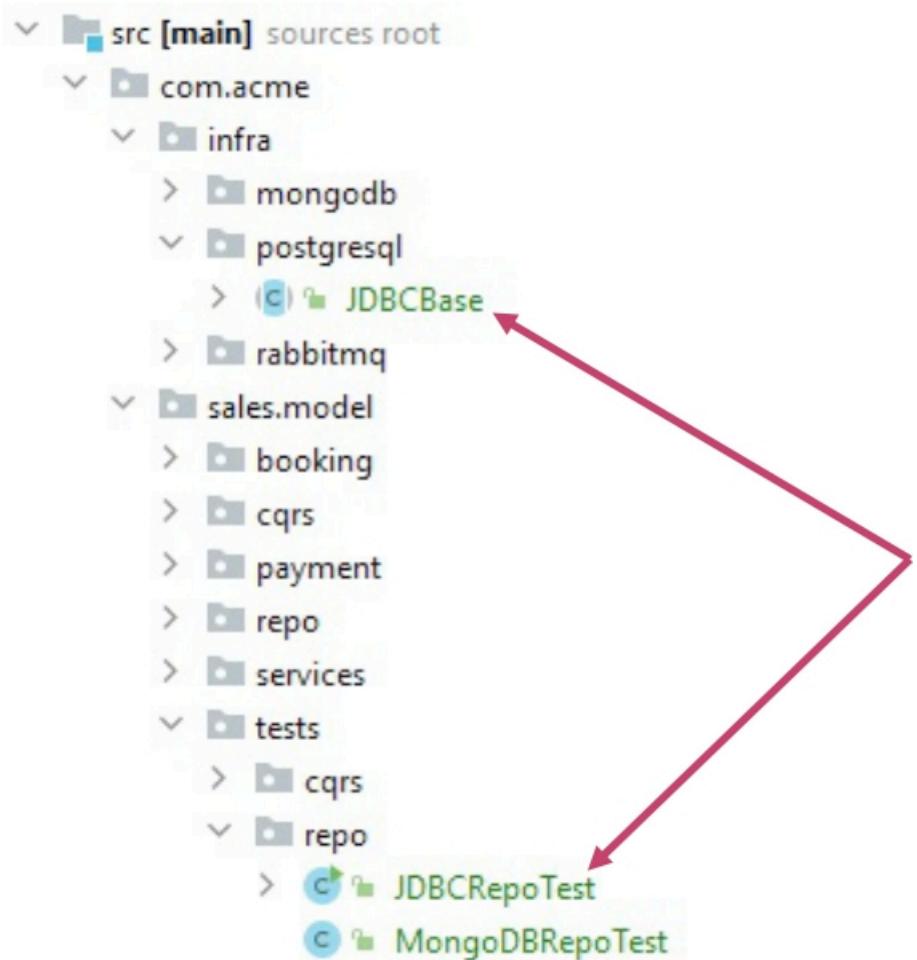
</>

- 1 Create an instance of PostgreSQL
- 2 JDBCBase , JDBCRepoTest
- 3 Testing

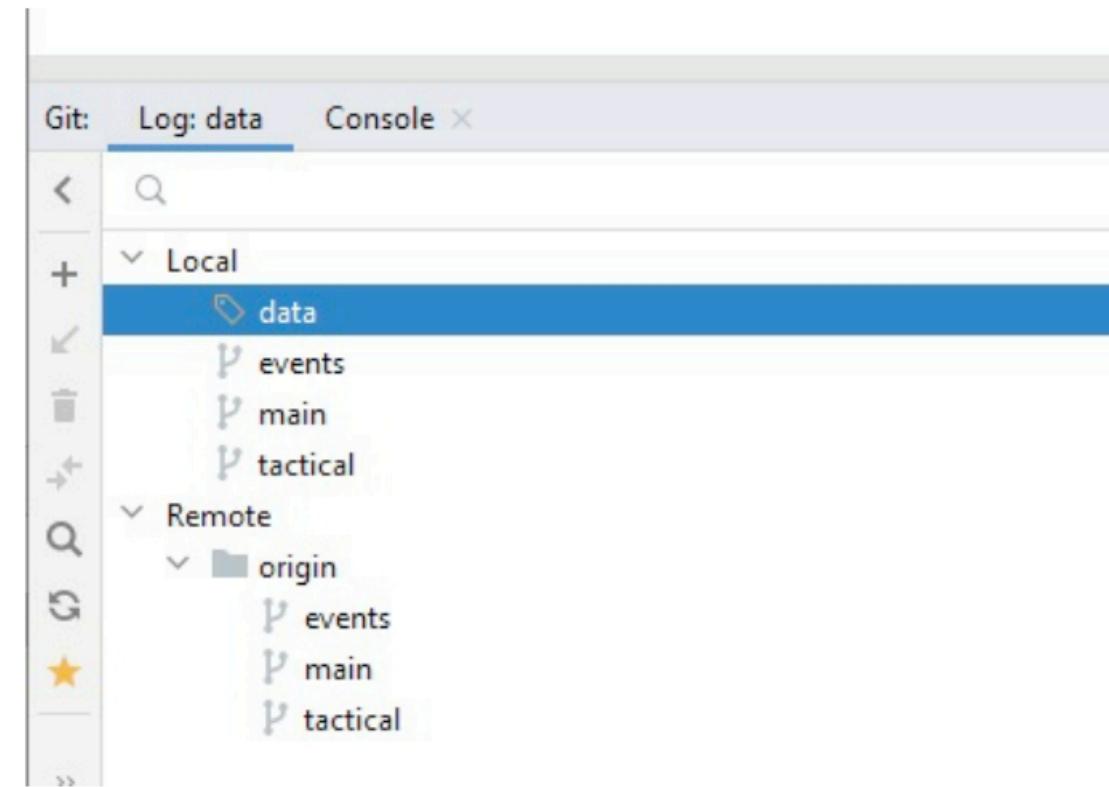
PS: Intent is NOT to teach JDBC

# Code

Checkout the branch : data



- Pre - Requisites: JDBC & JSON
- Instructions in README.md

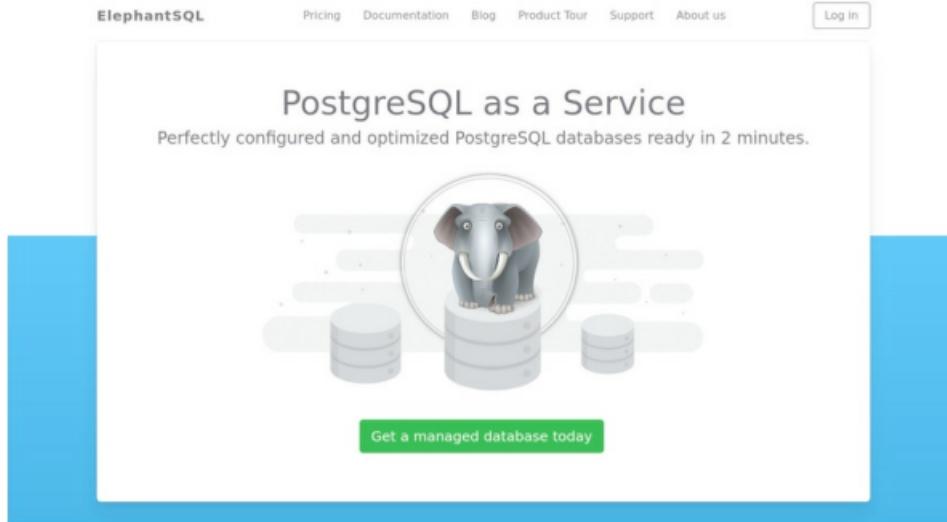


# PostgreSQL Database

## Database for ACME Sales management



Online FREE instance of the Database

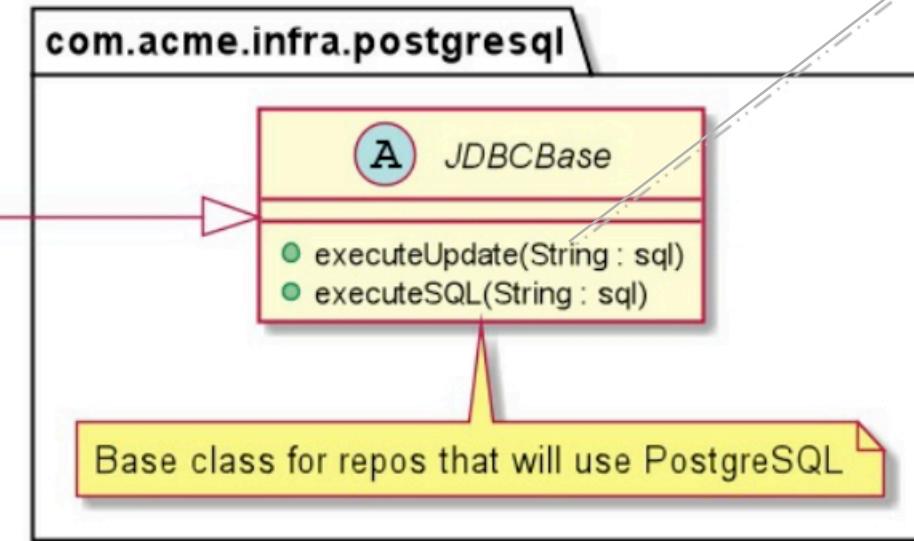
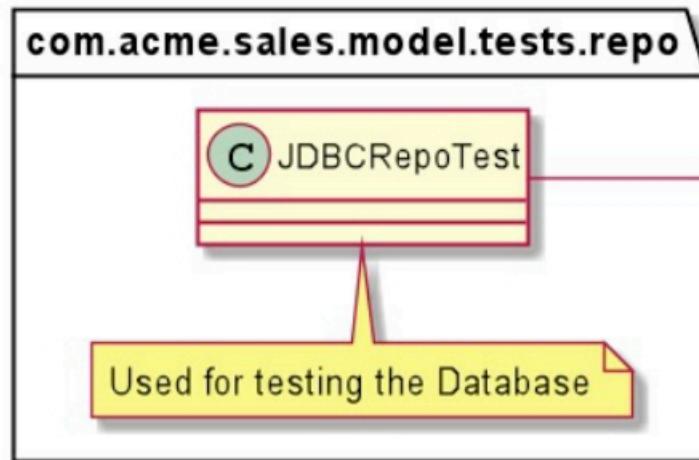


- Install PostgreSQL locally | cloud

# Testing from Java

uml / db/ jdbcbase.class.puml

Execute SQL functions return result in JSON format



## Details

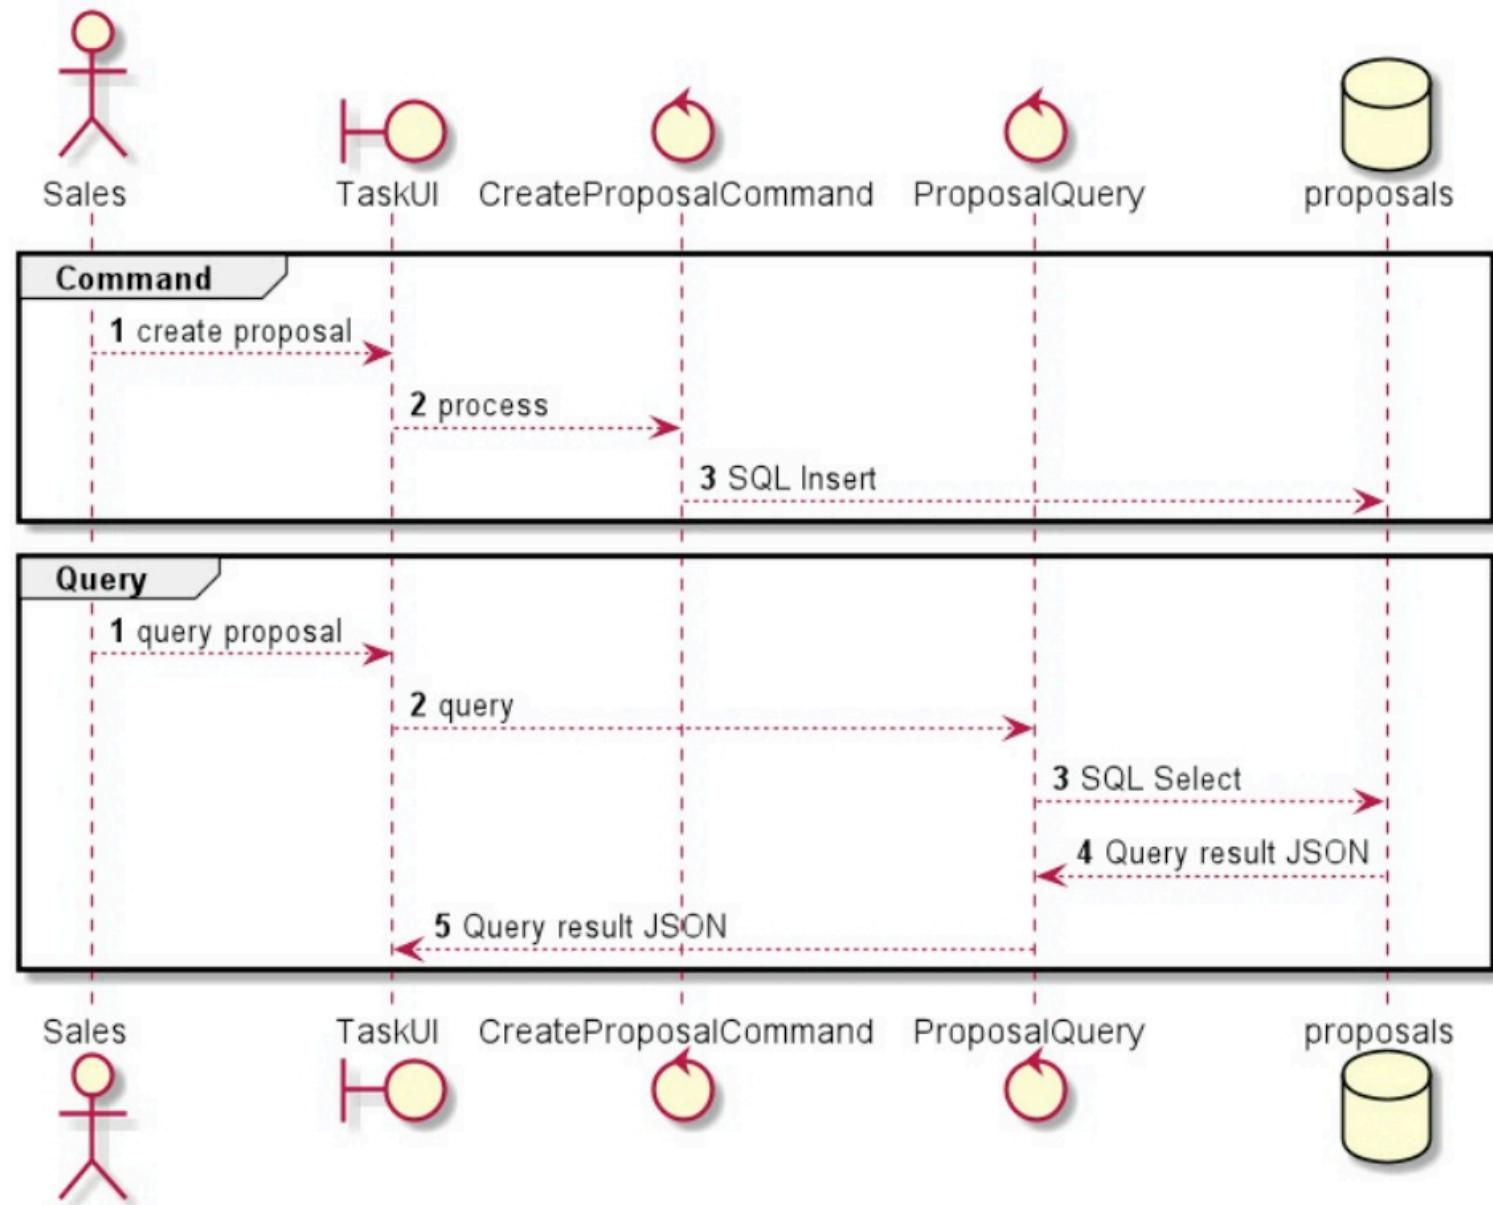
Server	ziggy.db.elephantsql.com (ziggy-01)
Region	amazon-web-services::us-east-1
Created at	2021-01-18 02:31 UTC+00:00
User & Default database	wzymokfo
Reset	
Password	AqkmHd...
URL	postgres://wzymokfo:AqkmHd...@ziggy.db.elephantsql.co m:5432/wzymokfo

MUST setup the PostgreSQL connection parameters and the credentials in JDBCBase class

OTHERWISE you would get Errors/ SQLException

# Command - Query Sequence

uml / cqrs /commandquery.v1.sequence.puml



# Build & Test Command and Query



Command will write to DB; Query will read from DB

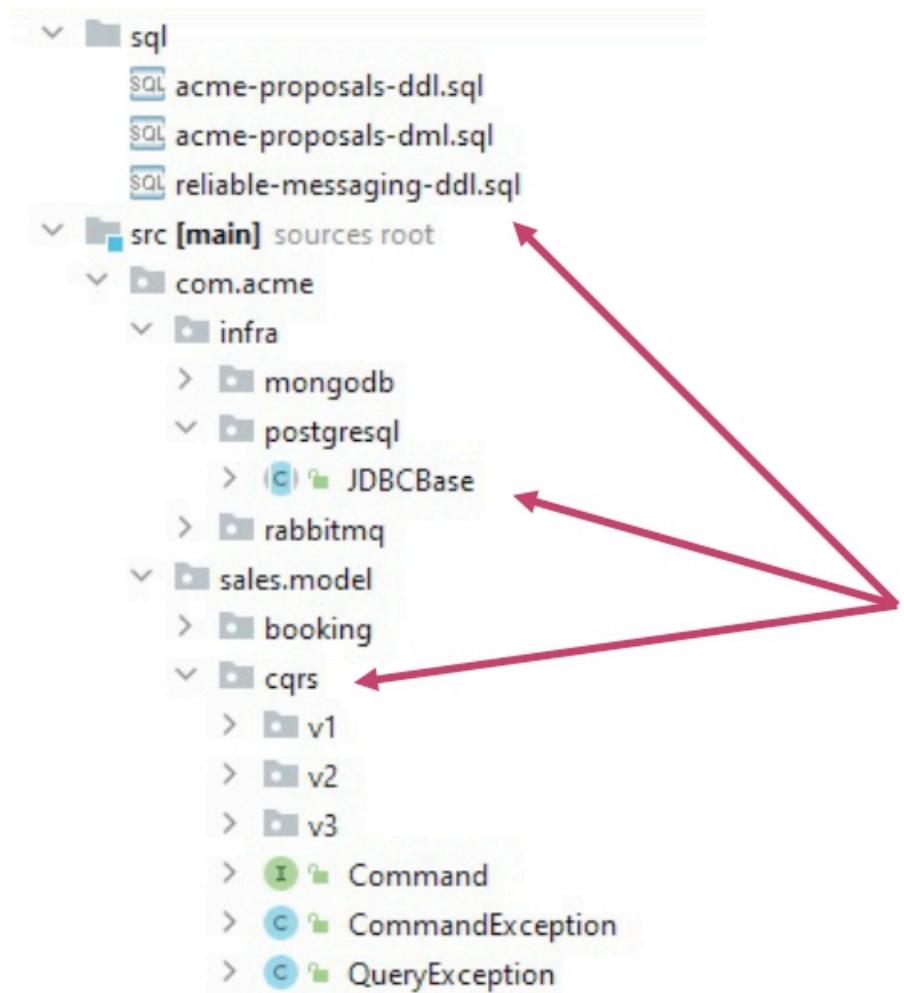


- 1 Need for a PoC
- 2 Classes for implementing the CQS
- 3 Test the Command & Query

Note: Use the PostgreSQL Database setup in previous lecture

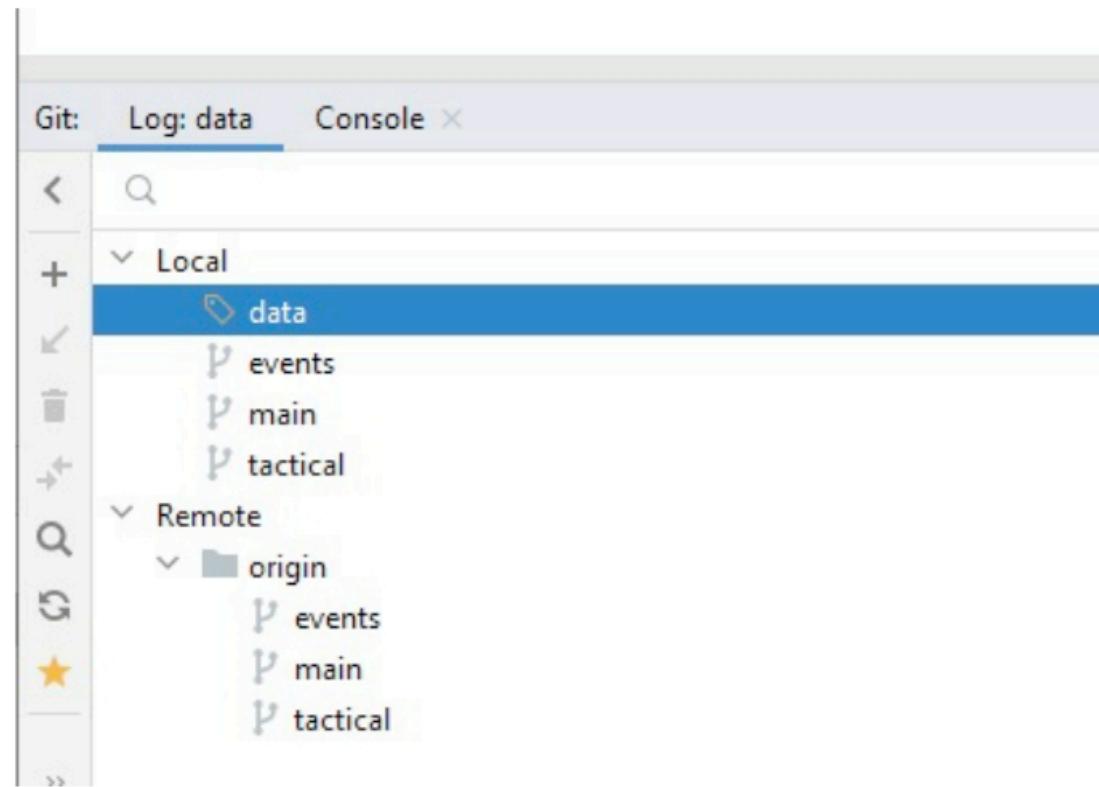
# Code

Checkout the branch : data



- Pre - Requisites: JDBC & JSON

- Instructions in README.md





## Need a working model to understand CQS

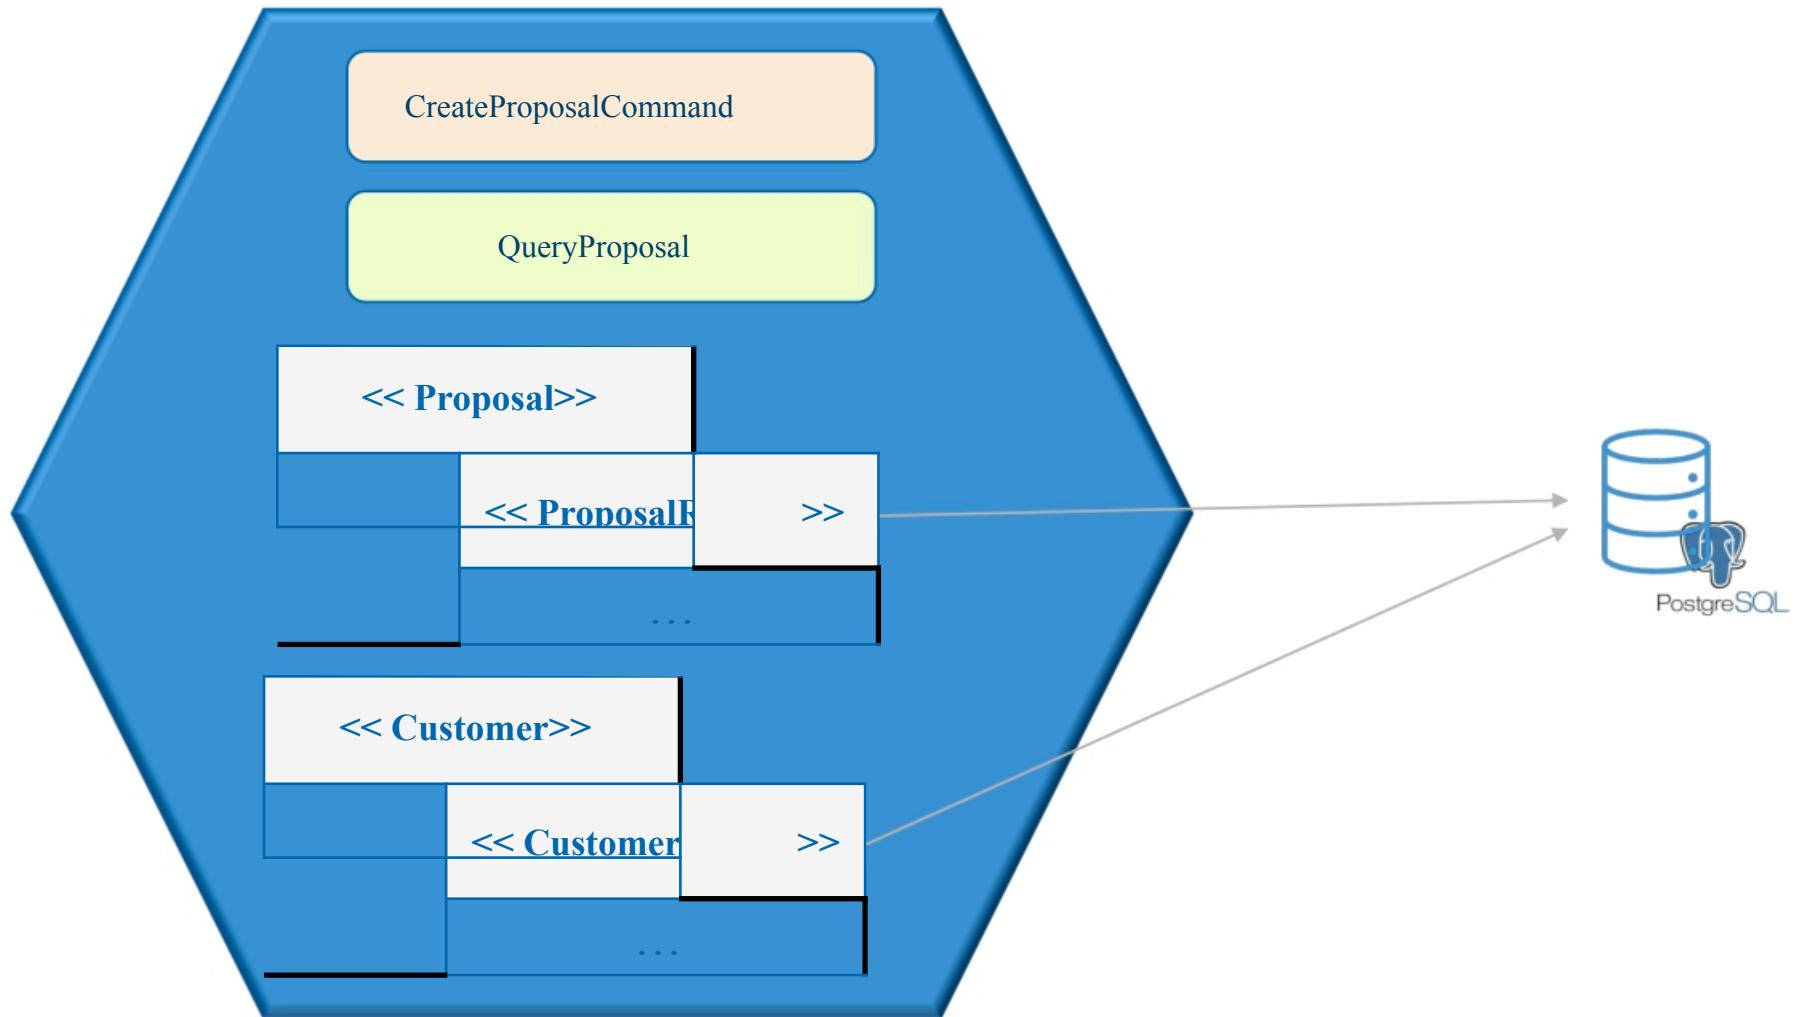
Create the command & query for proposals

Create the proposal for customer in Database

Start with some common queries but make it flexible for future queries



## Implementation components



## Setup the DB tables : customer table

sql/acme-proposals-ddl.sql

```
CREATE TABLE customers (
    customer_id      INT generated by default as identity,
    email_address    varchar(255) not null,
    fname            varchar(64) not null,
    mname            varchar(64) not null,
    lname            varchar(64) not null,
    phone            varchar(10) not null,
    PRIMARY KEY(customer_id)
);
```



## Setup the DB tables : proposals Table

sql/acme-proposals-ddl.sql

```
CREATE TABLE proposals (
    proposal_id      INT generated by default as identity,
    customer_id      INT not null,
    package_id        varchar(64) not null,
    pax_0              varchar(64) not null,
    pax_age_0          INT not null,
    pax_1              varchar(64),
    pax_age_1          INT,
    pax_2              varchar(64),
    pax_age_2          INT,
    pax_3              varchar(64),
    pax_age_3          INT,
    pax_4              varchar(64),
    pax_age_4          INT,
    PRIMARY KEY(proposal_id) ,
    CONSTRAINT fk_customer
        FOREIGN KEY(customer_id)
            REFERENCES customers(customer_id)
);
```



## Create & Populate the tables

1. Create tables      [sql/acme-proposals-ddl.sql](#)

2. Populate tables      [sql/acme-proposals-dml.sql](#)



## com.acme.sales.model.cqrs.v1.command

In this implementation each command implemented as a object \*but\* you may implement commands as operations on a single command object.

**C** CreateProposalCommand

implements uses

## com.acme.sales.model.cqrs

**C** QueryException

**I** Command

process()

throws

**C** CommandException

## com.acme.sales.model.tests.cqrs.v1

**C** TestCommand

**C** TestQuery

## com.acme.sales.model.cqrs.v1.query

**C** ProposalQuery

This is the Query implementation that uses the SAME model as the command but implemented as an independent component within the same unit

## com.acme.sales.model.cqrs.v1.repo.jdbc

**C** CustomerRepo

**C** ProposalRepo

Extends Extends

## com.acme.infra.postgresql

**A** JDBCBase

uml/ cqrs/cqrs.v1.class.puml

# CQRS



Command Query Responsibility Segregation

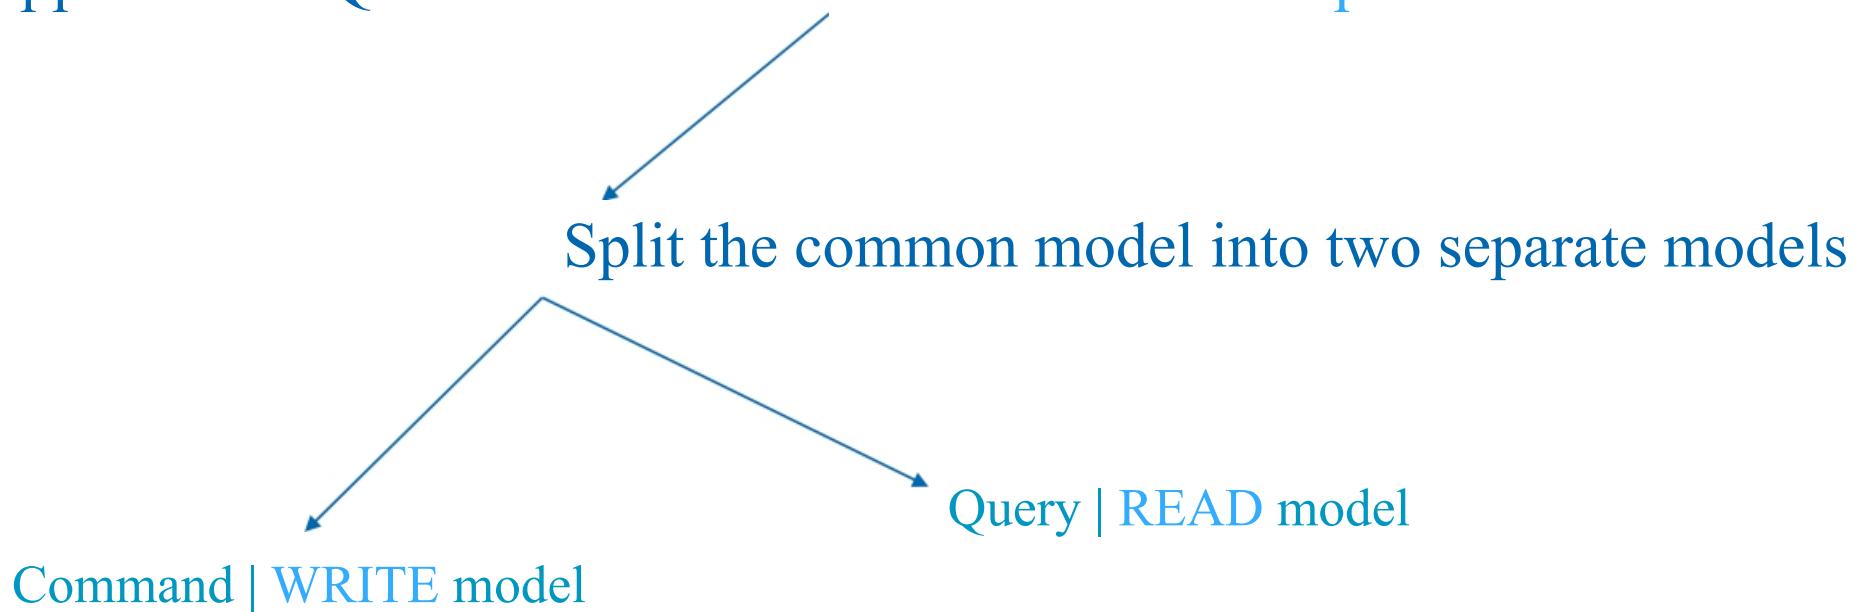


- 1 Define CQRS
- 2 Benefits
- 3 Realization

# CQRS

## Command Query Responsibility Segregation

- Applies the CQS idea to Domain model rather than operations



Suggested By: Greg Young

<https://martinfowler.com/bliki/CQRS.html>

“

Suggests the use of separate **READ** and **WRITE** Models for the realization of Collaborative domains



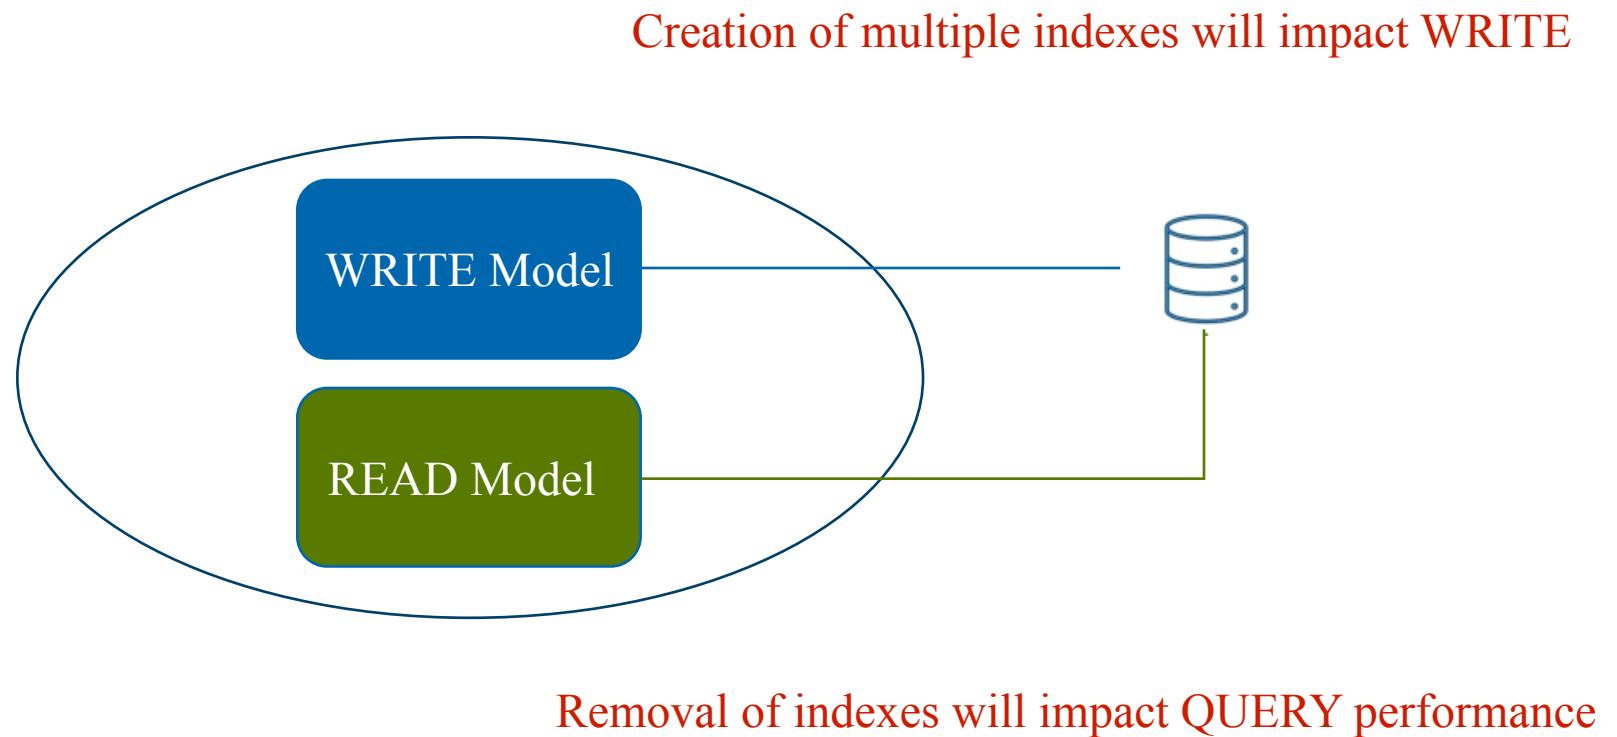
- Finer Granularity
- Frequent changes to Business Rules



- Meet different **READ** requirements
- High performance

## Common Datastore for READ & WRITE

READ | WRITE models may use the same instance of Datastore

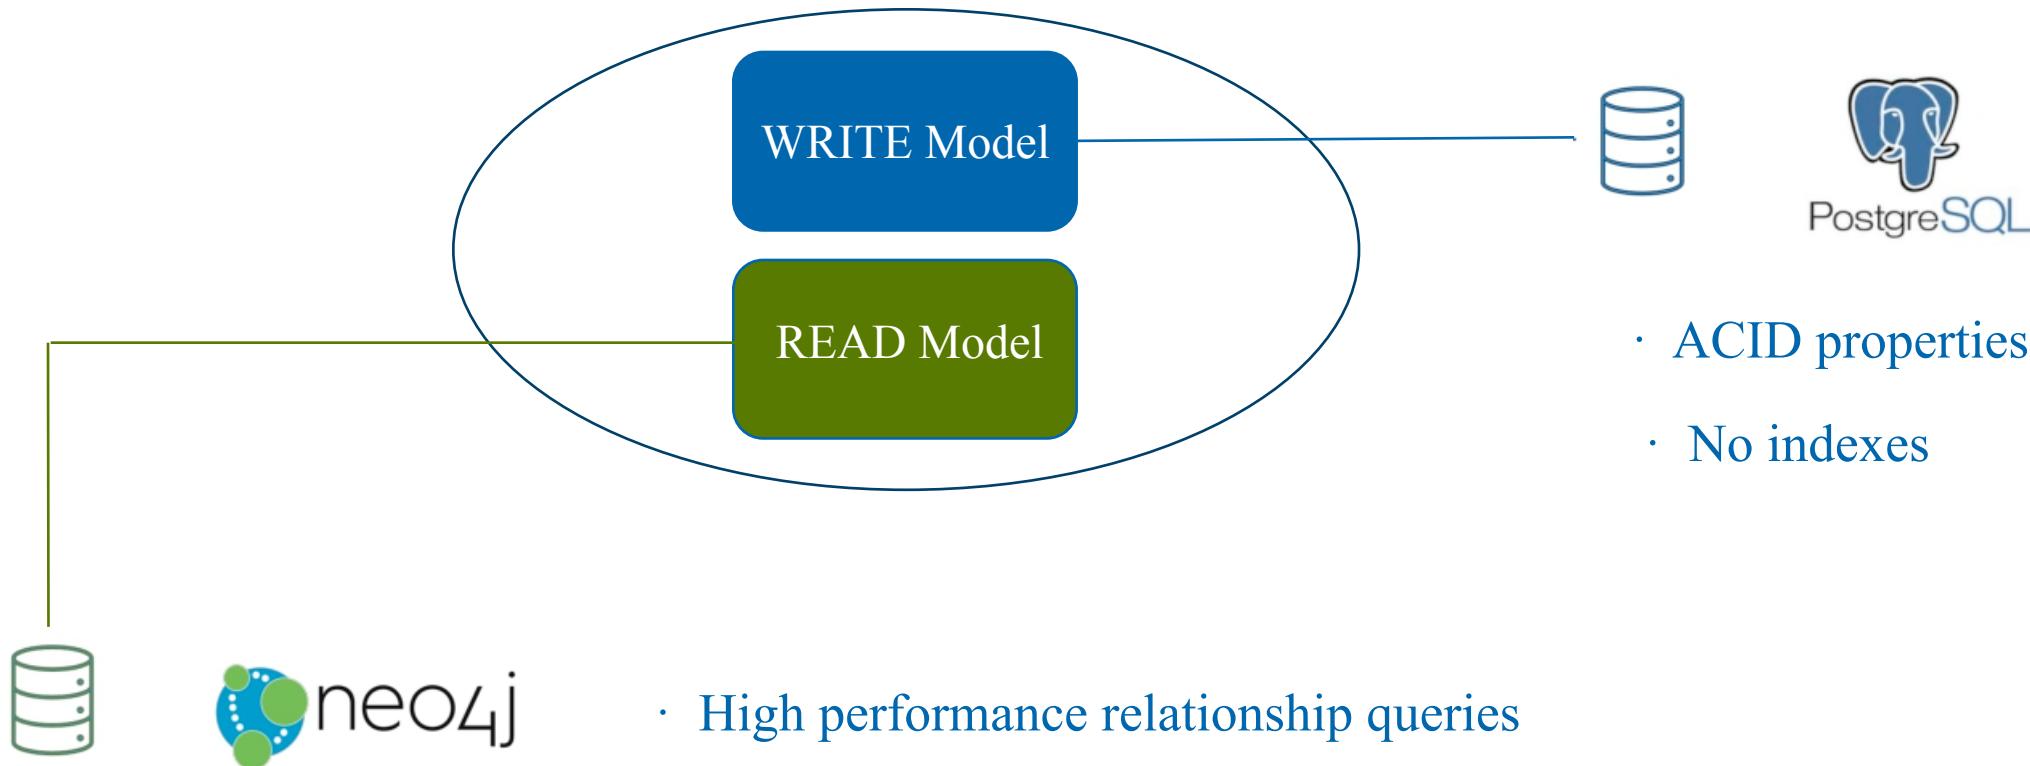


Loss of ability to optimize the DB for READ | WRITE

# Independent Datastores & Polyglot persistence

READ | WRITE models may use independent Datastores

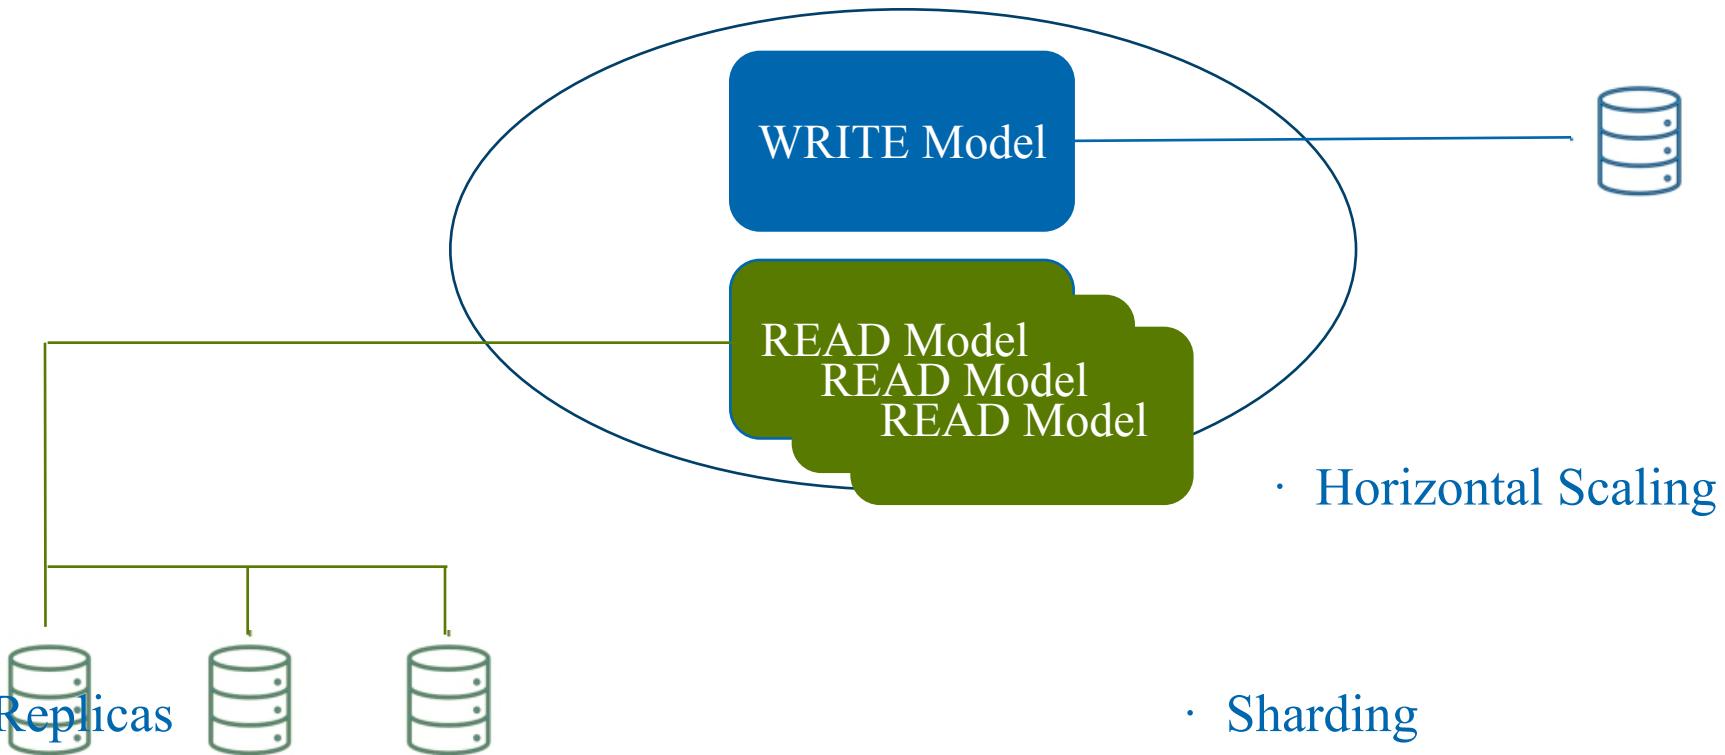
- May use different Database technologies



## Control on Scaling

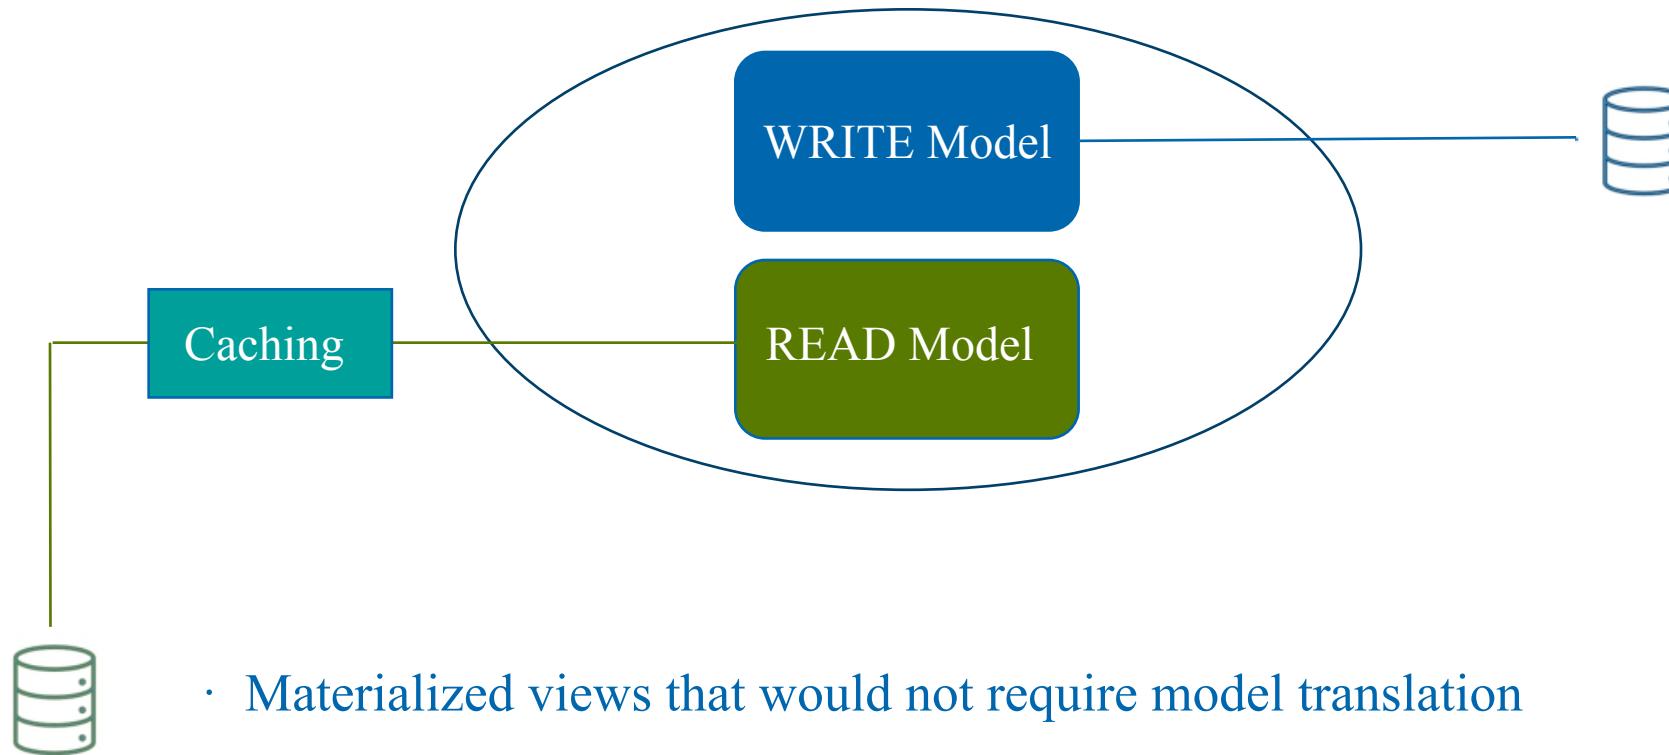
READ | WRITE side may scale independently

- Example: Read heavy application



## READ Performance considerations

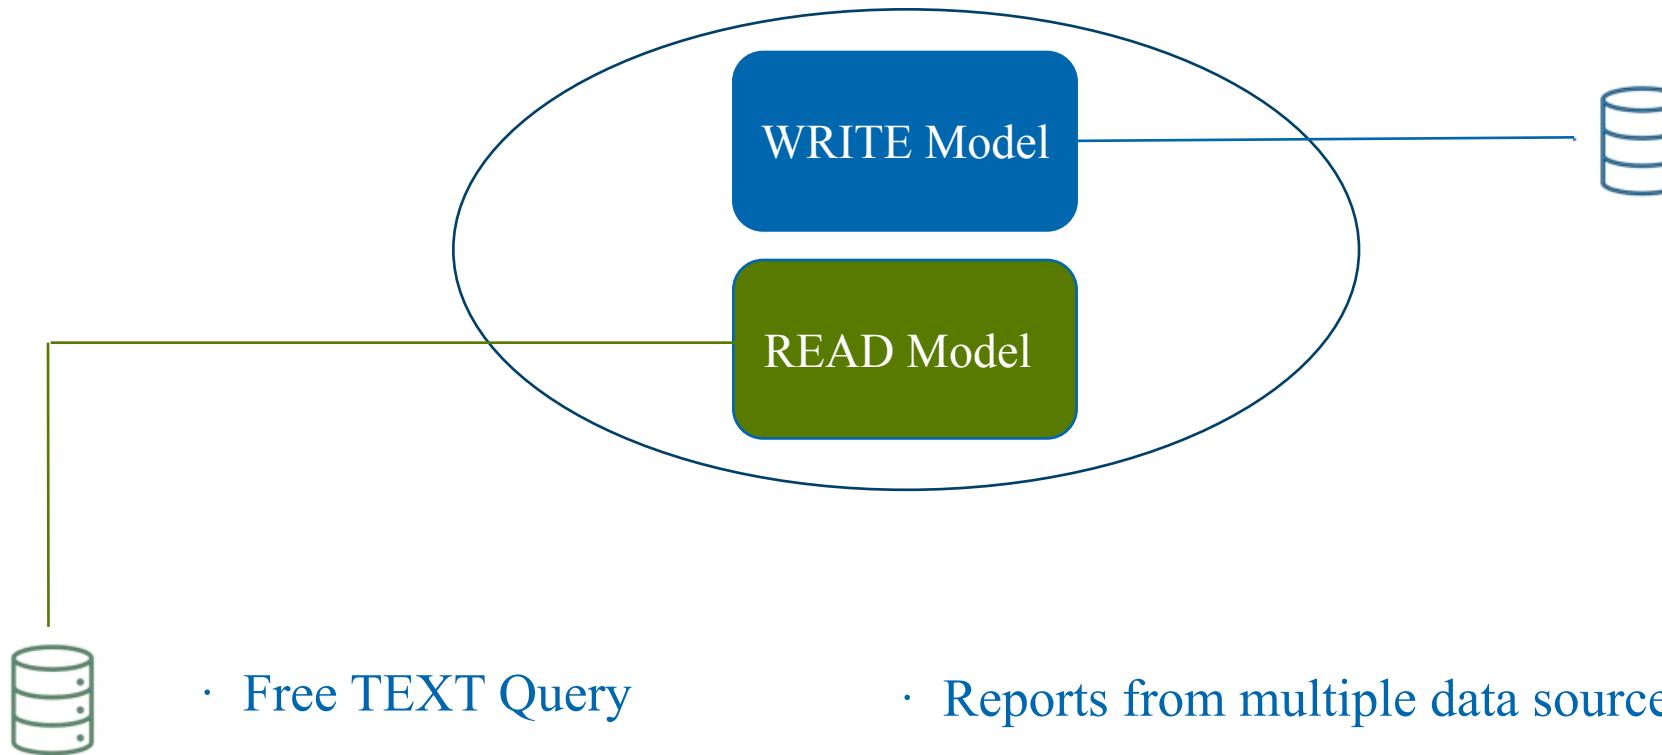
Typically, there are more READ than WRITE (~ 85% Reads)



## READ Performance considerations

Typically, there are more READ than WRITE (~ 85% Reads)

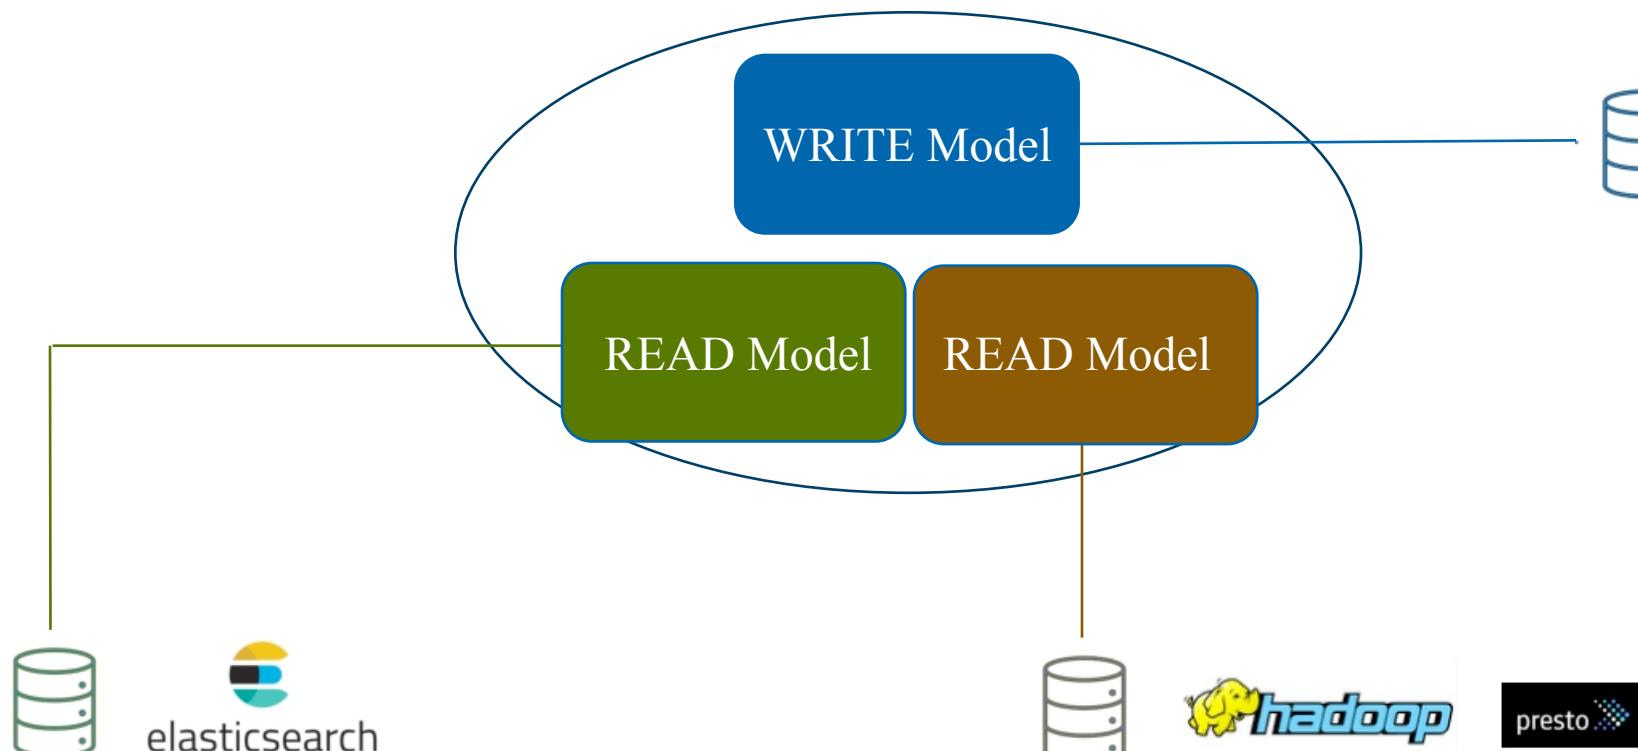
- Example: Multiple query requirements



## READ Performance considerations

Typically, there are more READ than WRITE (~ 85% Reads)

- Example: Multiple query requirements

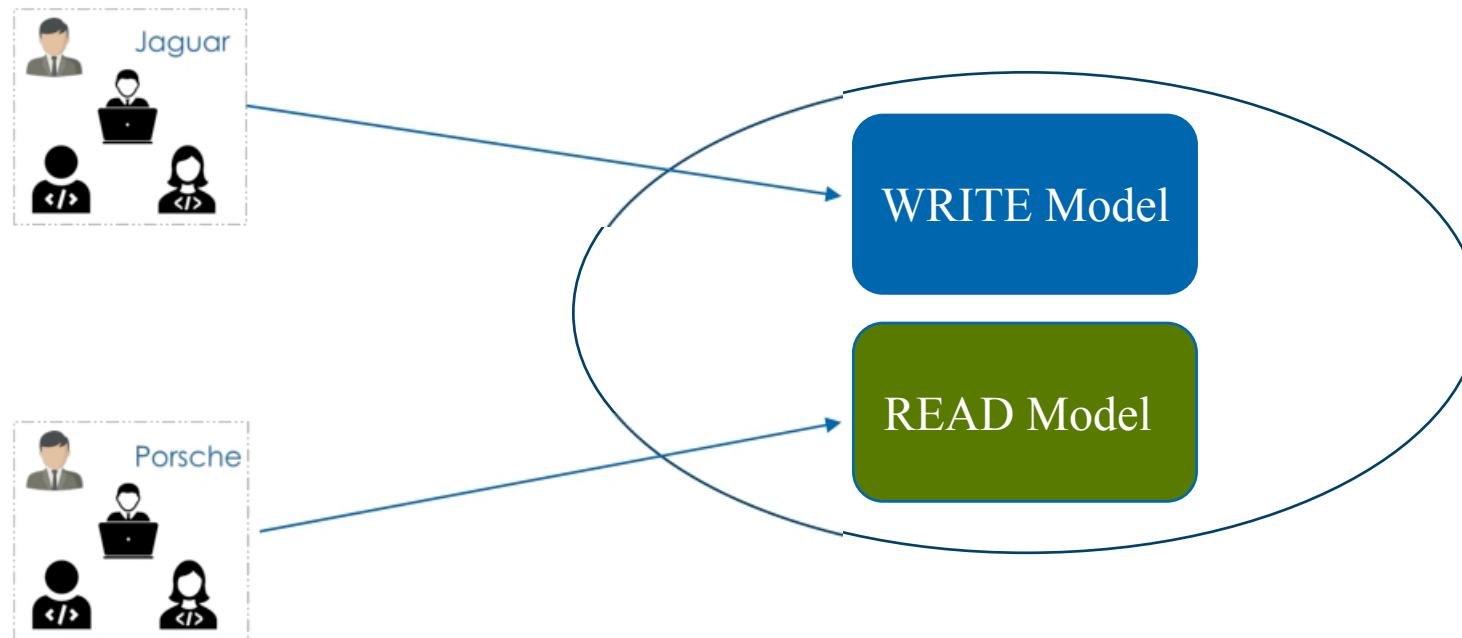


- Free TEXT Query

- Reports from multiple data sources

## Independent Change Management

READ | WRITE models may be managed by different teams



- Independent code bases
- Independent changes
- Independent deployments

Collaborative | Dynamic apps require frequent WRITE changes

## Misconception



In DDD you MUST always use CQRS

- Considerations
  - What is the advantage
  - HIGH Cost of solution
  - More moving parts | components



## Quick Review

Consider using CQRS for collaborative domains

- READ | WRITE independently control:
  - Performance
  - Scalability
  - Changes

# CQRS : Data Replication



Data needs to be replicated between WRITE & READ side

---

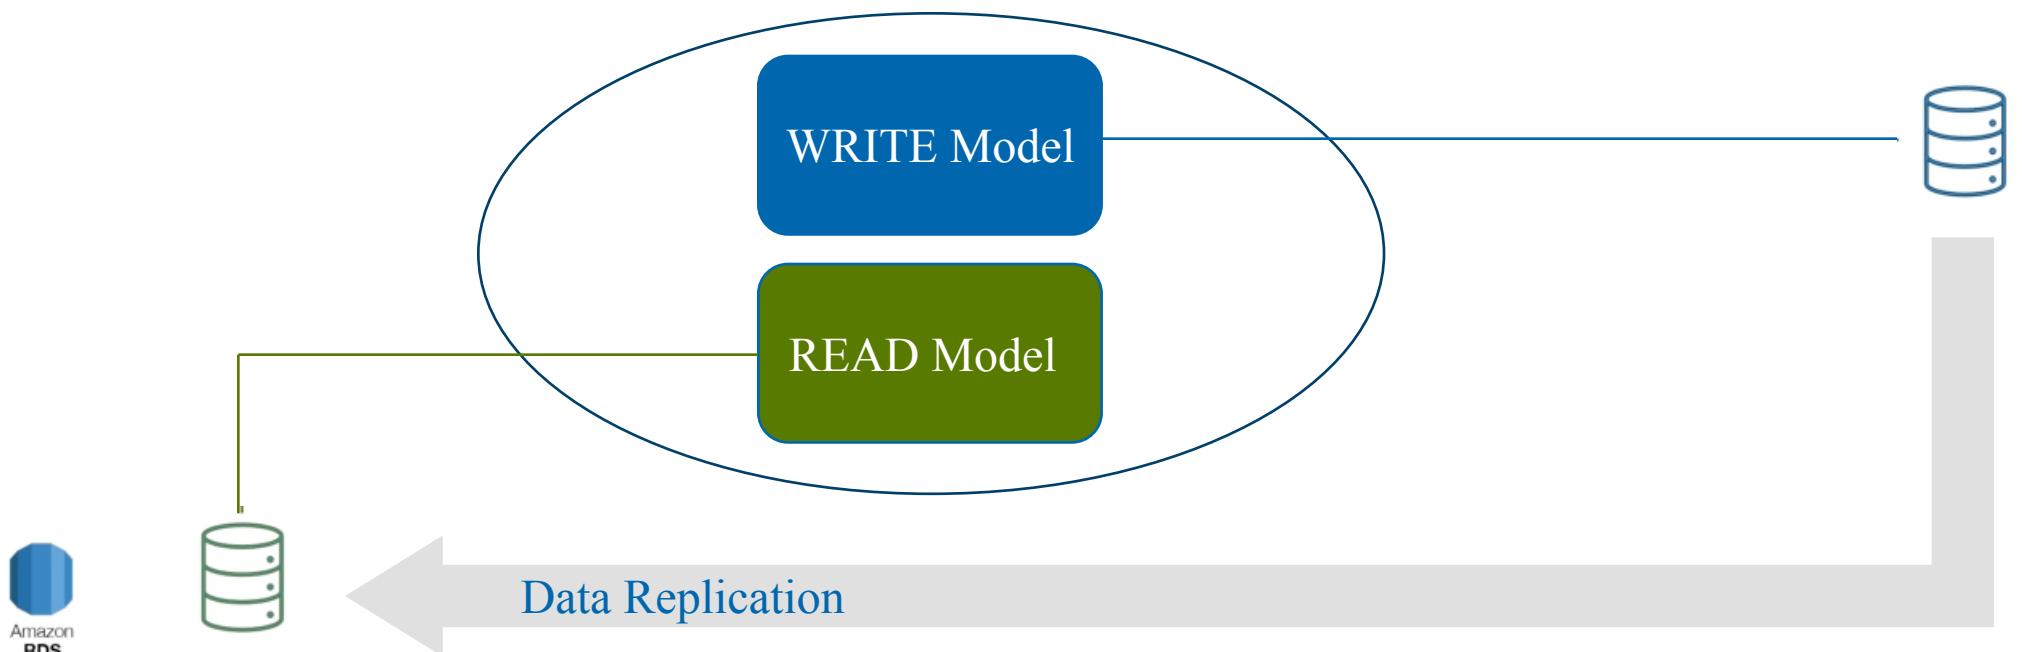


- 1 Need for Replication
- 2 Synchronous | Asynchronous Replication
- 3 Data Replication options

## Data Replication Considerations

NEEDED for independent READ | WRITE Data stores

### Synchronous Replication

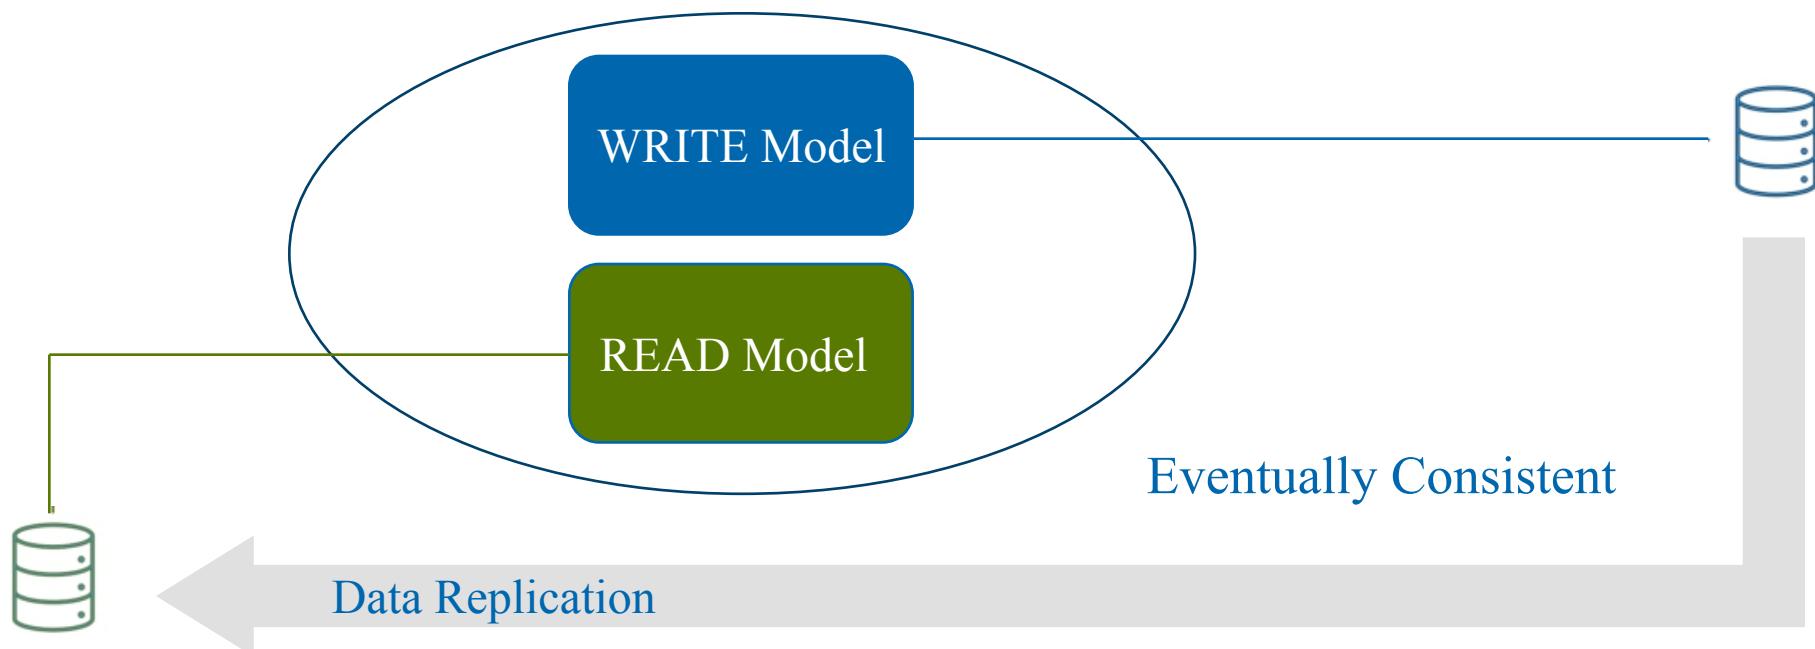


Technical feasibility | May impact WRITE Performance

## Data Replication Considerations

NEEDED for independent READ | WRITE Data stores

### Asynchronous Replication



READ side may NOT reflect the current state !!!

## Data Replication Options

Select based on use case | requirements

- Change Data Capture (CDC)



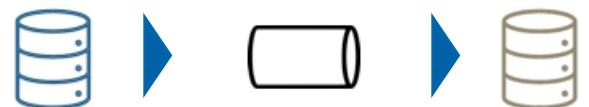
- Native replication



- 3rd party tools



- Messaging | Streaming | Events





## Quick Review

Replication needed between READ | WRITE data stores

- Synchronous | Asynchronous
- Option depends on requirements & supporting technologies

# CQRS



Command Query Responsibility Segregation

---



- 1 Define CQRS
- 2 Benefits & Considerations
- 3 Data Replication

## CQRS Frameworks



AxonIQ

eventuate

<https://github.com/topics/cqrs - framework>

# CQRS : Proposal Management

Implementation of CQRS for Proposals



- 1 Proposal Write/Read Requirements
- 2 Component Diagram
- 3 Sequence Diagram



IT Lead

## Proposal Management Design



Should it be implemented with CQRS



IT Lead



John, Travel Advisor



Travis, Business Intelligence



Jack, Marketing & Sales

We are in a highly competitive industry; performance matters

Business rules change almost every week

We need access to ALL the proposals data - new & historical

Management dashboards need to be real time and performant

We use current & historical data to understand market trends

Fast access to data is essential to be able to send offers

# Implement Proposal Management using CQRS

Dynamic Business Requirements

Multiple types of READ/Query requirements

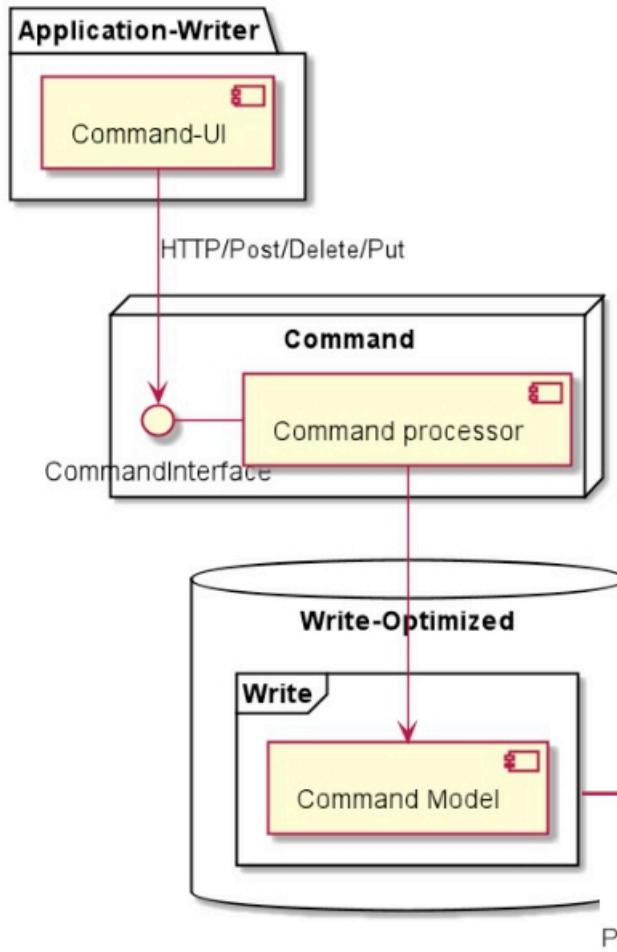
Expectation for High performance

Opportunity to achieve competitive advantage

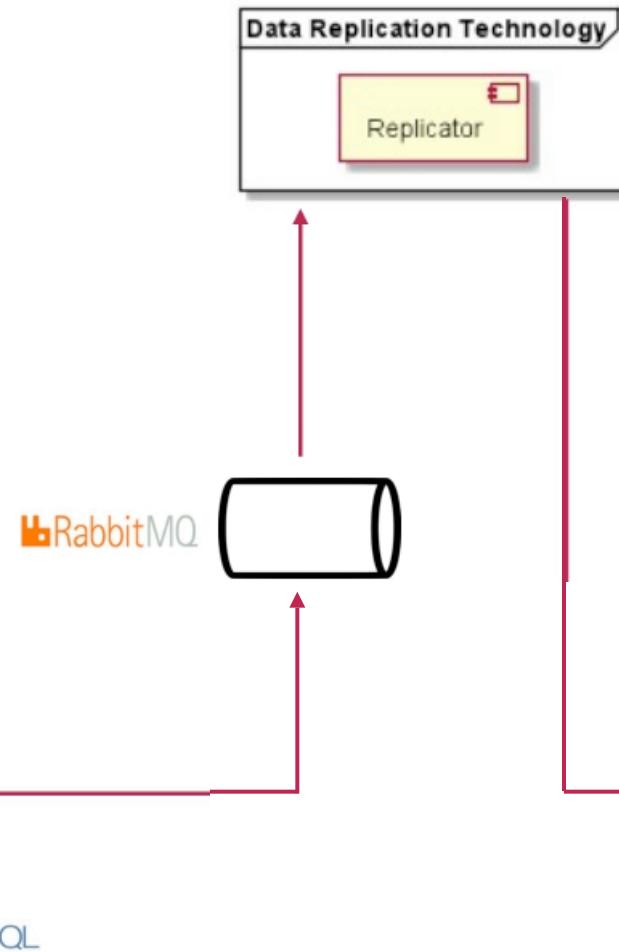
# Proposal CQRS Deployment

uml / cqrs /cqrsv2.component.puml

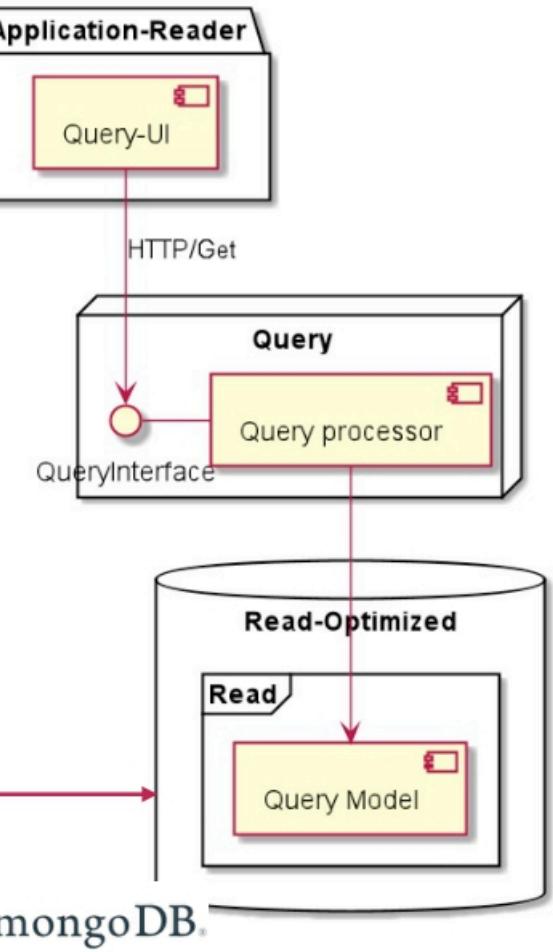
## Command



## Event Subscriber

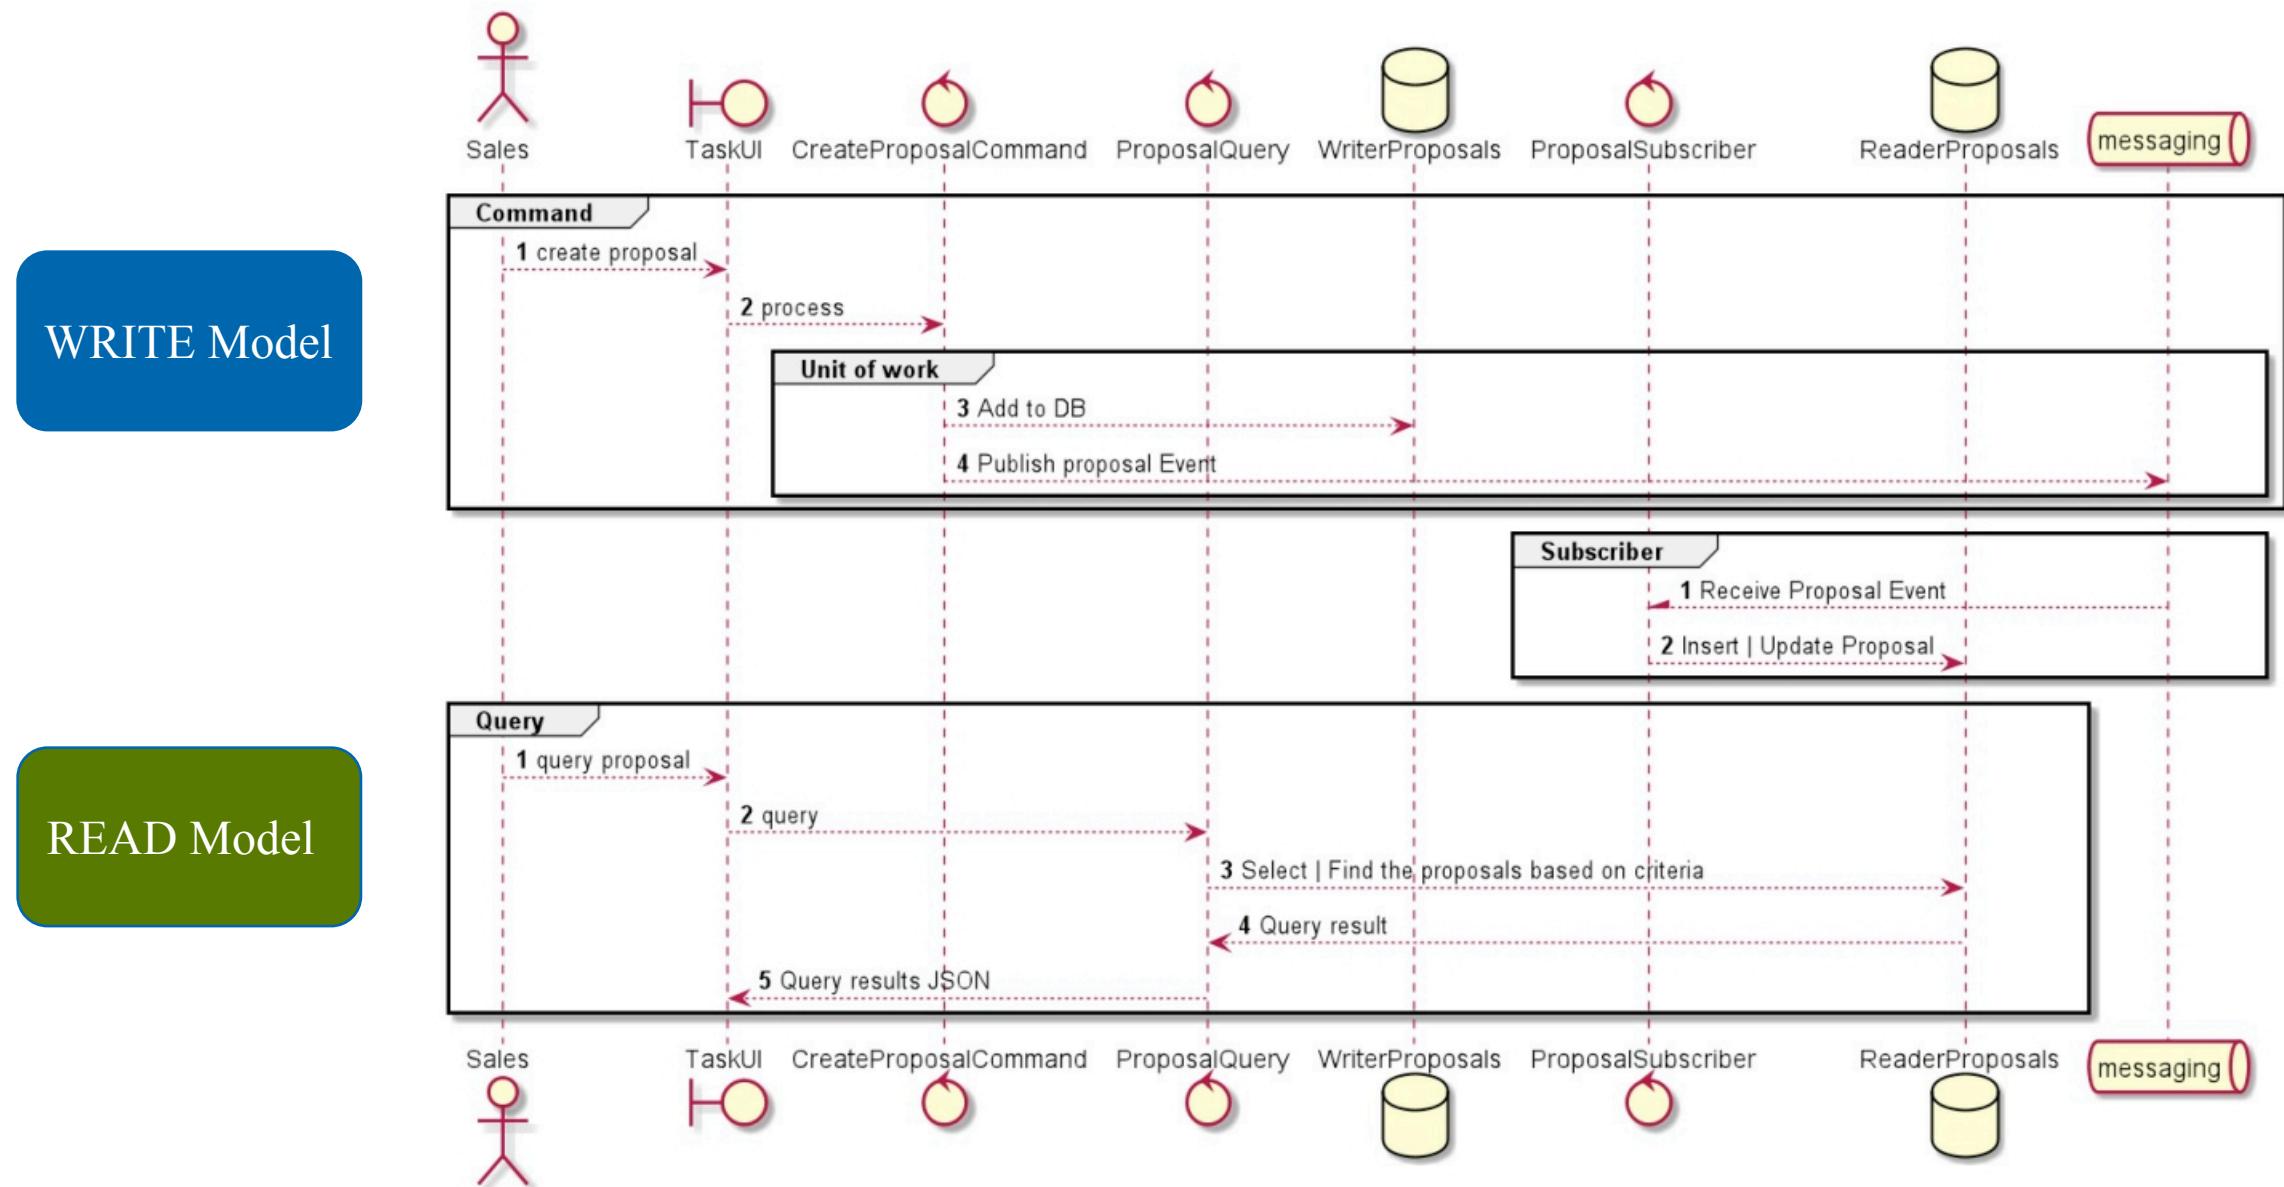


## Query



# Proposal CQRS : Replication with MQ

uml / cqrs /commandquery.v2.sequence.puml



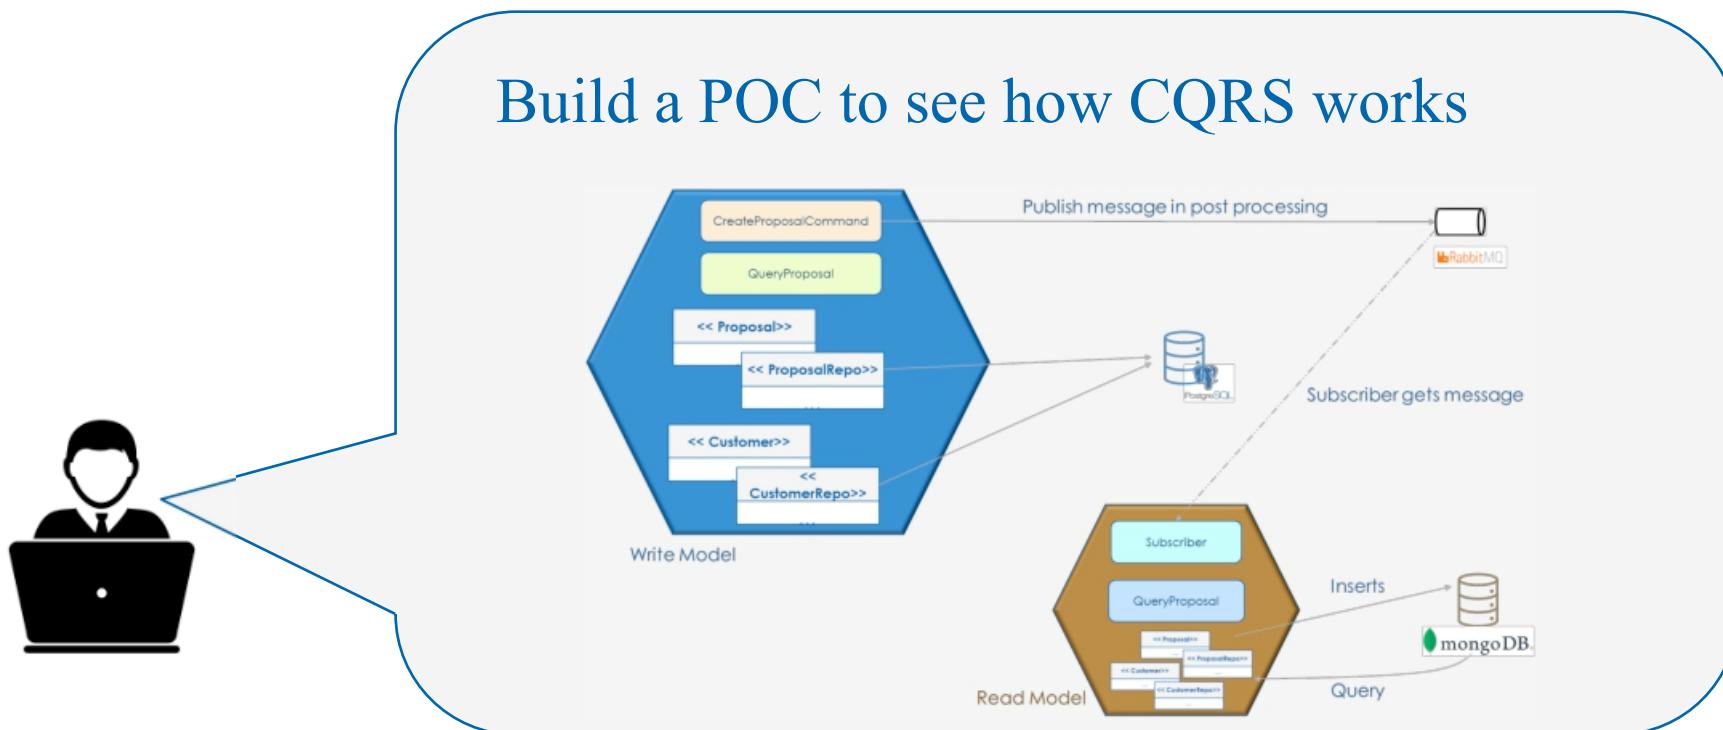


## Quick Review

### Decision to implement CQRS

- Depends on Use Case
- READ | WRITE characteristics

Build a POC to see how CQRS works



IT Lead

# CQRS: Build & Test Command

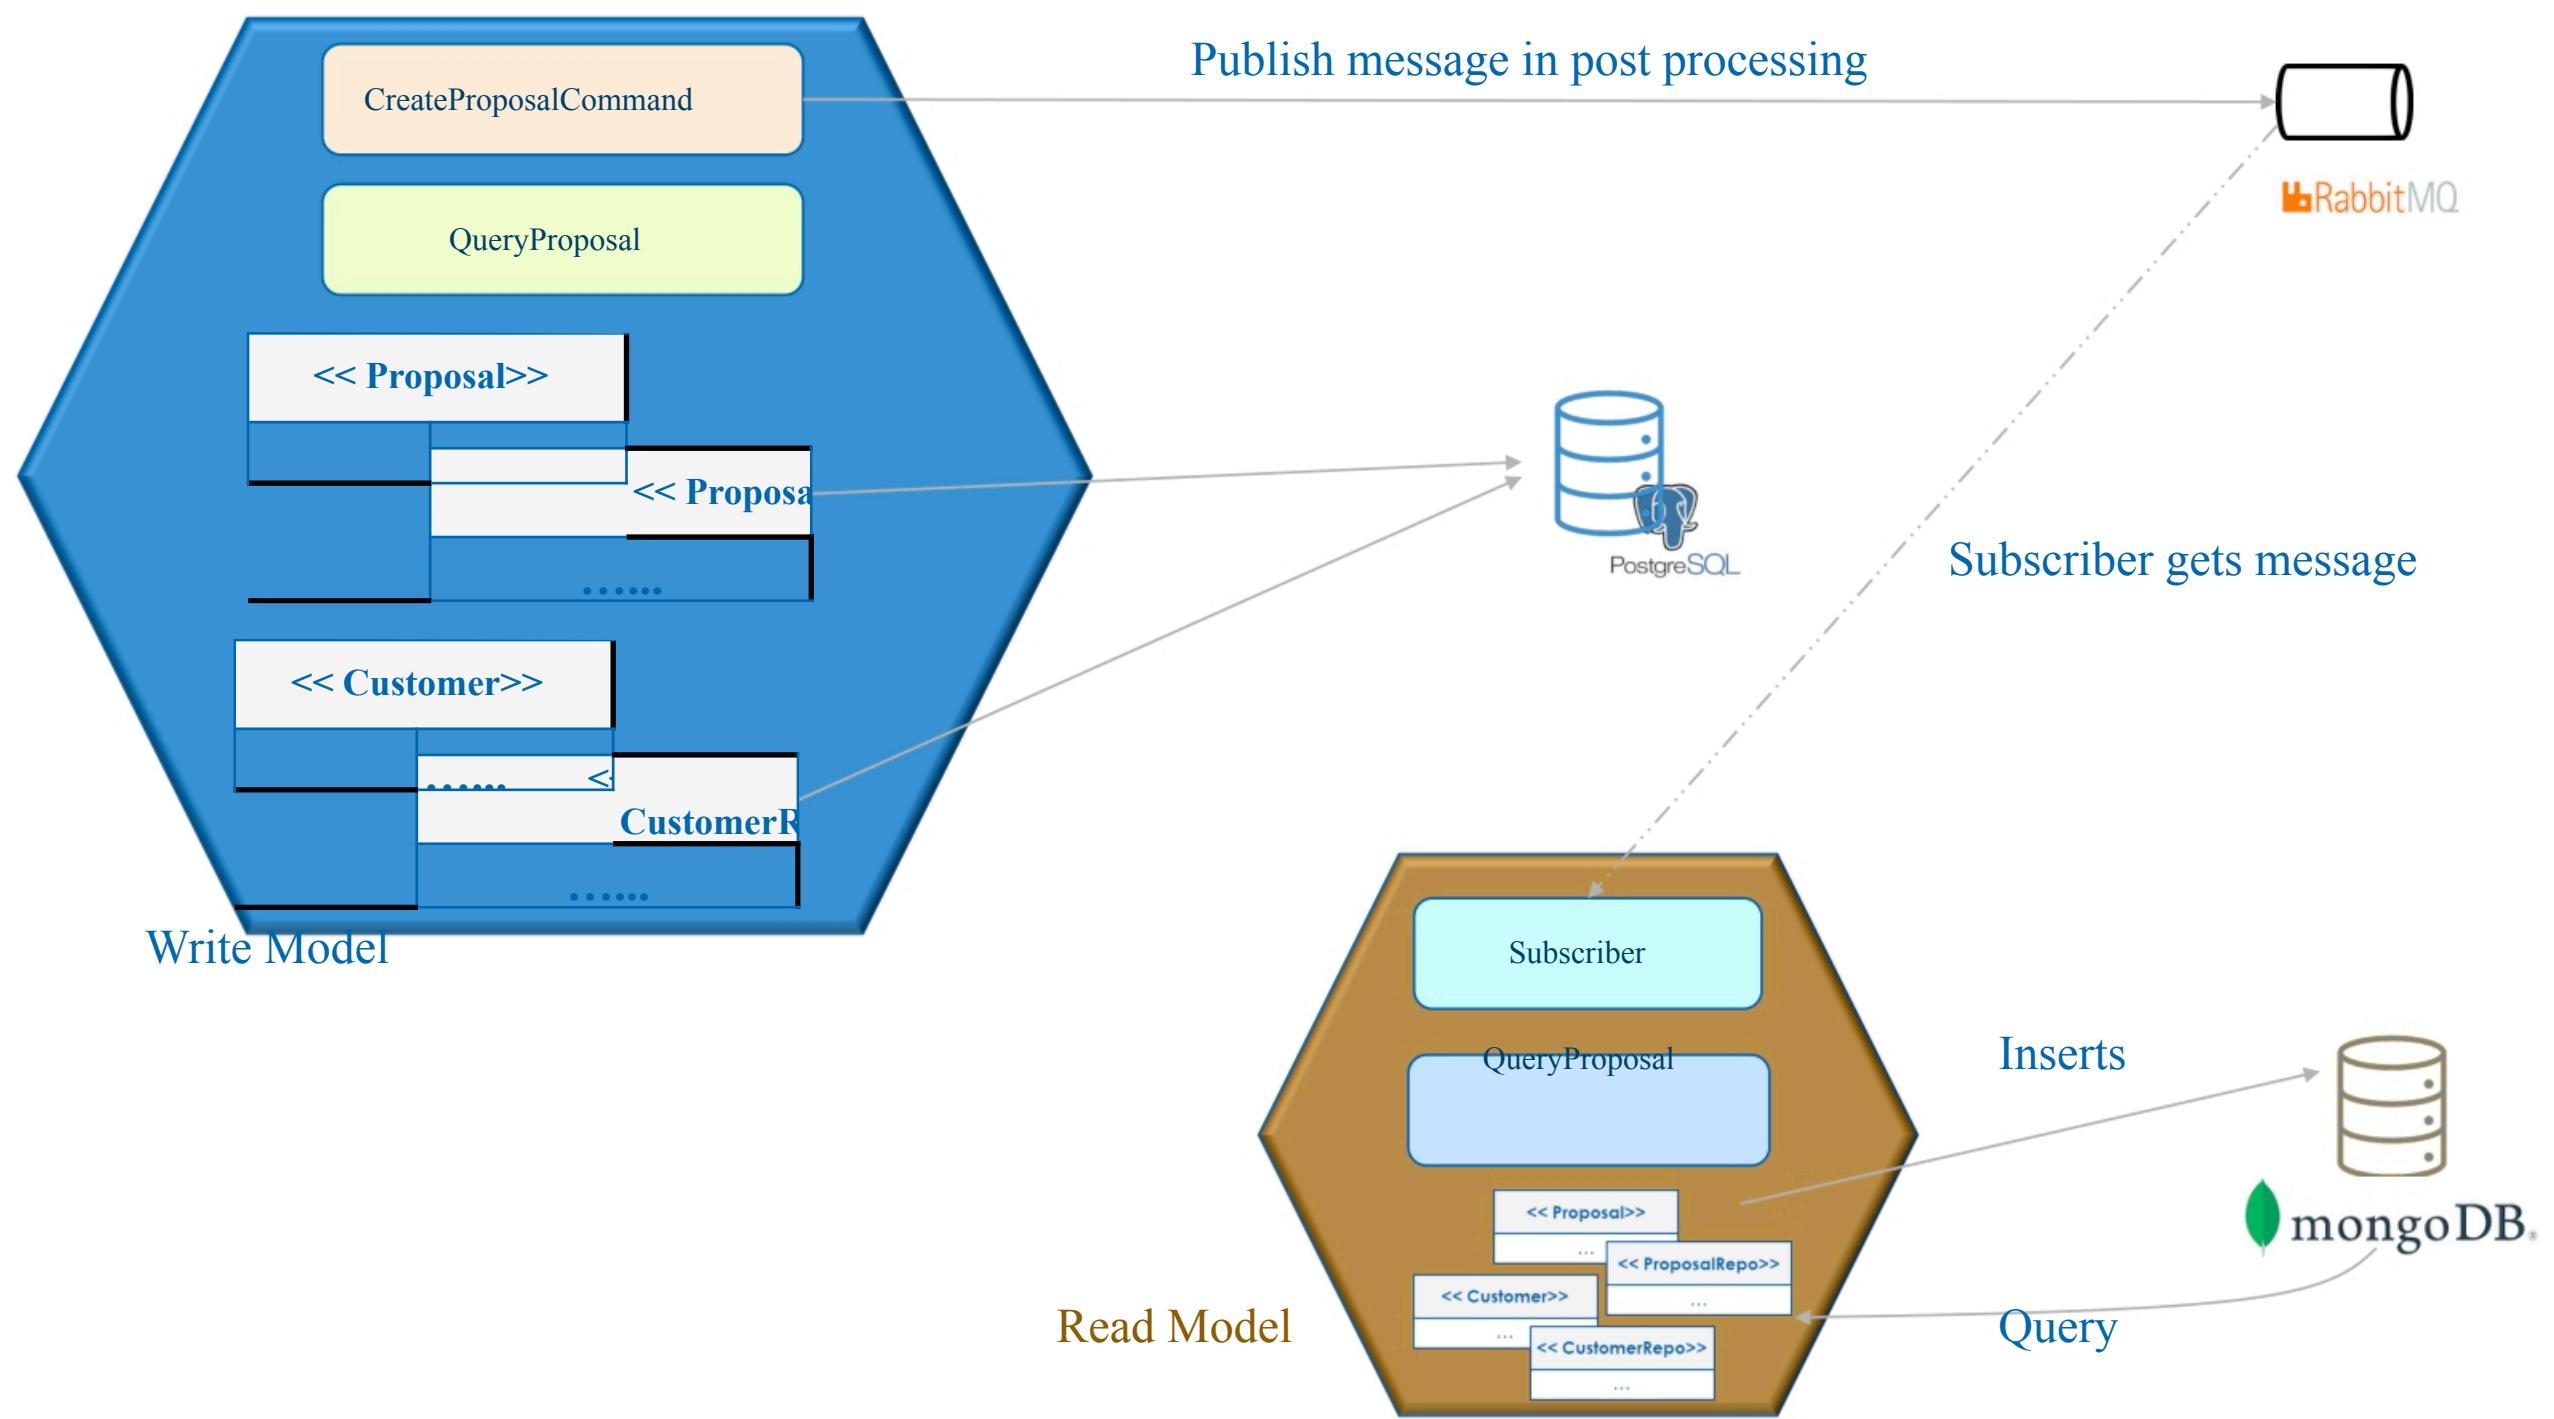


Command will write out an event to the queue



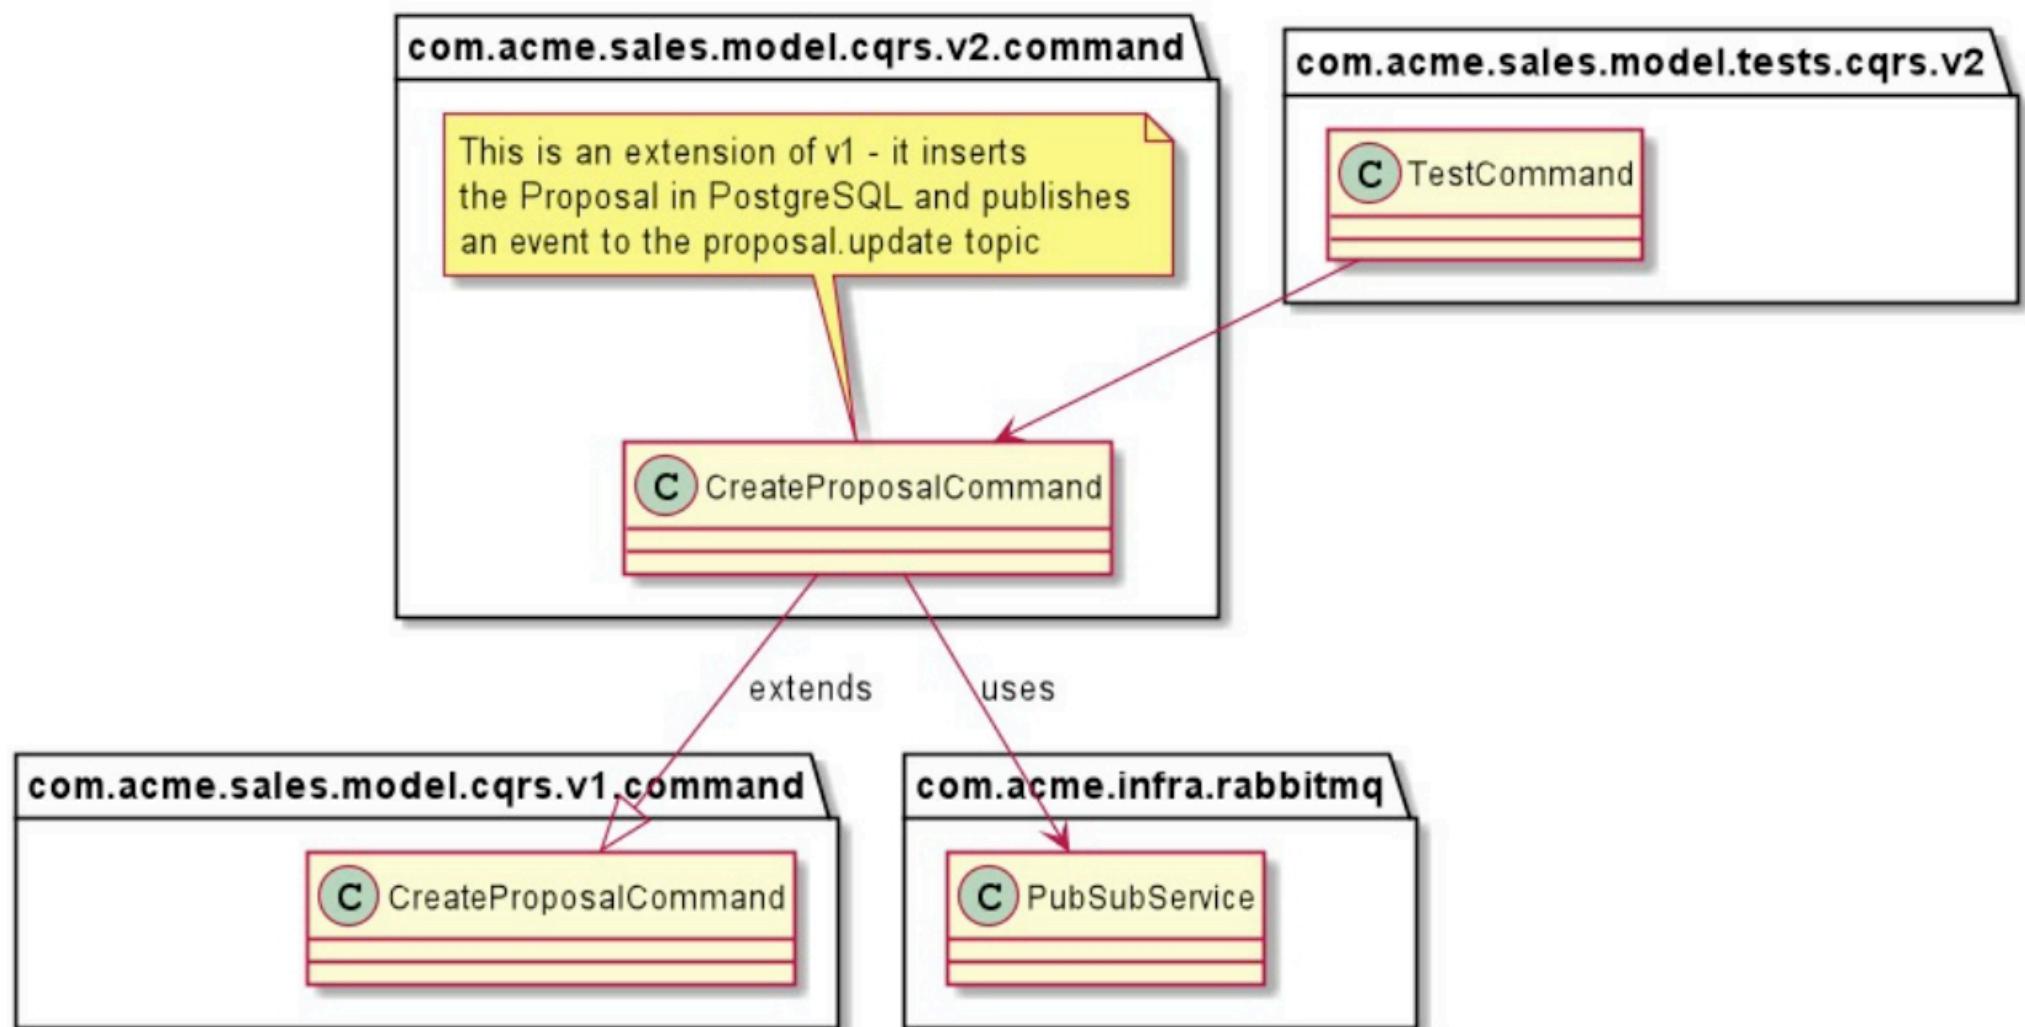
- 1 Setup Rabbit MQ for event messaging
- 2 Walkthrough of code for Command v2
- 3 Test the Command v2

Note: Uses the PostgreSQL Database & Rabbit MQ



# Command Classes

uml/ cqrs/cqrs.v2.class.puml



## RabbitMQ setup



Queue	proposal.reader.queue
Topic Exchange	acme.sales.topic
Queue Binding	proposal.update

# Setup MongoDB for READ side



Read Repository will use the MongoDB as datastore

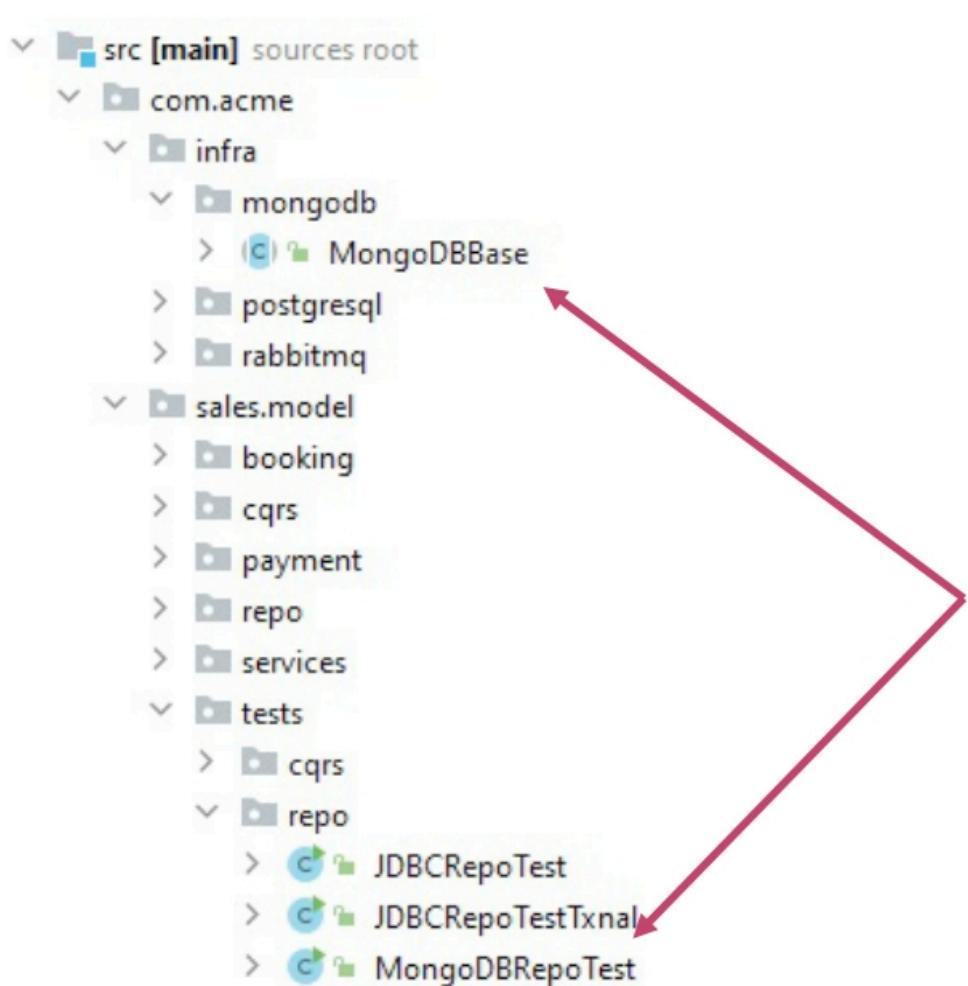


- 1 Setup MongoDB on [cloud.mongodb.com](http://cloud.mongodb.com)
- 2 Walkthrough of MongoDBBase class
- 3 Test the MongoDB setup

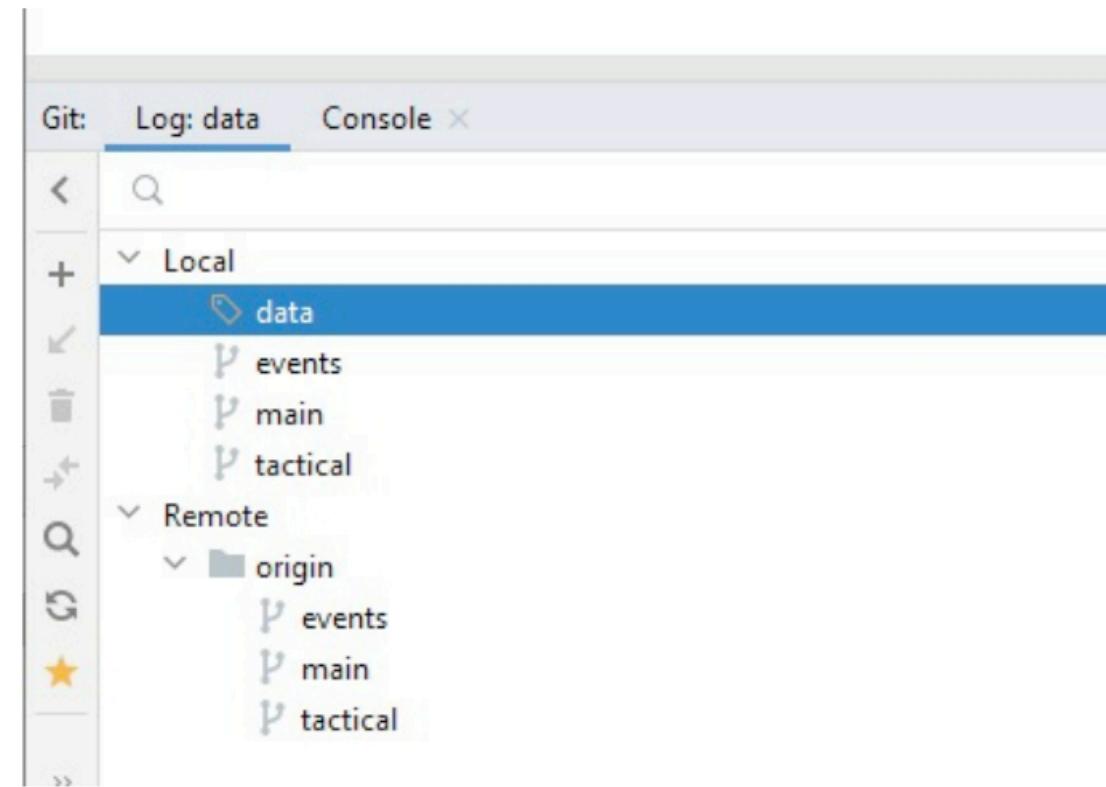
Note: Intent is NOT to teach MongoDB

# Code

Checkout the branch : data

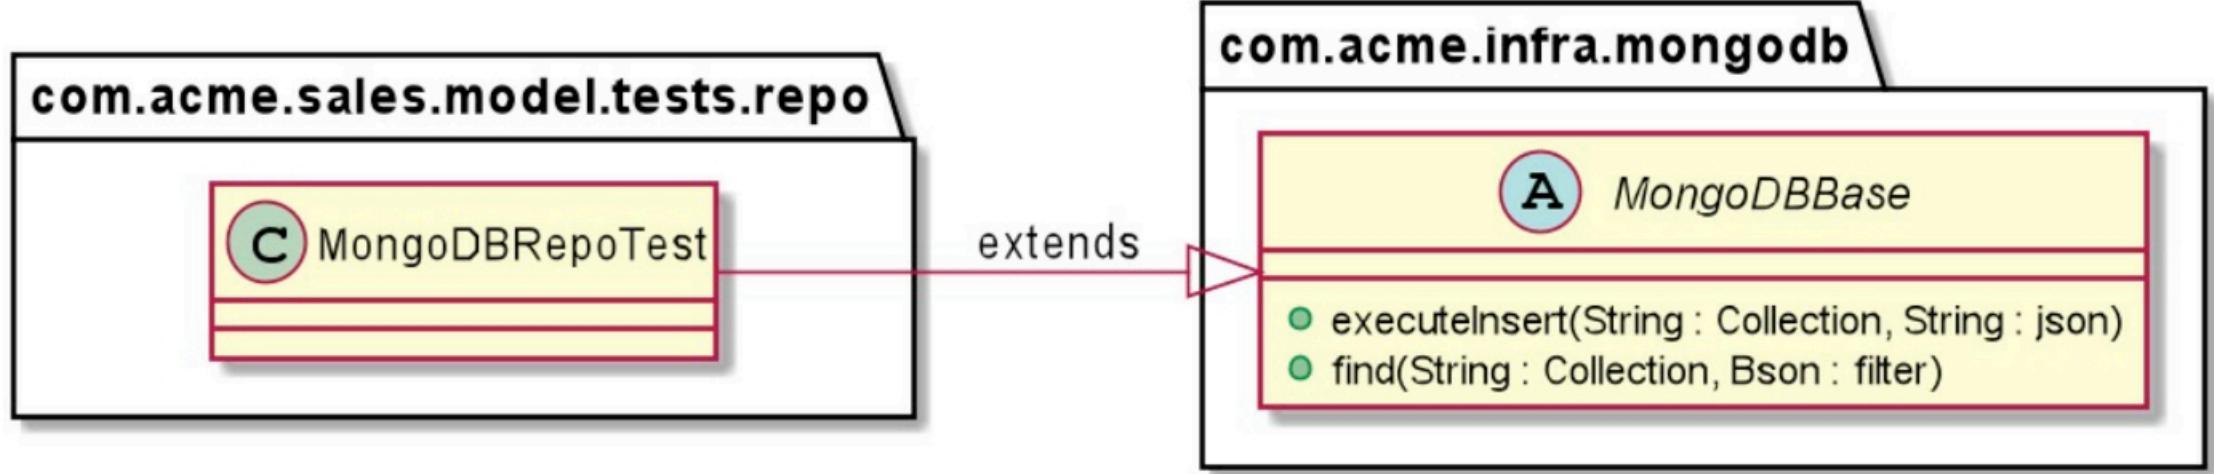


- Basic understanding of NoSQL Databases
- Instructions in README.md



# MongoDBBase Class

uml/ cqrs/cqrs.v2.mongodb.class.puml





Create an instance of DB on

[cloud.mongodb.com](https://cloud.mongodb.com)

Cluster Name      AcmeTravel

## Setup the database instance

Database user  
repouser

Database password  
repouser

# Event Sourcing



Persisting the events in a data store

---



- 1 State - Oriented versus Event- Sourced persistence
- 2 Event Stores
- 3 Benefits of Event Sourcing

## State Management

State -Oriented persistence system maintains only the current state

*Example:*



- Number of "Item x" in the inventory
- Item Added => Count increased
- Item Sold => Count decreased

Event -Sourced persistence systems persist all state changes



- Domain events are stored as they are received

## Event Sourcing

“

Event sourcing suggests persisting the domain events in an Event store and using it for recreating the state of the domain objects at any point in time

Event store is an append only event log used for persisting the received events

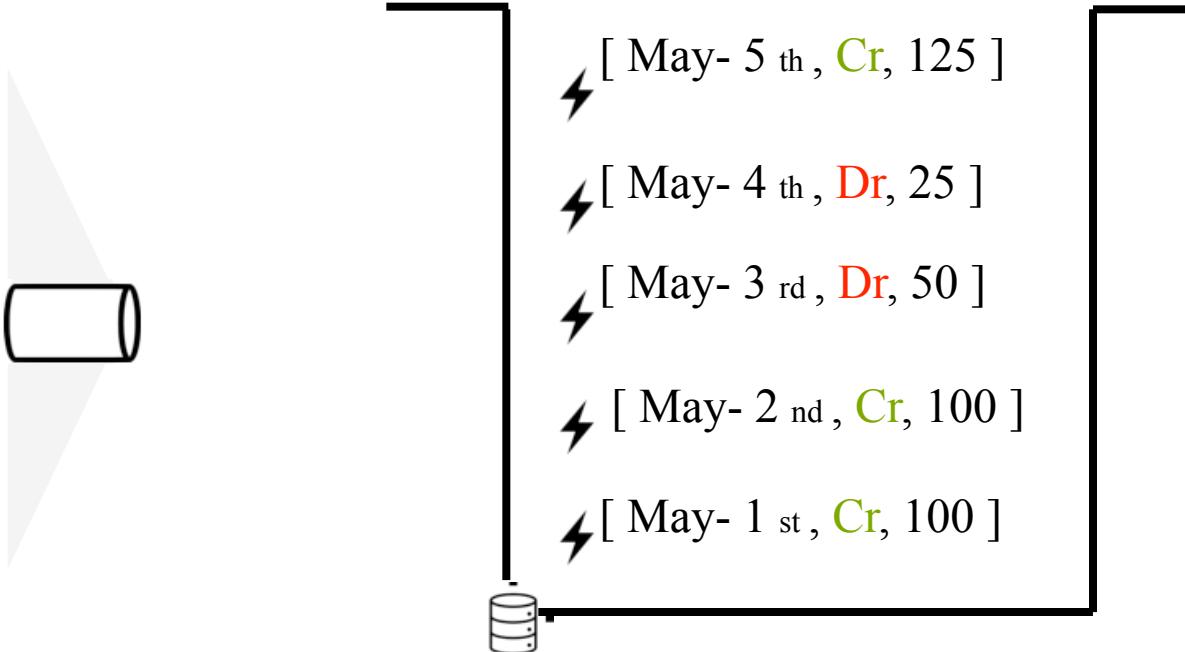
## Example : Bank Transactions

READ side persists events emitted by WRITE side in an Event Store

WRITE

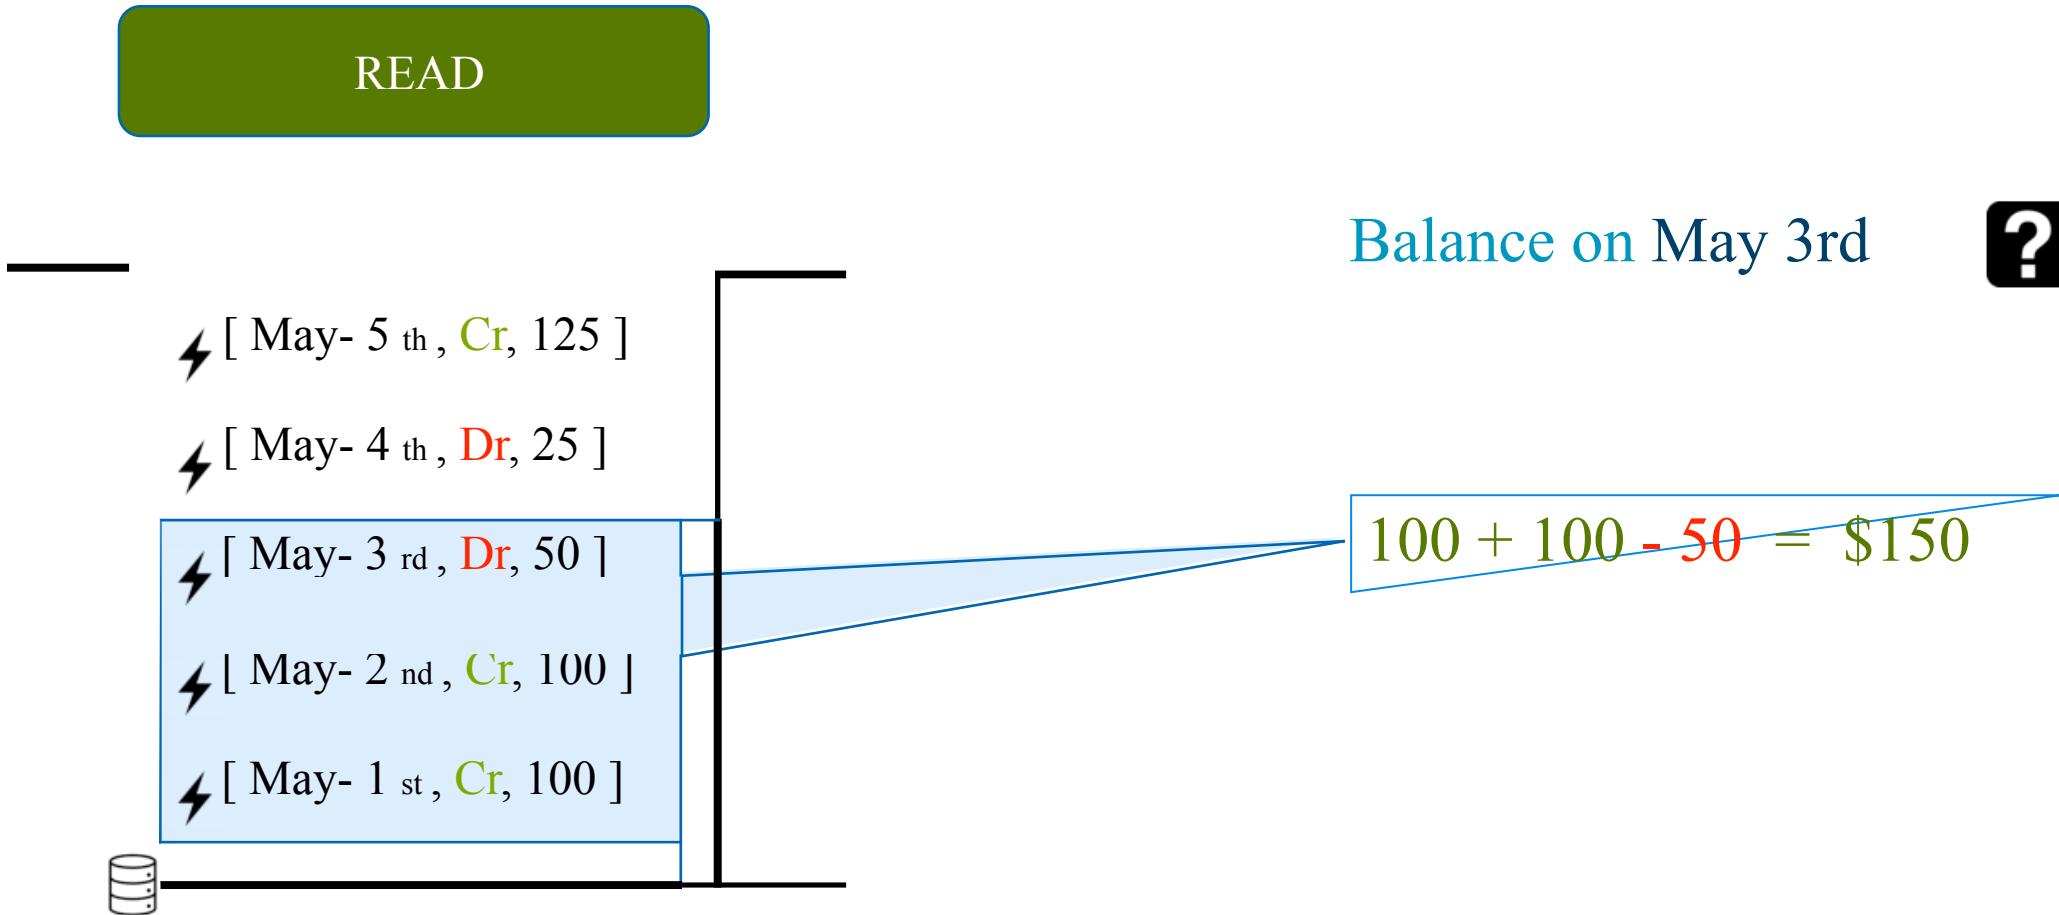
John's Bank Account			
			Database
100	May - 1 <sup>st</sup>	Deposit	\$100
101	May - 2 <sup>nd</sup>	Deposit	\$100
102	May - 3 <sup>rd</sup>	Withdraw	\$50
103	May - 4 <sup>th</sup>	Withdraw	\$25
104	May - 5 <sup>th</sup>	Deposit	\$125

READ

- 
- ⚡ [ May- 5<sup>th</sup>, Cr, 125 ]
  - ⚡ [ May- 4<sup>th</sup>, Dr, 25 ]
  - ⚡ [ May- 3<sup>rd</sup>, Dr, 50 ]
  - ⚡ [ May- 2<sup>nd</sup>, Cr, 100 ]
  - ⚡ [ May- 1<sup>st</sup>, Cr, 100 ]

## Example : Bank Transactions

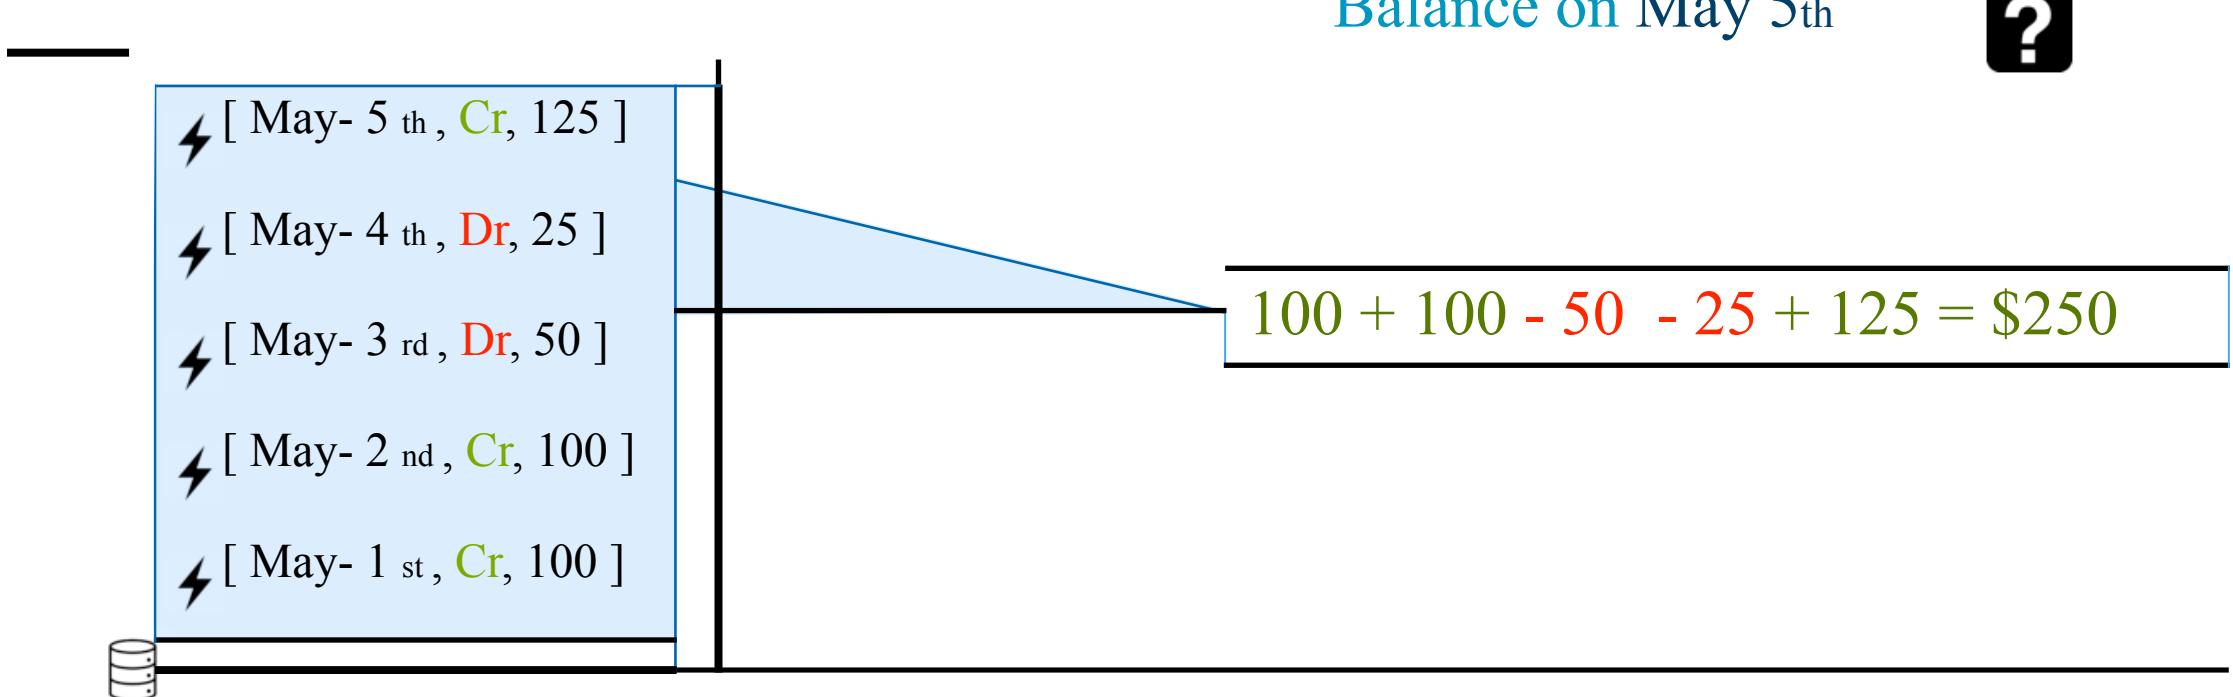
READ side can create the state for any date



## Example : Bank Transactions

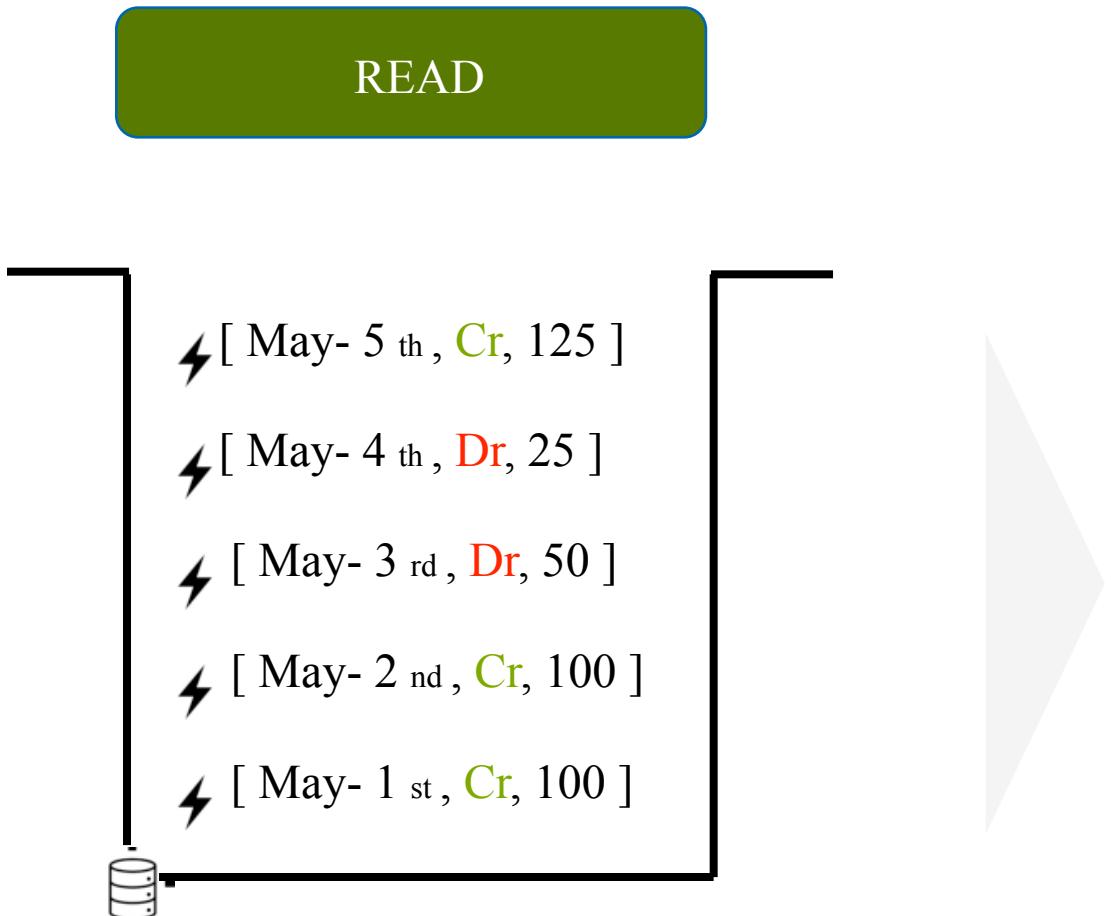
READ side can create the state for any date

READ



## Performance consideration

Recreating the state from events will lead to BAD performance

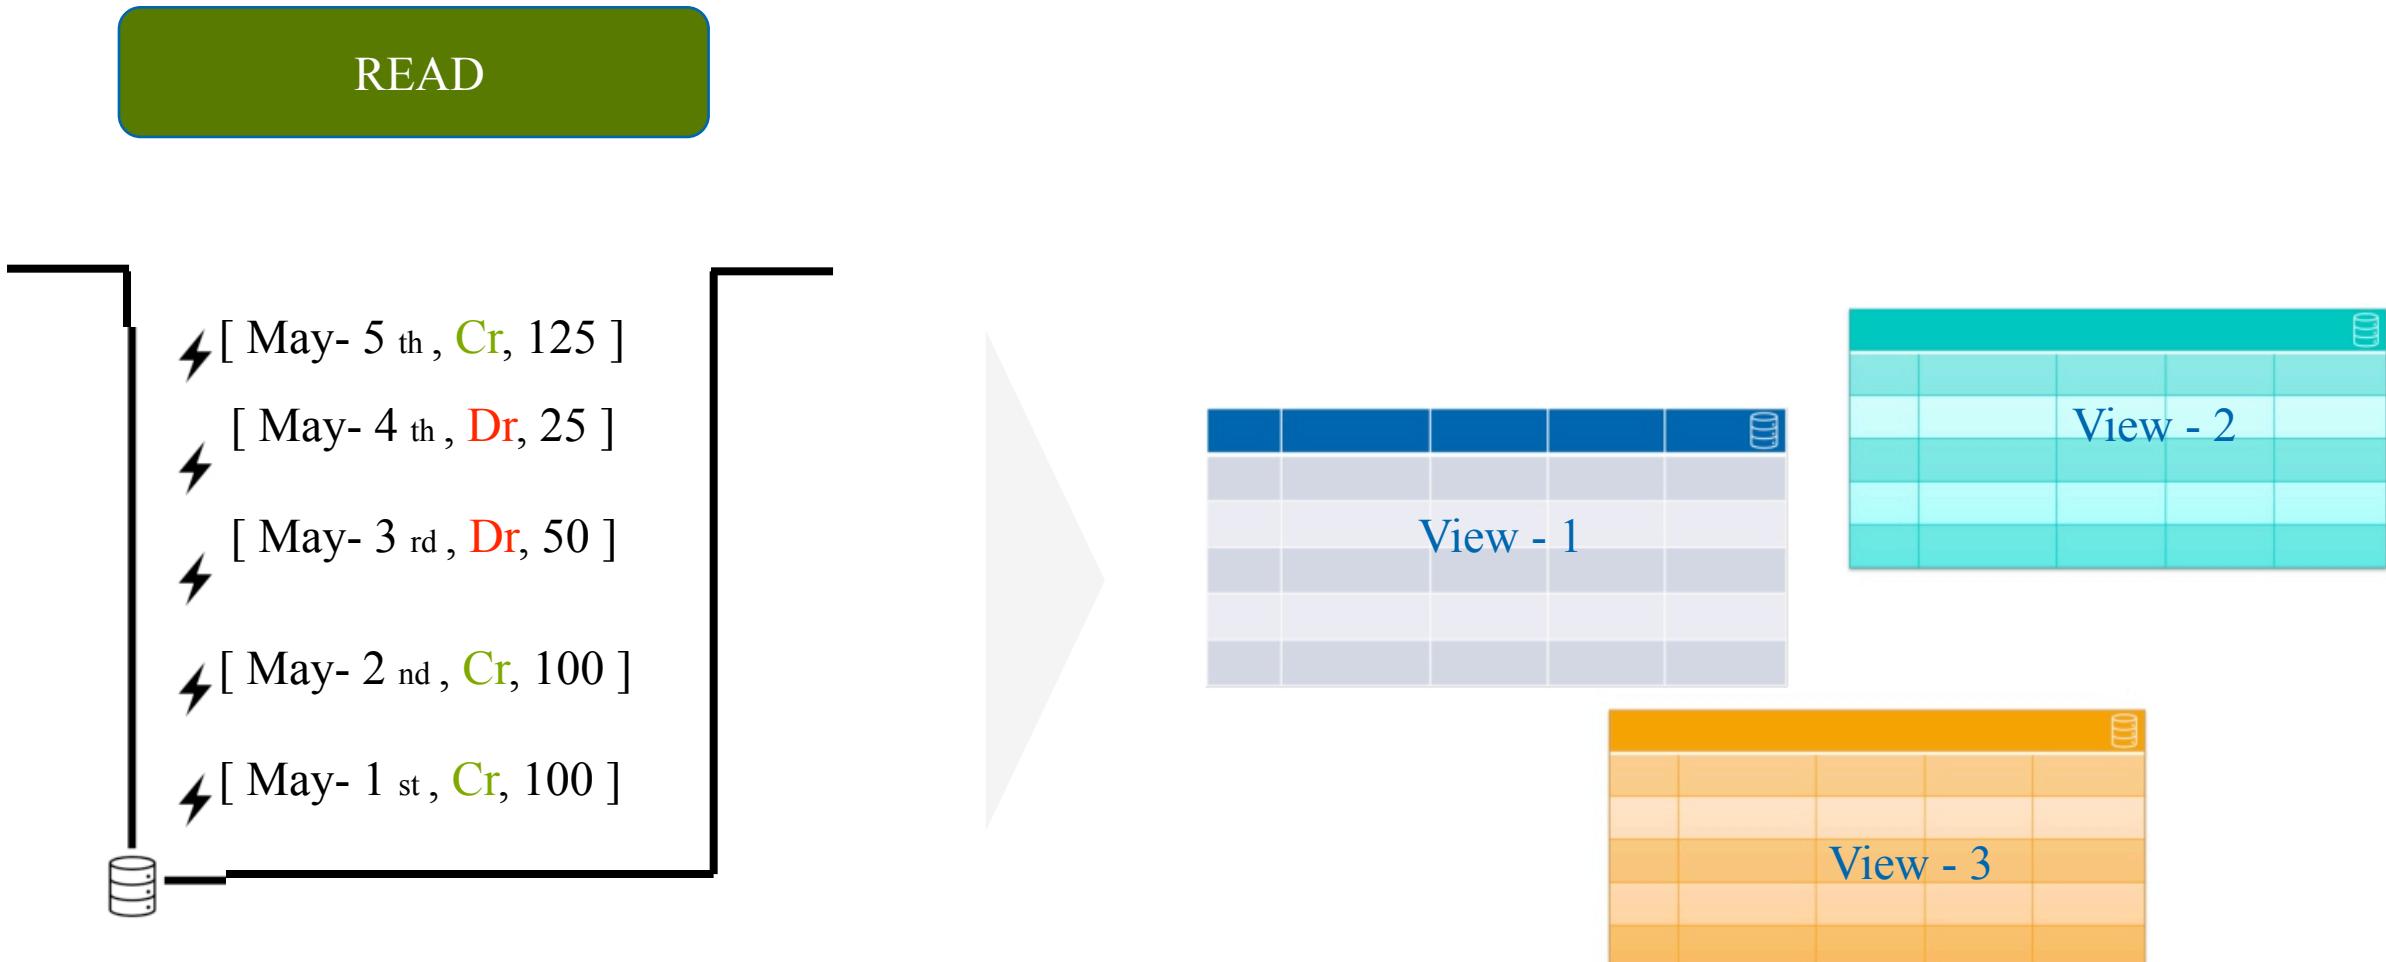


Current state in a separate Datastore  
*Lookup on Query*

		Database	
100	May 1 <sup>st</sup>	CR	100
101	May 2 <sup>nd</sup>	CR	200
102	May 3 <sup>rd</sup>	DR	50
103	May 4 <sup>th</sup>	DR	25
104	May 5 <sup>th</sup>	CR	125

## Performance consideration

Multiple data views may be created for optimizing queries



## Benefits

- 1 Event replay to create "*point in time state*"
- 2 Multiple read models | views
- 3 Accurate out of the box auditing
- 4 Simplified Reconciliation
- 5 Temporal & Complex historical queries

# Realization of Event Store

Traditional databases may be used

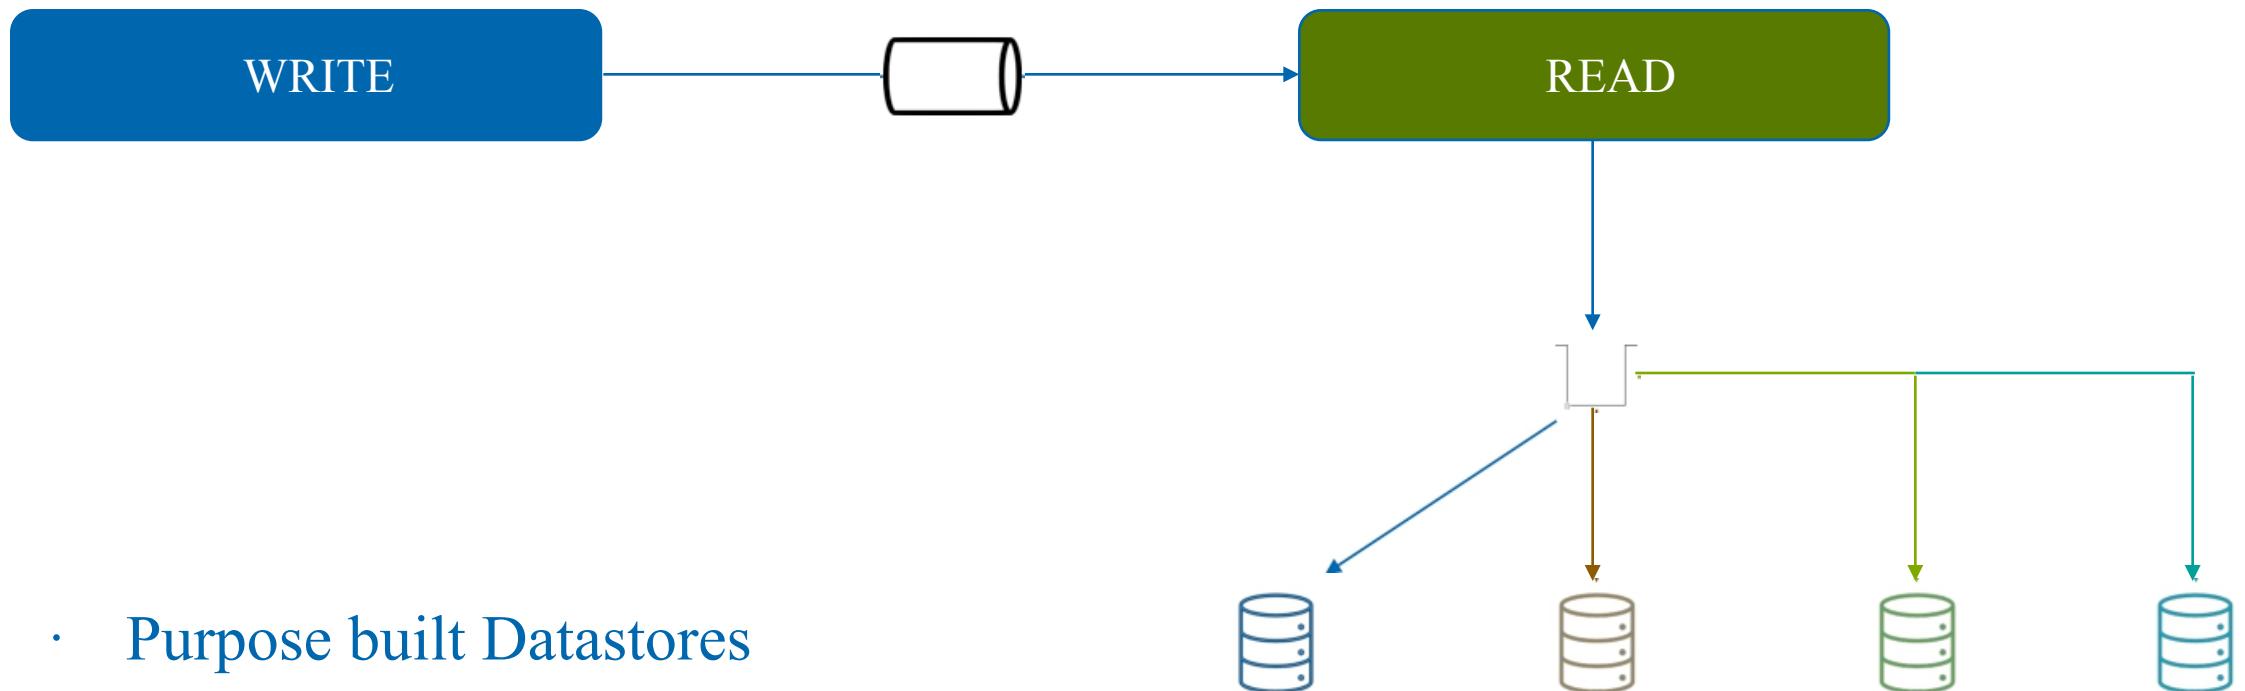
- RDBMS
- NoSQL Database
- Specialized databases



<https://www.eventstore.com>

# CQRS & Event Sourcing

Commonly used together



## When to use Event Sourcing?

Evaluate use case from the event sourcing benefits perspective

- ▶ 1 Event replay to create "*point in time state*"
- ▶ 2 Multiple read models | views
- ▶ 3 Accurate out of the box auditing
- ▶ 4 Simplified Reconciliation
- ▶ 5 Temporal & Complex historical queries



## Quick Review

Event -Sourced persistence systems persist all state changes

- Current state managed in a separate data store for performance
- Out of the box Audit, Reconciliation, Temporal queries

Events are stored in Event Stores that may be traditional DBs

# CQRS: Build ACL Subscriber

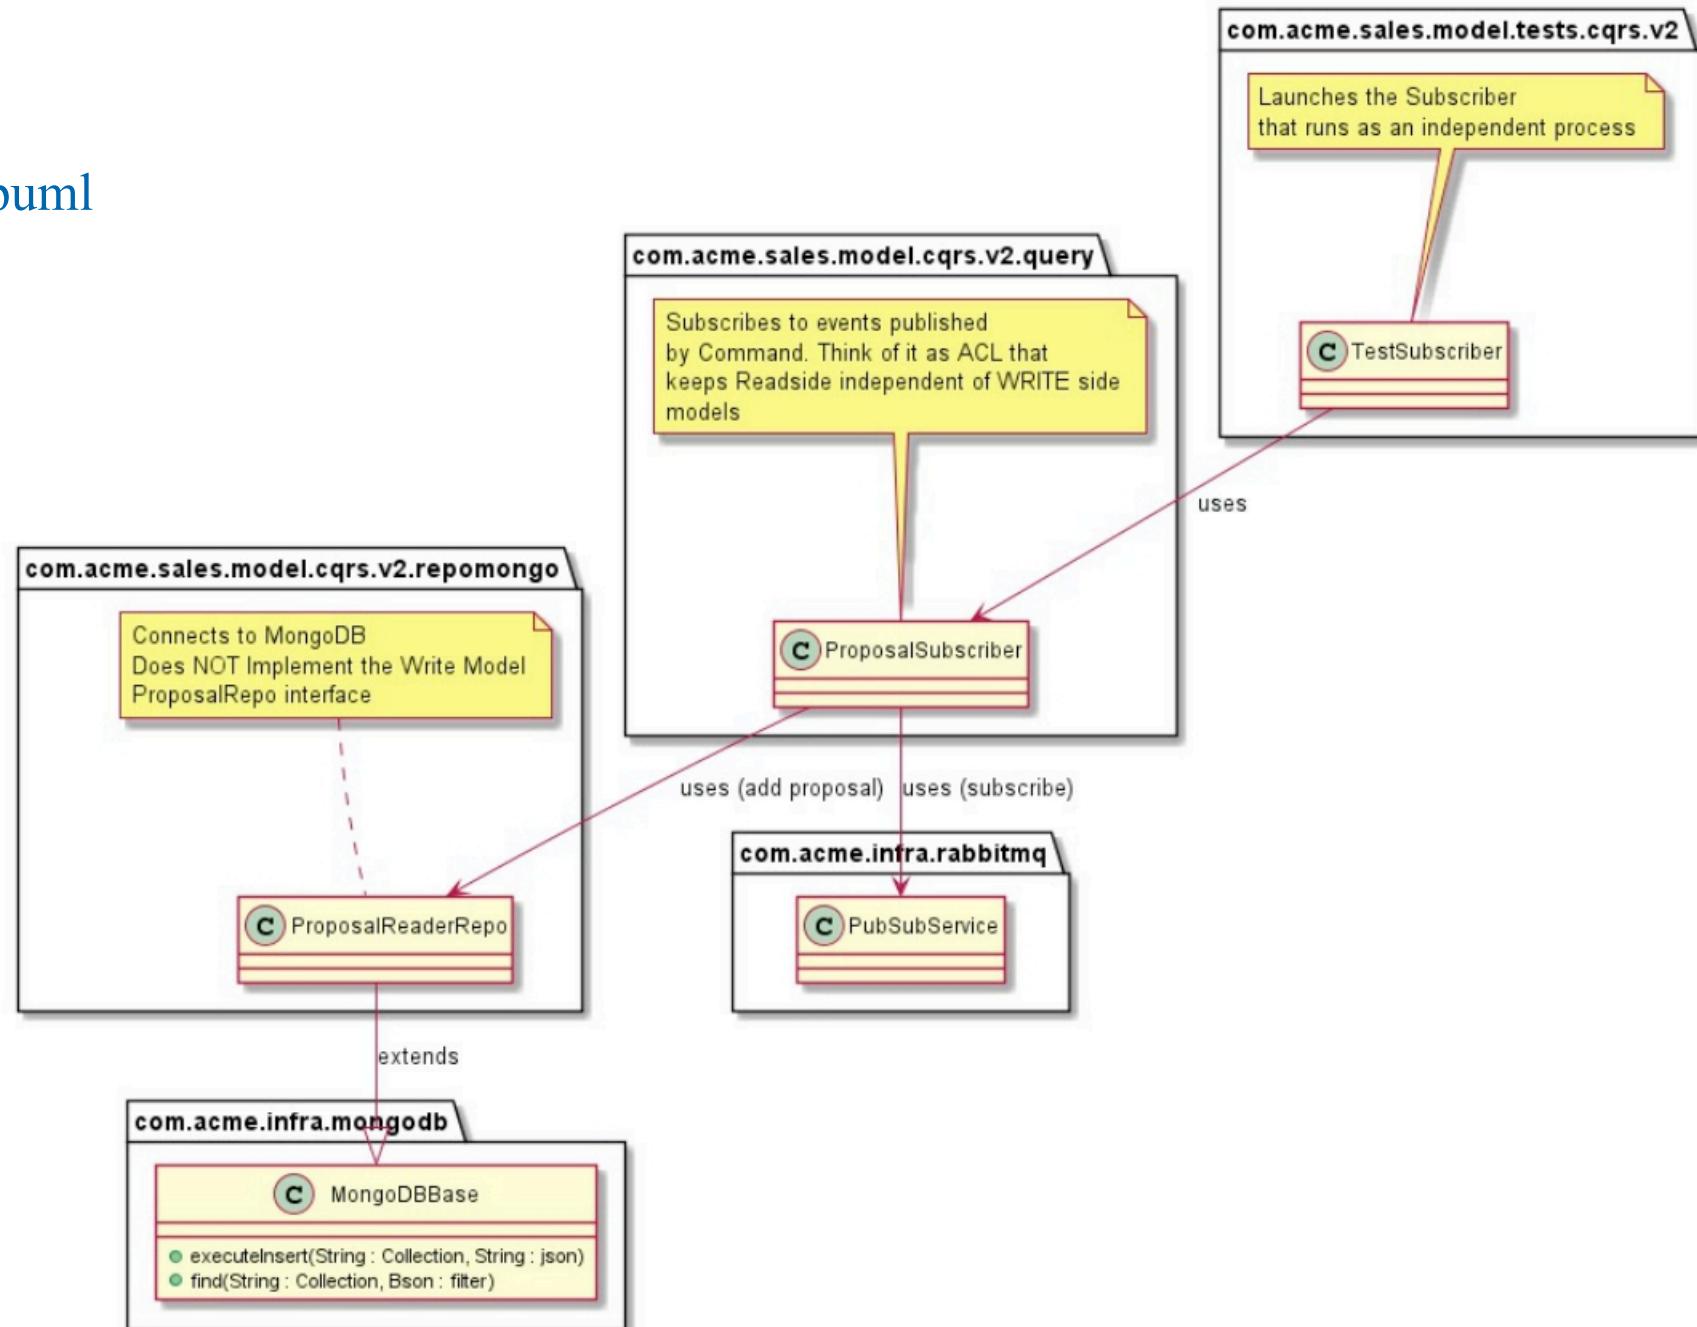
Stand alone Subscriber protects the READ side from corruption



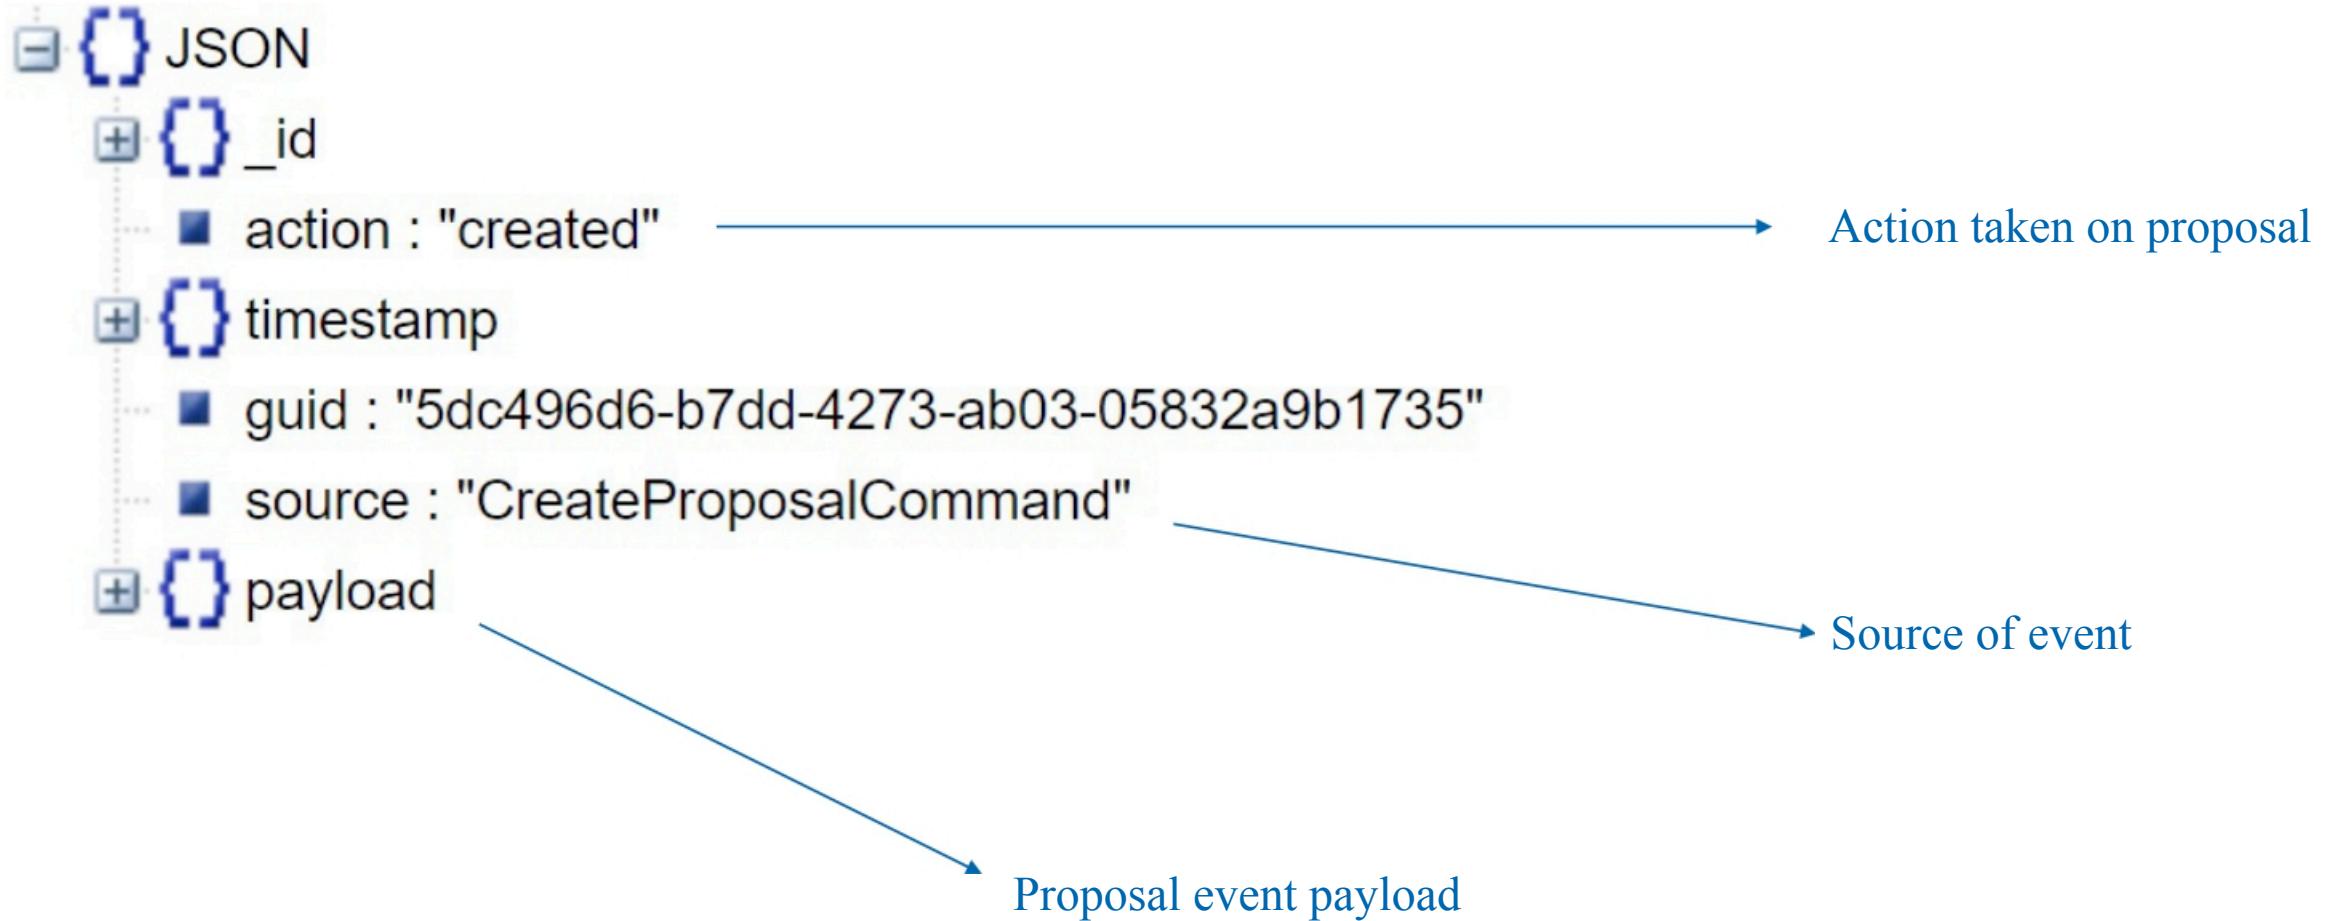
- 1 Walkthrough of ProposalReaderRepo
- 2 Walkthrough of ProposalSubscriber
- 3 Test the Proposal Subscriber

# READ Side Classes

uml/ cqrs/cqrs.v2.readside.class.puml



## Proposal Event



# Testing

1. Launch TestSubscriber
2. Execute the TestCommand
3. Check data in MongoDB

# CQRS: READ Side Query

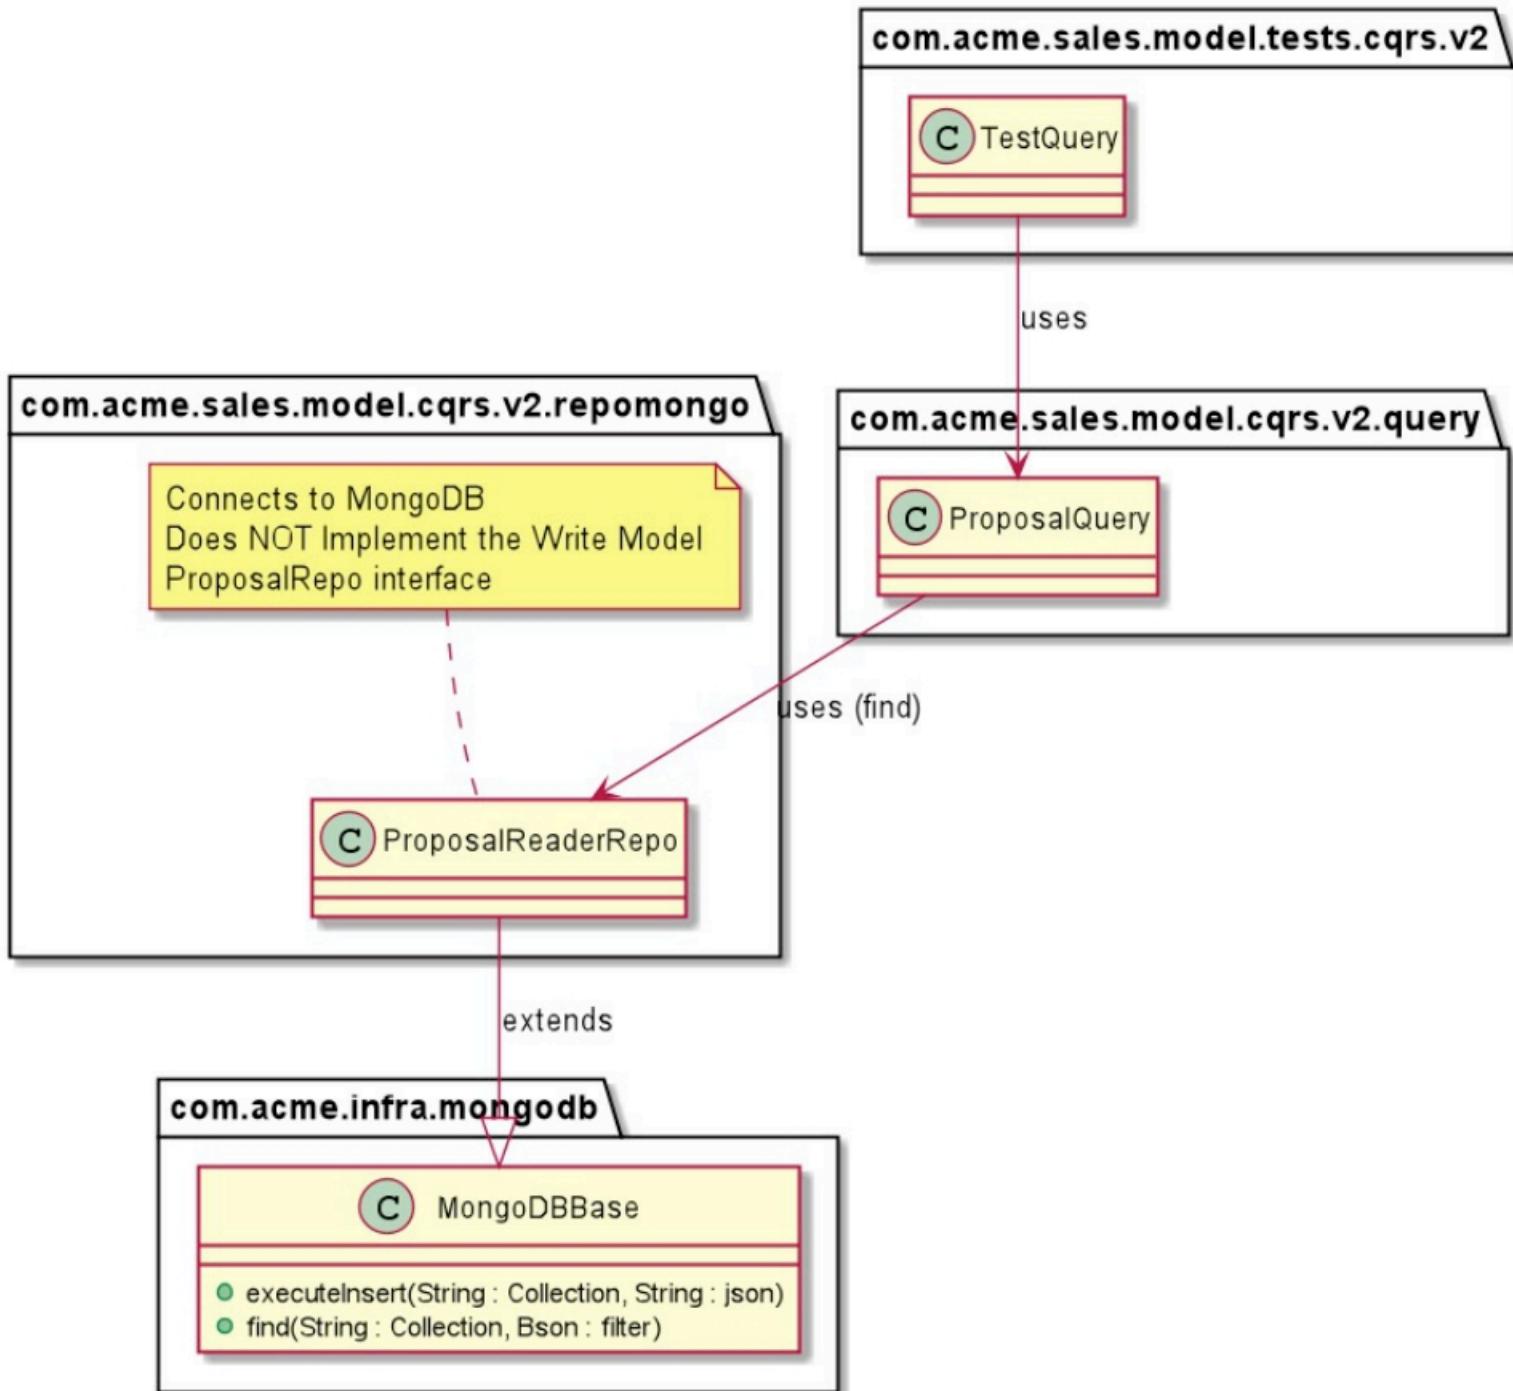
READ Side Query will use the MongoDB Data Store



- 1 Walkthrough of ProposalQuery
- 2 Test the Proposal Query

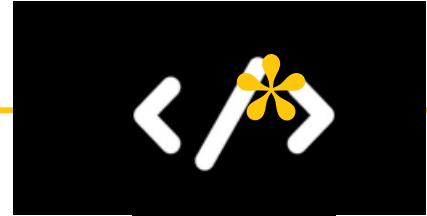
## READ side Query class

uml/ cqrs/cqrs.v2.readside.class.puml

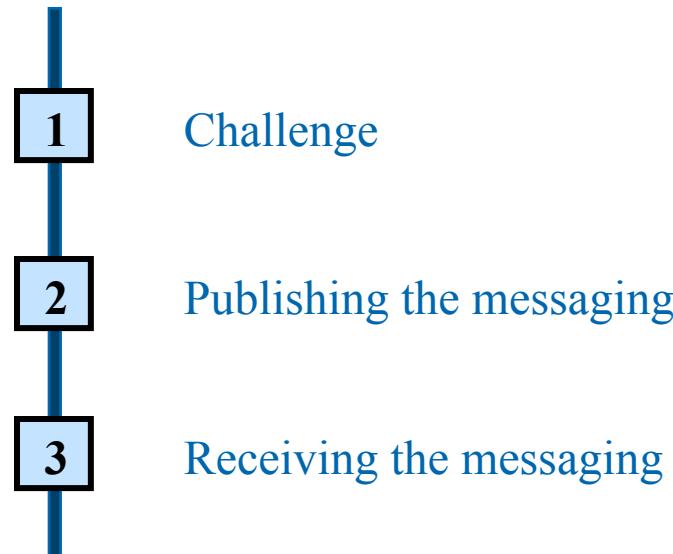


# Reliable Messaging | Dist Txn

What are they?

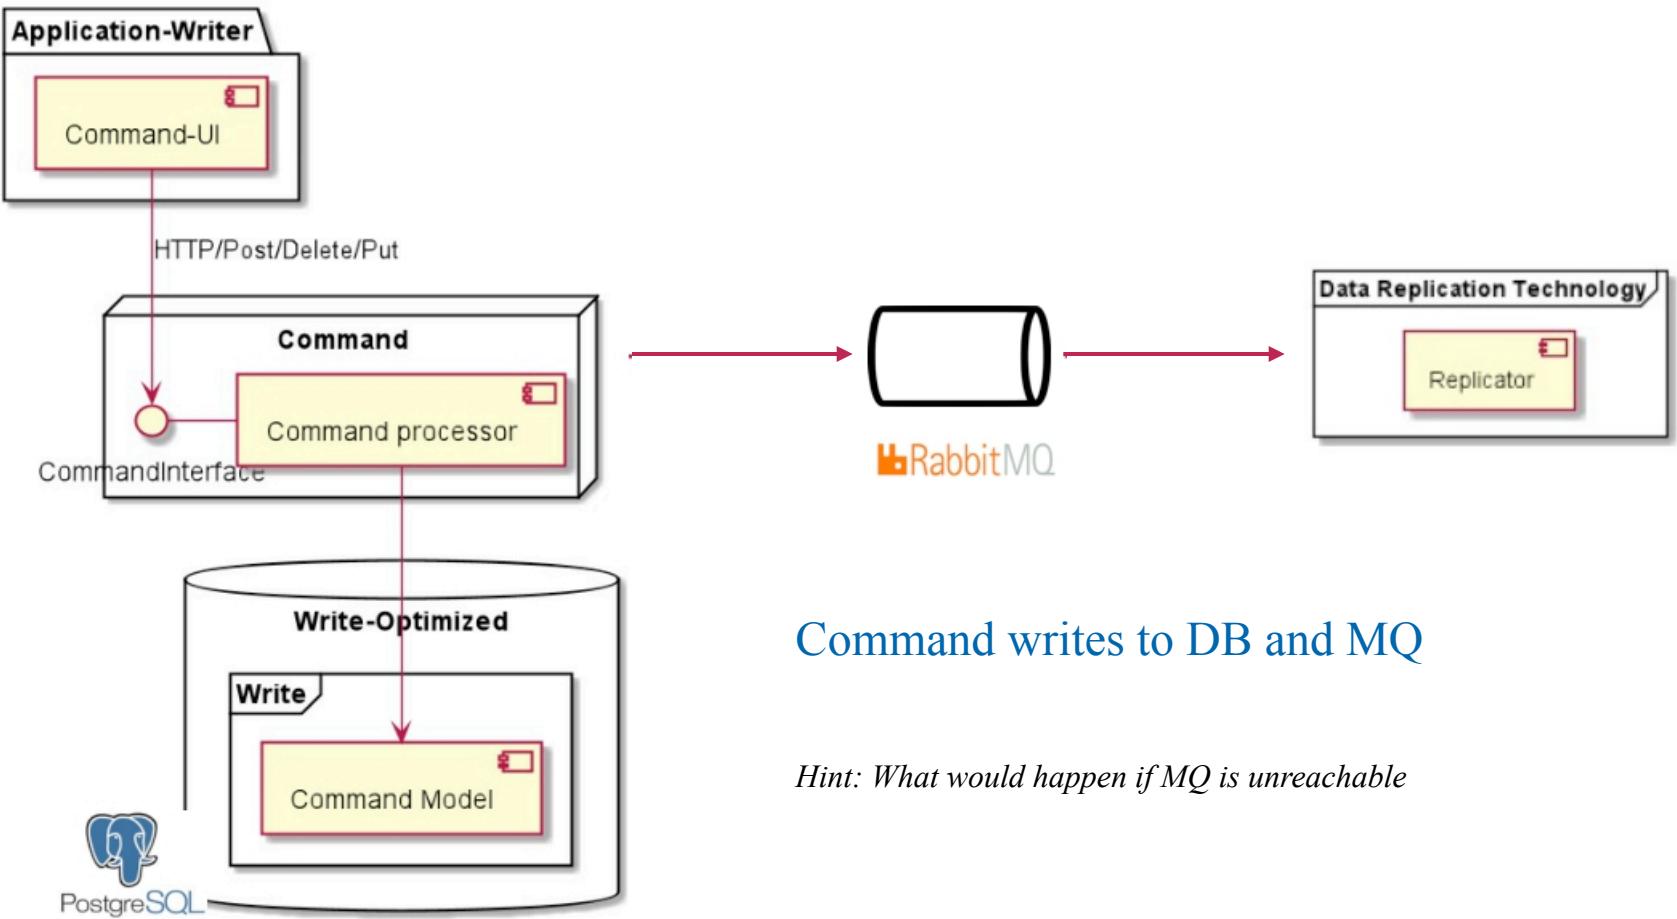


<https://vimeo.com/111998645>





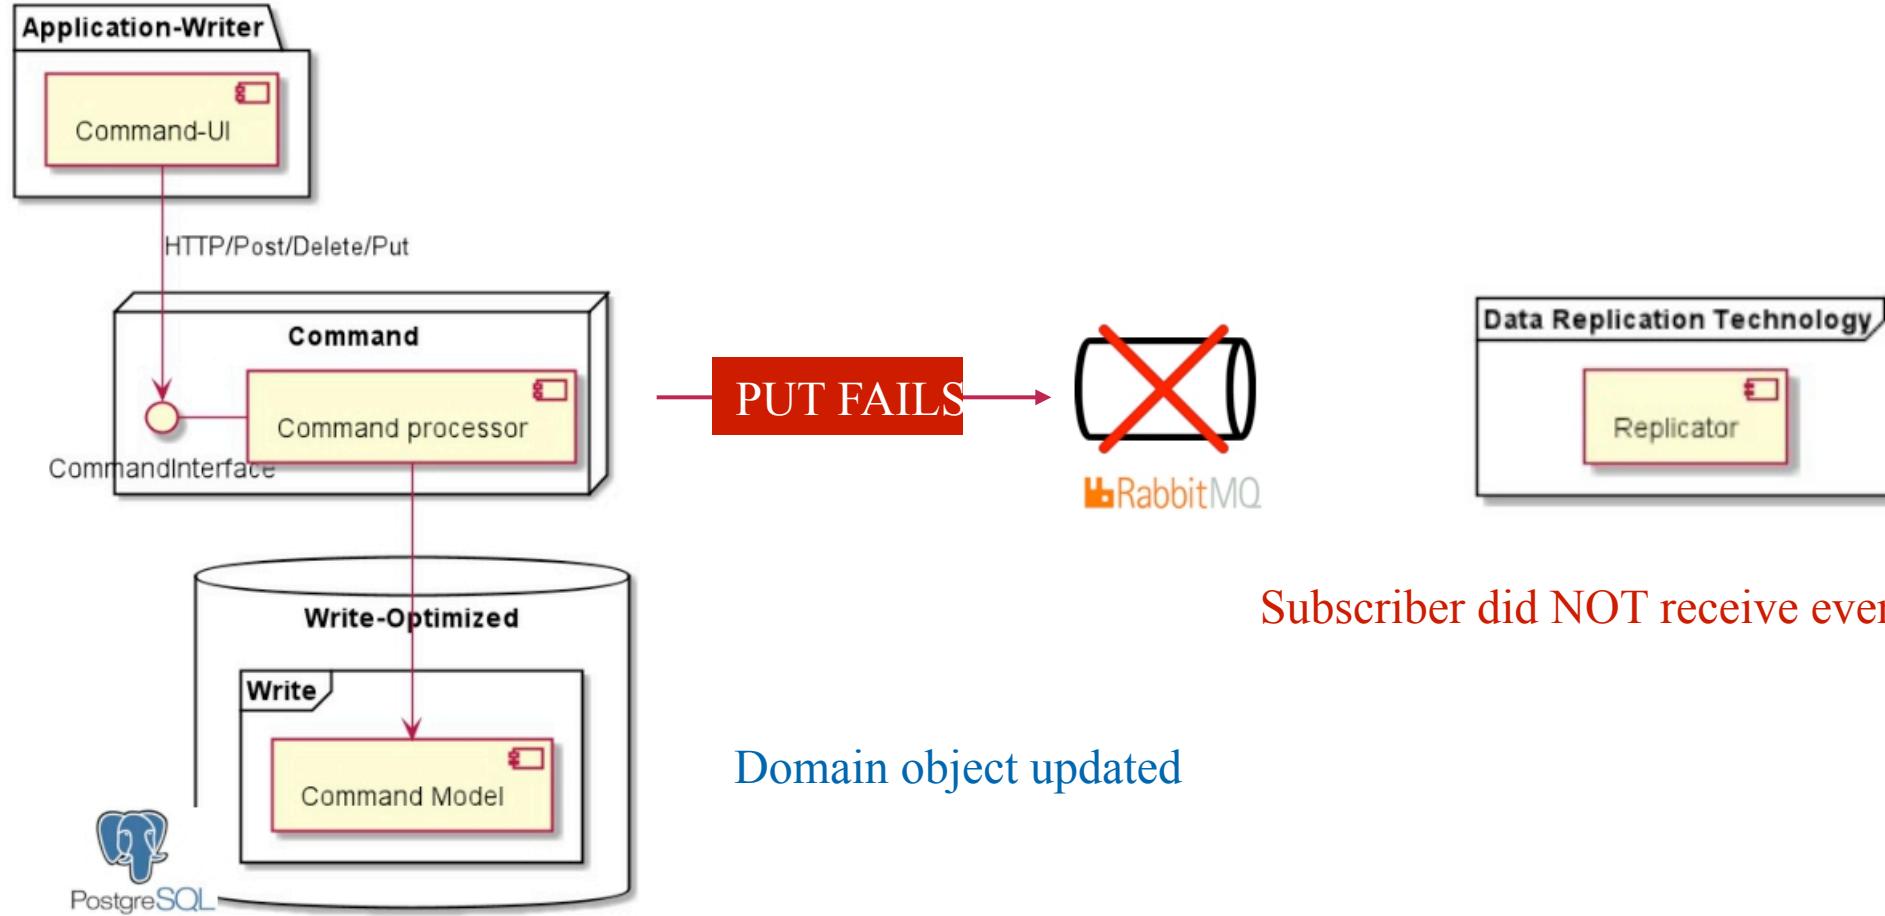
## Command Writes to DB & MQ



Do you see a problem here?

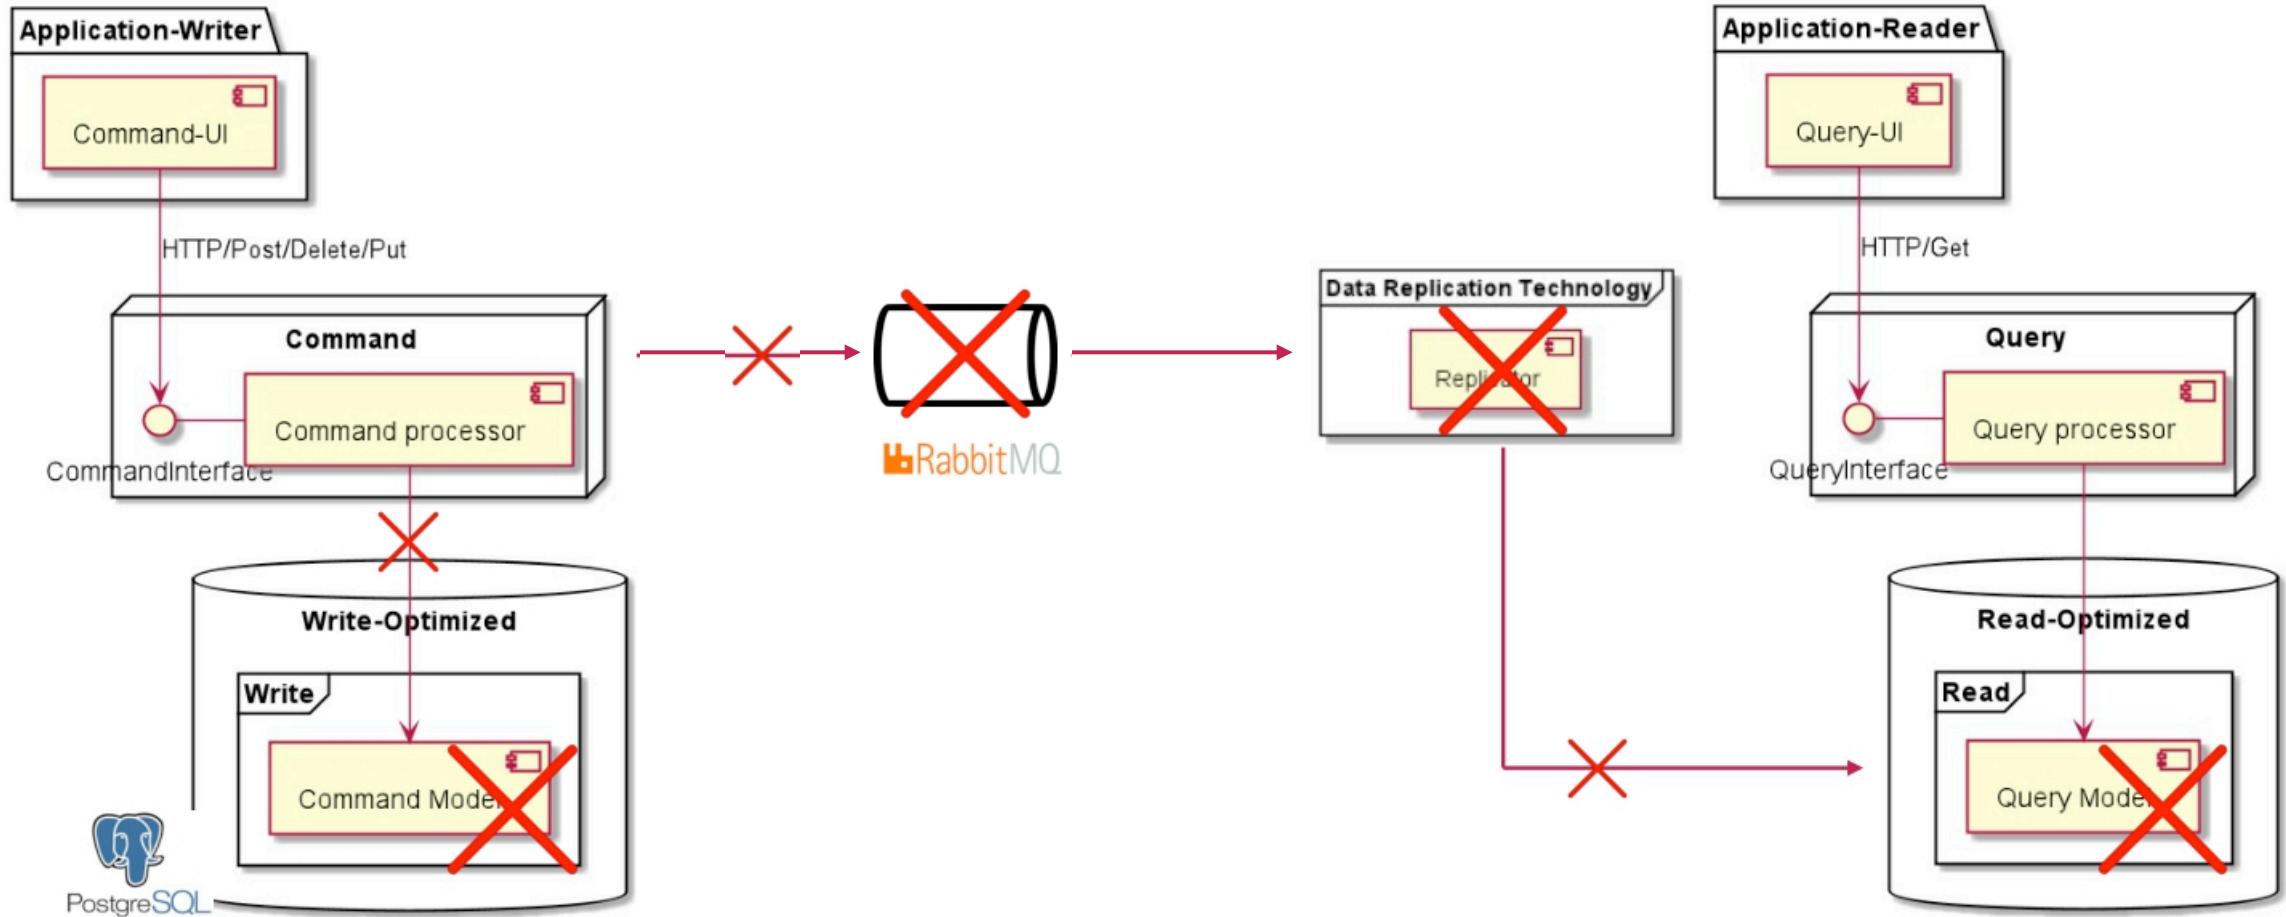
# Events will be LOST !!!!

UNLESS DB Write and MQ PUT are carried out in a single Transaction



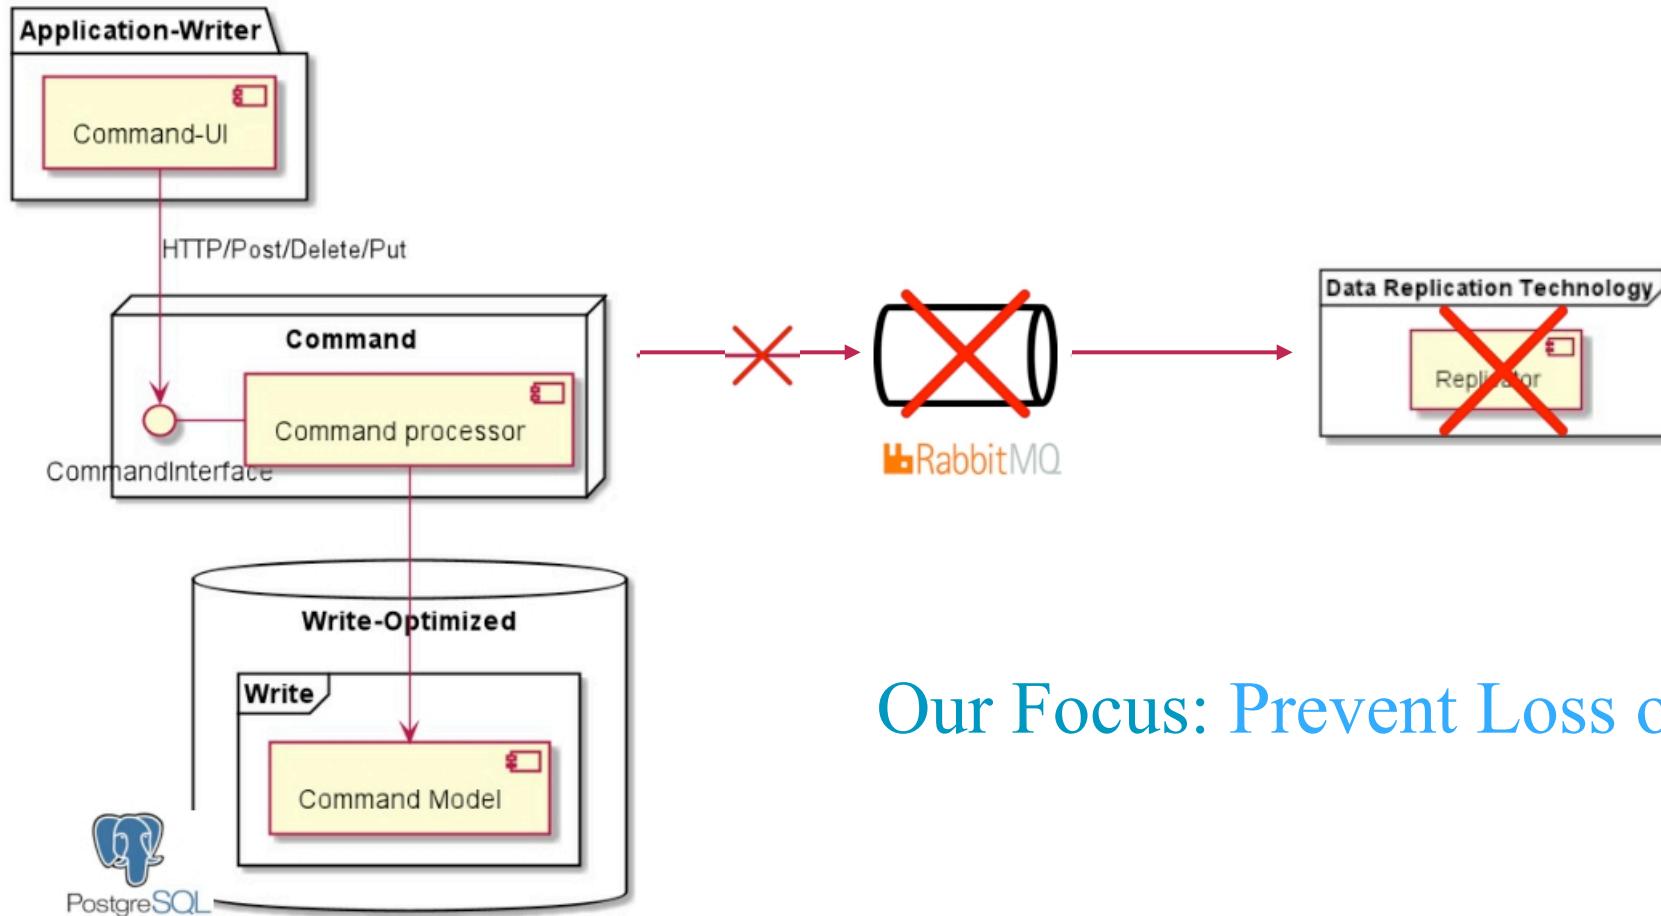
# Failures in Distributed systems

Just ASSUME that there will be failures & prepare



# Failures in Distributed systems

Just ASSUME that there will be failures & prepare



Our Focus: Prevent Loss of Events

## Solutions

DB Write & MQ Put in a single unit of work (transaction)

2 Phase Commit

Distributed algorithm that coordinates all the processes involved in distributed transaction

- Also known as eXtended Architecture or XA
- Transaction manager coordinates with resources

[https://en.wikipedia.org/wiki/Two-phase\\_commit\\_protocol](https://en.wikipedia.org/wiki/Two-phase_commit_protocol)

## Solutions

DB Write & MQ Put in a single unit of work (transaction)

2 Phase Commit

Distributed algorithm that coordinates all the processes involved in distributed transaction

- Many commonly used distributed technologies do NOT support it !!!!



[https://en.wikipedia.org/wiki/Two-phase\\_commit\\_protocol](https://en.wikipedia.org/wiki/Two-phase_commit_protocol)

## Solutions

DB Write & MQ Put in a single unit of work (transaction)

Reliable Messaging

Write aggregate data & event data in the database with a local transaction and replay the events against queueing system in a subsequent step.

# Reliable Messaging in action

An implementation that prevents loss of event messages



- 1 Walkthrough of ProposalReaderRepo
- 2 Walkthrough of ProposalSubscriber
- 3 Test the Proposal Subscriber



IT Lead

## Proposal Events CQRS Testing

Events are LOST when Rabbit MQ is down!!

Build Reliable Messaging to fix it !!!!

## Setup the DB tables : proposal events table

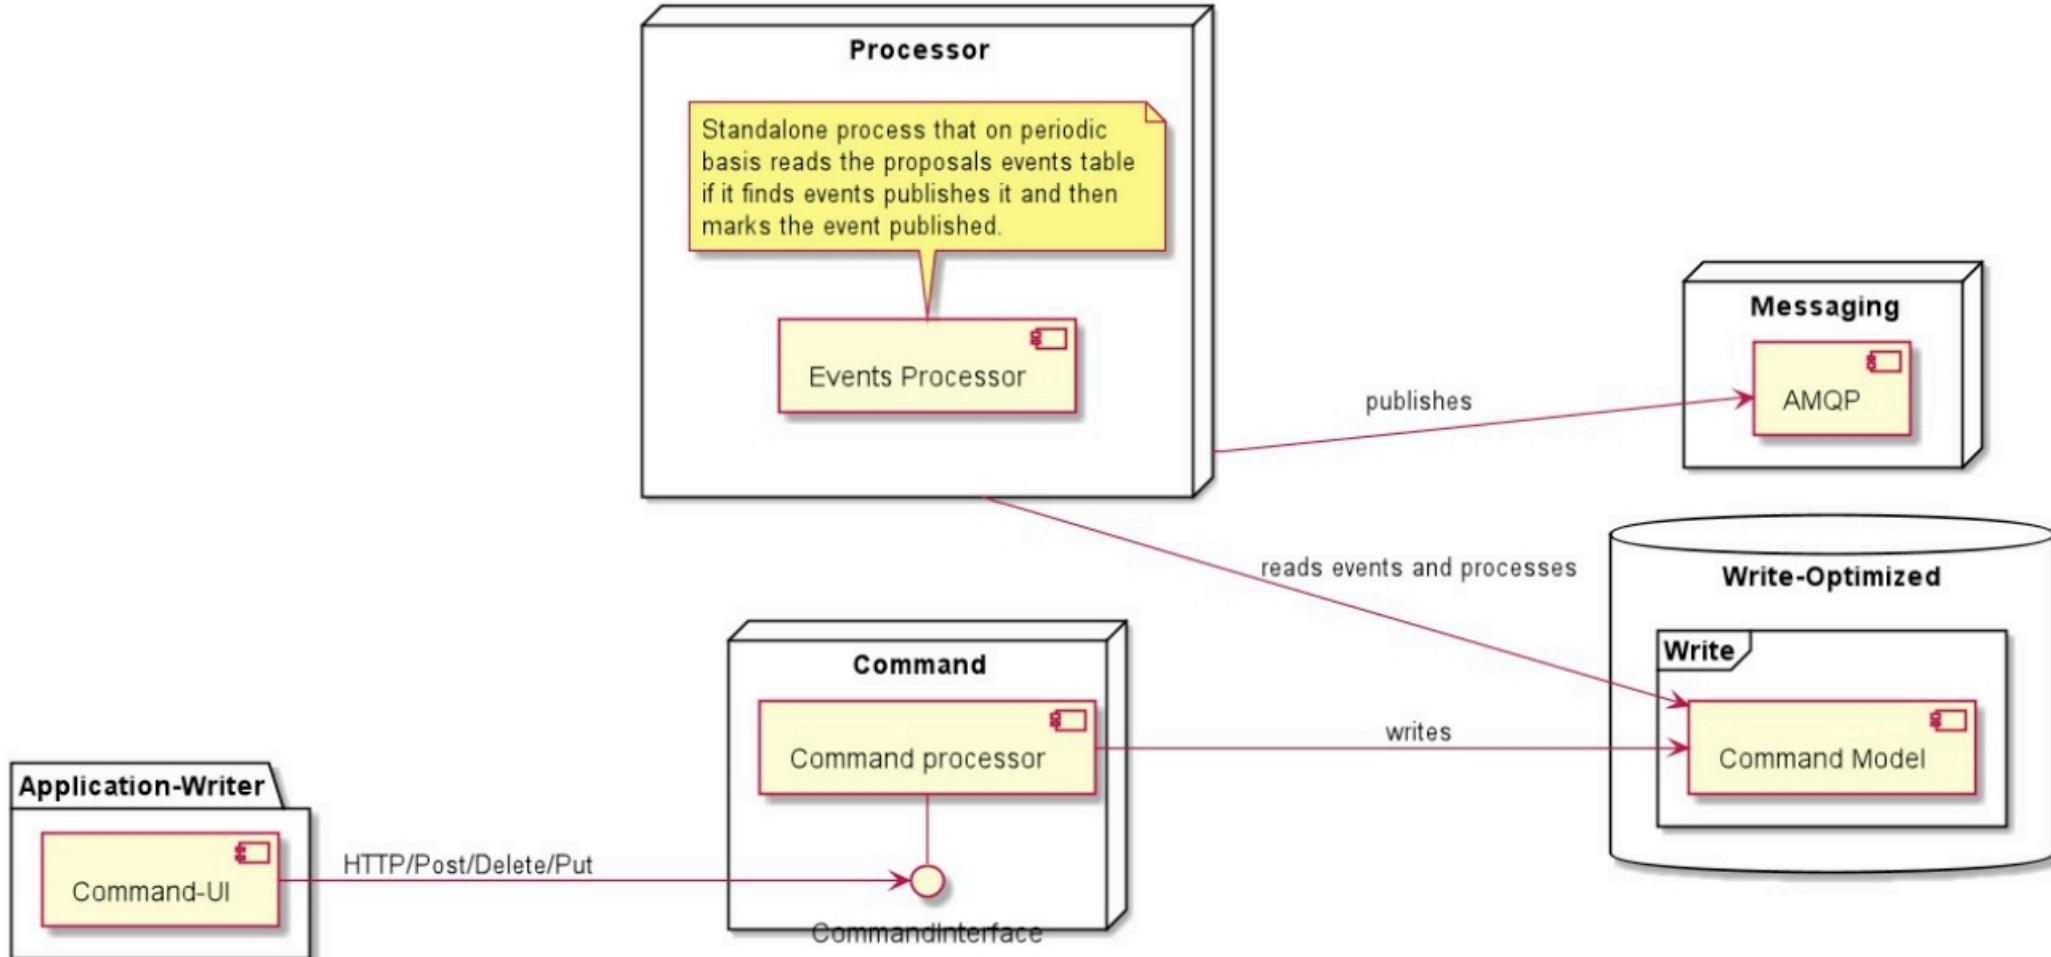
sql/reliable -messaging -ddl.sql

```
CREATE TABLE proposalevents (
    event_id      INT generated by default as identity,
    proposal_id   INT,
    event_guid    VARCHAR(36) not null,
    payload        VARCHAR(1024) not null,
    processed     BOOLEAN DEFAULT FALSE,
    created_date  TIMESTAMP DEFAULT now(),
    processed_date TIMESTAMP,
    PRIMARY KEY(event_id),
    CONSTRAINT fk_customer
        FOREIGN KEY(proposal_id)
            REFERENCES proposals(proposal_id)
);
```



# Write side Components

uml/ cqrs/cqrs.writeside.v3.component.puml



# Event processing classes

uml/ cqrs/cqrs.writeside.v3.class.puml

